



UNIVERSIDAD DE BUENOS AIRES  
FACULTAD DE CIENCIAS EXACTAS Y NATURALES  
DEPARTAMENTO DE COMPUTACIÓN

[your title here]

Tesis de Licenciatura en Ciencias de la Computación

Rafael Romani

Director: Santiago Cifuentes

Codirector: Santiago Figueira

Buenos Aires, 2026



**[YOUR TITLE HERE]**

[Acá va el resumen]

**Palabras claves:** [acá van las palabras clave]



[YOUR TITLE HERE]

[write your abstract here]

**Keywords:** [write your keywords here]



*[acá van los agradecimientos]*



## CONTENTS



# 1. INTRODUCTION

## 1.1 Uso y poder de los LLMs

In 2017, a group of researchers working at Google published a paper introducing a technology that revolutionized society. In *Attention is All You Need*, they presented the transformers, today's core of large language models (LLMs) and generative AI.

This architecture was build intended as text sequence-to-sequence model for machine translation but today has applications in various fields as natural language processing, computer vision or speech recognition.

The first model implementing transformers open for public use was ChatGPT by OpenAI. It made its debut in November 2022 and by December 2024 it reported 300 million global weekly users. Now there are many models competing for this market as Claude, from Anthropic, or Gemini from Google.

These models can be all be grouped under the generative AI tools. This is a technology that is arguably the most rapid technology adopted by society. In the second half of 2025, roughly one in six people worldwide were using generative AI tools. [not sure how to cite this https://www.microsoft.com/en-us/research/wp-content/uploads/2026/01/Microsoft-AI-Diffusion-Report-2025-H2.pdf](https://www.microsoft.com/en-us/research/wp-content/uploads/2026/01/Microsoft-AI-Diffusion-Report-2025-H2.pdf)

LLMs are being used not only as chatbots for asking small questions, but are shaping the industry and future. For example, [?] found that more than 17% of the Computer Science's papers published between January 2020 and February 2024 on the *arXiv*, *bioRxiv*, and *Nature* portfolio journals had sentences heavily edited by ChatGPT.

### Missing usage in industry

So naturally, the question about their computing power pops up. Are this kind of models really capable of doing the stuff for what they are being used?

- well-suited to parallelism enables transformer models to take full advantage of the power and speed offered by GPUs during both training and inference. This possibility, in turn, unlocked the opportunity to train transformer models on unprecedentedly massive datasets through self-supervised learning.
- chatgpt, claude, gemini,
- every day users
- number of users?
- use in the industry?
- other uses besides the chatbots?

**1.2** Pero qué es un LLM?

**1.3** Formalizaciones previas y resultados que ignoran los aspectos probabilísticos

**1.4** Objetivos y contribuciones

## 2. THEORETICAL FRAMEWORK

### 2.1 Transformers

In this work we are going to analyze transformers as language recognizers. A transformer is a machine that works over an alphabet  $\Sigma$ . It has only one tape from where it reads the input and in an autoregressive way writes in it more tokens. There are two distinguished tokens called stopping symbols. Once the transformer writes one of those it halts. We say that a transformer accepts a word if when it halts the stopping symbol printed is the accepting one ( $q_{accept}$ ). Correspondingly, a transformer rejects a word if the stopping symbol printed is the rejecting one ( $q_{reject}$ ).

In each step the transformer maps the content of the tape to a probability distribution over the alphabet. Depending on this distribution the transformer prints the next token. The transformer iterates this process until it halts. These iterations are called **chain of thought**.

The process of generating the next token consists of two separated steps. First a probability distribution over the alphabet for the token is computed by a process called distribution generation. After that, it is necessary to decide the specific token by a process called token selection. These two algorithms together is what defines the transformer. During this work we will analyze and compare two different options for token selection.

For a more realistic analysis of the transformers, we don't allow an infinite tape. The transformers studied in this work, for each input size  $n$ , will have a tape of size  $\text{budget}(n) + n$ . It's important to remark that, since in every step of the transformer a symbol is written, the maximum number of steps a transformer can do for an input of length  $n$  is  $\text{budget}(n)$ .

The process of distribution generation is dependent on many parameters. These parameters, in the real models, are the ones calculated during the training. We refer to them as  $\theta$ . During this work we assume that the parameters force the transformer to halt exactly when the budget runs out. We consider the parameters as part of the definition of a transformer.

We note  $\text{TF}(\alpha)$  to the token outputted by one step of the transformer TF when the tape content is  $\alpha \in \Sigma^*$ . The autoregressive token generation is defined recursively. The token outputted after  $i$  steps is  $\text{TF}^i(\alpha) := \text{TF}^{i-1}(\alpha \text{TF}(\alpha))$  with base case  $\text{TF}^1(\alpha) := \text{TF}(\alpha)$ . In other words, the output of the  $i$ th step is  $\sigma_i = \text{TF}^i(\alpha) = \text{TF}(\alpha \sigma_1 \dots \sigma_{i-1})$  where  $\sigma_j \in \Sigma$  for all  $j$ .

We note  $\text{TF}^*(\alpha)$  for the token written by the transformer when it halts, i.e. at the end of the computation, when the input is  $\alpha$ . As was said before, the transformer only halts when it prints a stopping symbol ( $q_{accept}$  or  $q_{reject}$ ). Therefore,  $\text{TF}^*(\alpha) \in \{q_{accept}, q_{reject}\}$ .

As stated before, the **distribution generation** is a mapping of the content of the tape to a probability distribution over  $\Sigma$ . There are different algorithms for this process in the literature. All the transformers considered in this work use the one presented by [?]. This algorithm consists of four parts: a token embedding layer (TE), a position encoding layer (PE), an output linear layer (OUTPUT), and a stack of  $L$  identical decoder layers, where  $L$  is also called the depth of the model. Each decoder layer has two sub-layers: a multi-head self-attention layer (ATTN) and a position-wise fully-connected feed-forward network

Write an appendix about this decision. We can always add an extra step or artifact forcing

add somewhere def of softmax

(FF). Each part and layer has its own trainable parameters. Then, we can define the distribution generation algorithm by its parameters  $\theta = \{\theta_{\text{TE}}, \theta_{\text{PE}}, \{\theta_{\text{ATTN}}^l, \theta_{\text{FF}}^l\}_{l=0}^{L-1}, \theta_{\text{OUTPUT}}\}$ . Let's define each component.

**Token Embedding:** It's a mapping from  $\Sigma$  to vectors in  $\mathbb{R}^d$ . This mapping is defined by  $\theta_{\text{TE}} \in \mathbb{R}^{d \times |\Sigma|}$ . For all  $\sigma \in \Sigma$ , we note  $\theta_{\text{TE}}(\sigma) \in \mathbb{R}^d$  to the mapping of token  $\sigma$ .

**Position Encoding:** It's a mapping from the positions of the tape to vectors in  $\mathbb{R}^d$ . This mapping is defined by  $\theta_{\text{PE}} \in \mathbb{R}^{d \times \text{budget}(n)+n}$ . For all position of the tape  $1 \leq i \leq \text{budget}(n) + n$  we note  $\theta_{\text{PE}}(i) \in \mathbb{R}^d$  to the mapping of position  $i$ .

**Self Attention Layer:** It's a mapping from  $(\mathbb{R}^d)^a$  to  $(\mathbb{R}^d)^a$ . Given parameters  $\theta_{\text{ATTN}} = (W_Q, W_K, W_V, W_O) \in \mathbb{R}^{d \times d} \times \mathbb{R}^{d \times d} \times \mathbb{R}^{d \times d} \times \mathbb{R}^{d \times d}$ , the layer it's defined by the algorithm ???. This layer is defined only as a single head attention layer because we can simulate 1 multi-head attention layer with any constantly many heads with multiple single-head attention layers. [?]

---

**Algorithm 1** Causal Self-Attention, ATTN

---

**Require:** Parameter  $\theta_{\text{ATTN}} = (W_Q, W_K, W_V, W_O)$ , Input embedding  $h = (h_1, \dots, h_a) \in \mathbb{R}^d \times \dots \times \mathbb{R}^d$ .

**Ensure:** Output embedding  $h' = (h'_1, \dots, h'_a) := \text{ATTN}_{\theta_{\text{ATTN}}}(h_1, \dots, h_a)$ .

- 1:  $q_i := W_Q h_i, k_i := W_K h_i, v_i := W_V h_i, \forall i \in [1, \dots, a]$
  - 2:  $s_i := \text{softmax}(\langle q_i, k_1 \rangle, \dots, \langle q_i, k_i \rangle) \parallel (0, \dots, 0)$
  - 3:  $h'_i := W_O \sum_{j=1}^a (s_i)_j v_j$
- 

**Feed Forward Network:** It's a mapping from  $\mathbb{R}^d$  to  $\mathbb{R}^d$ . Given  $\theta_{\text{FF}} = (W_1, b_1, W_2, b_2) \in \mathbb{R}^{d \times d} \times \mathbb{R}^d \times \mathbb{R}^{d \times d} \times \mathbb{R}^d$ , it's defined as  $\text{FF}_{\theta_{\text{FF}}}(h) := W_2 \text{relu}(W_1 h + b_1) + b_2$ , where  $\text{relu} : \mathbb{R}^d \rightarrow \mathbb{R}^d$  replaces each position  $x_i$  by  $\max(0, x_i)$ .

**Output Layer:** It's a mapping from  $\mathbb{R}^d$  to a probability distribution over  $\Sigma$ . Given  $\theta_{\text{OUTPUT}} \in \mathbb{R}^{|\Sigma| \times d}$ , it's defined as  $\text{OUTPUT}_{\theta_{\text{OUTPUT}}}(h) := \text{softmax}(\theta_{\text{OUTPUT}} h)$ .

All these parts come together for the distribution generator process in the algorithm ???.

---

**Algorithm 2** Distribution generator

---

**Require:** Transformer parameter  $\theta = (\theta_{\text{TE}}, \theta_{\text{PE}}, \{\theta_{\text{ATTN}}^l, \theta_{\text{FF}}^l\}_{l=0}^{L-1}, \theta_{\text{OUTPUT}})$  and tokens of the tape  $\alpha \in \Sigma^*$ .

**Ensure:** Output distribution  $p_\theta(\cdot | \alpha)$ .

- 1:  $h_i^{(0)} \leftarrow \theta_{\text{TE}}(\alpha_i) + \theta_{\text{PE}}(i), \forall 1 \leq i \leq |\alpha|$
  - 2: **for**  $l = 0, \dots, L - 1$  **do**
  - 3:    $(h_1^{(l+0.5)}, \dots, h_{|\alpha|}^{(l+0.5)}) \leftarrow (h_1^{(l)}, \dots, h_{|\alpha|}^{(l)}) + \text{ATTN}_{\theta_{\text{ATTN}}^{(l)}}(h_1^{(l)}, \dots, h_{|\alpha|}^{(l)})$
  - 4:    $h_i^{(l+1)} \leftarrow h_i^{(l+0.5)} + \text{FF}_{\theta_{\text{FF}}^{(l)}}(h_i^{(l+0.5)}), \forall 1 \leq i \leq |\alpha|$
  - 5: **end for**
  - 6:  $p_\theta(\cdot | \alpha) \leftarrow \text{OUTPUT}_{\theta_{\text{OUTPUT}}}(h_{|\alpha|}^{(L)})$
- 

The second component of a transformer is the **token selector**. This algorithm consumes the probability distribution computed by the distribution generator and returns a token.

Lets  $p : \Sigma \rightarrow [0, 1]$  be the output of the distribution generator. In this work we will study two selectors. The first one always returns the token with the largest probability

(argmax over  $p$ , where  $p$  is the output of the distribution generator). The other, works non-deterministically. It returns the token  $\sigma \in \Sigma$  with probability  $p(\sigma)$ .

I forgot to talk about finite representation. Define good symbols for  $\infty$  y  $-\infty$ . Properties i use later on that i should explain here:

- $\exp(-\infty) = 0$
- $\exp(\infty) = \infty$
- $-\infty * \infty = -\infty$
- $\infty * \infty = \infty$
- $x + 0 = x$
- $x - x = 0$

**Definition 2.1.1** (Deterministic transformer). We call a transformer deterministic when the token selector is the one taking argmax over the distribution. We say that a family of deterministic transformers  $\{\text{TF}_n\}_{n \in \mathbb{N}}$  recognize a language  $\mathcal{L}$  when:

$$\alpha \in \mathcal{L} \iff \text{TF}_{|\alpha|}^*(\alpha) = q_{\text{accept}}$$

As mentioned in chapter 1, there is a theoretical gap in the area. *está bien dicho theoretical gap?* A study of non-determinism in transformers is missing.

When TF is a deterministic transformer,  $\text{TF}^i(\alpha)$  is an exact symbol for all  $\alpha \in \Sigma^*$  and  $i \in \mathbb{N}$ . In counterpart, when the token selection process of the transformer samples the distribution, it becomes a random variable. Therefore, we introduce the following definition.

**Definition 2.1.2** (Probabilistic transformer). We call probabilistic transformers to the transformers that sample the distribution for selecting the next token. We say that a family of probabilistic transformers  $\{\text{TF}_n\}_{n \in \mathbb{N}}$  recognize a language  $\mathcal{L}$  when:

$$\alpha \in \mathcal{L} \iff \Pr[\text{TF}_{|\alpha|}^*(\alpha) = q_{\text{accept}}] > 2/3$$

## 2.2 Transformer Complexity classes

*Maybe rethink this title?*

As discussed in the first chapter, there are different resources to bound in the transformers to limit their expressive power. In this thesis, we will analyze the effect of bounding the number of steps of chain of thought that the transformers are allowed to make before answering. With that in mind, we define the following complexity classes.

Given two functions  $T(n)$  and  $d(n)$  we define the complexity class  $\text{CoT}[T(n), d(n)]$ . Informally, a language  $\mathcal{L}$  is in  $\text{CoT}[T(n), d(n)]$  if and only if there is a deterministic family of transformers with constant depth, embedding size  $d(n)$  and budget  $T(n)$  that recognizes it.

Add an image representing the transformer architecture

**Definition 2.2.1** ( $\text{CoT}[T(n), d(n)]$ ). A language  $\mathcal{L}$  is in  $\text{CoT}[T(n), d(n)]$  if and only if there is an integer  $L$  and two functions  $T'(n) = O(T(n))$ ,  $d'(n) = O(d(n))$ , such that for every positive integer  $n$ , there is an  $L$ -layer deterministic transformer, denoted by  $\text{TF}_n$  with embedding size  $d'(n)$ , that can output  $\mathcal{L}(\alpha)$  given any input  $\alpha$  in  $\Sigma^n$ , using  $T'(n)$  steps of chain of thought.

In the same way, we define the classes of languages recognized by probabilistic transformers. Given two functions  $T(n)$  and  $d(n)$  we define the complexity class  $\text{RCoT}[T(n), d(n)]$ . Informally, a language  $\mathcal{L}$  is in  $\text{RCoT}[T(n), d(n)]$  if and only if there is a probabilistic family of transformers with constant depth, embedding size  $d(n)$  and budget  $T(n)$  that recognizes it.

**Definition 2.2.2** ( $\text{RCoT}[T(n), d(n)]$ ). A language  $\mathcal{L}$  is in  $\text{RCoT}[T(n), d(n)]$  if and only if there is an integer  $L$  and two functions  $T'(n) = O(T(n))$ ,  $d'(n) = O(d(n))$ , such that for every positive integer  $n$ , there is an  $L$ -layer probabilistic transformer, denoted by  $\text{TF}_n$  with embedding size  $d'(n)$ , that can output  $\mathcal{L}(\alpha)$  given any input  $\alpha$  in  $\Sigma^n$ , using  $T'(n)$  steps of chain of thought with probability higher than  $2/3$ .

### 2.3 Circuit Complexity

### 2.4 Known results

### 3. GENERAL PROPERTIES

In order to better understand  $\text{CoT}[T(n), d(n)]$  classes, it was needed to fill a gap in the literature related to the properties of these classes. For that, in this section we introduce a notion of normalization for transformers and prove that  $\text{CoT}[T(n), d(n)]$  is closed under boolean operations. As middle step, it was needed to build certain primitives operations which are also explained in this chapter.

#### 3.1 Transformers Primitives

##### Flagging Tokens and Positions

We will see later on that we will want to be able to identify when a certain token appeared in the tape. Also, we would like to know when a certain position is reached or if a certain vector corresponds to a specific position. This last task is not naturally done because of the extreme parallelism in the model.

For flagging a token we use the token embedding, and in the same way, for flagging a position we use the position encoding. Either way, the strategy is exactly the same.

For each flag that we are looking to add, we will add one to the dimension embedding. Notice that if we are working with a constant number of flags, the  $\text{CoT}[T(n), d(n)]$  class does not change. *Not sure about this last sentence. Maybe I should expand the idea if I want to say it*

The idea for flagging is straight forward. We will have an extra coordinate that will work as the flag. This coordinate will have two possible values to be defined. It will have the on value  $f_{\text{on}}$  if we are flagging its vector or the off value  $f_{\text{off}}$ , otherwise.

For example, let's be TF a transformer and  $\theta_{TE}$  its token embedding function. We will define a transformer  $\text{TF}'$  with the same behavior, but it will have flagged every vector that came from the token  $\sigma_f$ .

For the sake of clarity, we will add the flag in the last coordinate of the embedding, but this can be done in any position.

Let's define each component of  $\text{TF}'$  based on the components of TF. First, let's see the token embedding function. It will be the same except that in the new coordinate it will put a  $V_{\text{on}}$  if the token that is consuming is the one to be flagged and a 0 otherwise.

$$(\theta'_{TE}(\sigma))_j = \begin{cases} (\theta_{TE}(\sigma))_j & \text{if } j \neq d+1 \\ f_{\text{on}} & \text{if } j = d+1 \text{ and } \sigma = \sigma_f \\ f_{\text{off}} & \text{if } j = d+1 \text{ and } \sigma \neq \sigma_f \end{cases}$$

For the position embedding is enough to always put a 0 in the new coordinate. Then is only necessary to extend the matrix of the other layers to the new dimension. This can be done with a new row or column with 0s. The behavior of  $\text{TF}'$  will be the same as TF.

This mechanism clearly works for adding more dimensions to the embedding. We just add new flags with the same value for  $f_{\text{on}}$  and  $f_{\text{off}}$ .

## Preserve and ignore coordinates through ATTN and FF

When defining new transformers over existing ones we will want to extend their embedding dimension. We would like to preserve the behavior of the old transformer in the existing coordinates but ignore the new ones. Let's see how to extend the dimensions of the ATTN and FF layers in a way that we accomplish this.

If we want to ignore the  $i$ th coordinate through a layer of attention, it's enough to put zeros in the  $i$ th column and row of  $W_Q, W_K, W_V, W_O$ . This can be done with as many coordinates as desired. Note that because the  $i$ th row of  $W_O$  is zero, the output vector of the ATTN layer will have a zero in the  $i$ th coordinate. This is the point that makes it possible to not just ignore the value of the  $i$ th coordinate, but to not affect it. This happens because the result of the attention is added to the original vector.

The same applies in the feed forward layer, putting zeros in the  $i$ th column and row of the matrix and vectors of the FF makes it ignore the value of the  $i$ th coordinate. Also in the result of the FF will appear a zero in the  $i$ th coordinate.

## Add vector

I have this one in my notes, but I don't think I use it anywhere. Should I write it?

Make a vector 0

## Linear transformations

Although it looks that transformers are a model that should be able to perform linear transformations in a native way, it is not the case. Applying a linear transformation to a vector is not direct since in every step of ATTN and FF the result of the layer is added to the original vector instead of replacing it. A sharp reader maybe would think that this can be solved applying the transformation  $M - I$  when we want to transform the vector  $x$  to  $Mx$ , since  $x + (M - I)x = x + Mx - Ix = Mx$ . This does not hold in our model since the addition and multiplication in the finite representation of the numbers used is not commutative. For avoiding the representation error, we will extend the embedding dimension to get more memory per vector such that we can store in new coordinates the value of  $Mx$  and then replace it in the original ones. After applying this process we will have corrupted all vectors except the one we are interested in.

The strategy is the following. We will apply the transformation to only one vector which was flagged with 1 as  $f_{\text{on}}$  and  $-1$  as  $f_{\text{off}}$  in the embedding step. This flag will appear duplicated. Let's assume the flag is in the last two coordinates. For having a clearer notation, we will say that the vector to which we are applying the linear transformation is the  $f$ th. Also, we will call  $M$  to the matrix of the linear transformation. The linear transformation will be done with two layers of ATTN.

The first layer will compute the result of the transformation and store it in new coordinates added as memory. The second layer will move the result to the right place.

For doing this we will need to extend the embedding dimension with enough coordinates to store the result of the transformation. Therefore, if we are working with a transformer of embedding dimension  $d$  and want to do a linear transformation of a vector, we will need to extend the dimension to  $2d + 2$ . The first  $d$  coordinates will be for the original vector, the second  $d$  coordinates will be the ones using as memory, and the extra two are for the flags. We will assume that the  $d$  coordinates used as memory are set to 0 by the

layers previous to this mechanism. This can be achieved maintaining them as 0 since the beginning or transformed to 0 with the primitive described in the previous section.

$$h_f = \begin{pmatrix} x_1 \\ \vdots \\ x_d \\ 0 \\ \vdots \\ 0 \\ 1 \\ 1 \end{pmatrix} \xrightarrow{\text{first ATTN layer}} h_f = \begin{pmatrix} 0 \\ \vdots \\ 0 \\ (Mx)_1 \\ \vdots \\ (Mx)_d \\ 1 \\ 1 \end{pmatrix} \xrightarrow{\text{second ATTN layer}} h_f = \begin{pmatrix} (Mx)_1 \\ \vdots \\ (Mx)_d \\ 0 \\ \vdots \\ 0 \\ 1 \\ 1 \end{pmatrix}$$

#### a graphic representation of the technique

Let's construct each of the ATTN layers. In both layers we will need for  $h_f$  to only attend to itself. For accomplishing that we will make use of the fact that the flag has value 1 only in  $h_f$  and  $-1$  in the other vectors. We asked for the flag by duplicate since we need two ATTN layers and in each layer we are corrupting all vectors except  $h_f$ . Therefore, making use of the mechanism explained before to preserve and ignore coordinates through ATTN layers, we will write all the matrix as if we are working with embedding dimension  $2d + 1$  but in reality, in the first layer we are using the first copy of the flag (and ignoring/preserving the second) and in the second we are doing it in the other way. Thus, adding the missing row and column of zeros will get us the real values of the matrix.

Let's define the parameters of the ATTN layers for  $h_f$  only attending to itself. If we take as query and key matrix as

$$W_Q = W_K = \begin{pmatrix} 0 & \dots & 0 & 0 \\ \vdots & \ddots & \vdots & \vdots \\ 0 & \dots & 0 & 0 \\ 0 & \dots & 0 & \infty \end{pmatrix}$$

then the keys and values of the vectors are

$$W_Q h_i = q_i = W_K h_i = k_i = \begin{cases} (0, \dots, 0, \infty * 1) & \text{if } i = f \\ (0, \dots, 0, \infty * -1) & \text{if } i \neq f \end{cases} = \begin{cases} (0, \dots, 0, \infty) & \text{if } i = f \\ (0, \dots, 0, -\infty) & \text{if } i \neq f \end{cases}$$

so the attention values for  $h_f$  are 1 for itself and 0 for the rest.

$$\begin{aligned} s_f &= \text{softmax}(\langle q_f, k_1 \rangle, \dots, \langle q_f, k_{f-1} \rangle, \langle q_f, k_f \rangle) \parallel (0, \dots, 0) \\ &= \text{softmax}(\infty * -\infty, \dots, \infty * -\infty, \infty * \infty) \parallel (0, \dots, 0) \\ &= \text{softmax}(-\infty, \dots, -\infty, \infty) \parallel (0, \dots, 0) \\ &= 0, \dots, 0, 1, 0, \dots, 0 \end{aligned}$$

$$(s_f)_i = \begin{cases} 1 & \text{if } i = f \\ 0 & \text{if } i \neq f \end{cases}$$

Now that we got for  $h_f$  to attend only to itself, let's remember from the definition of the ATTN mechanism (??) that the output of the ATTN layer is  $h'_i = W_O \sum_{j=1}^a (s_i)_j v_j$ . Let's analyze what  $h'_f$  looks by now:

$$h'_f = W_O \sum_{j=1}^a (s_f)_j v_j = W_O v_f = W_O (W_V h_f)$$

Then, defining  $W_O$  and  $W_V$  as follows

$$W_V = id^{\mathbb{R}^{(2d+1) \times (2d+1)}} \quad W_O = \begin{pmatrix} -id^{\mathbb{R}^{d \times d}} & 0^{\mathbb{R}^{d \times d}} & 0^{\mathbb{R}^d} \\ M & 0^{\mathbb{R}^{d \times d}} & 0^{\mathbb{R}^d} \\ 0^{\mathbb{R}^{1 \times d}} & 0^{\mathbb{R}^{1 \times d}} & id^{\mathbb{R}} \end{pmatrix}$$

makes us get

$$h'_f = \begin{pmatrix} -x_1 \\ \vdots \\ -x_d \\ (Mx)_1 \\ \vdots \\ (Mx)_d \\ 1 \end{pmatrix}$$

which added to  $h_f$  gives us what we were looking for after the first layer.

For the second ATTN layer, we can use the same mechanism for making  $h_f$  only attend to itself. And defining  $W_V$  and  $W_O$  as follows

$$W_V = id^{\mathbb{R}^{(2d+1) \times (2d+1)}} \quad W_O = \begin{pmatrix} 0^{\mathbb{R}^{d \times d}} & id^{\mathbb{R}^{d \times d}} & 0^{\mathbb{R}^d} \\ 0^{\mathbb{R}^{d \times d}} & -id^{\mathbb{R}^{d \times d}} & 0^{\mathbb{R}^d} \\ 0^{\mathbb{R}^{1 \times d}} & 0^{\mathbb{R}^{1 \times d}} & id^{\mathbb{R}} \end{pmatrix}$$

we get

$$h'_f = \begin{pmatrix} (Mx)_1 \\ \vdots \\ -(Mx)_d \\ -(Mx)_1 \\ \vdots \\ -(Mx)_d \\ 1 \end{pmatrix}$$

which added to  $h_f$  gives us what we were looking for after the second layer.

If we wanted to remove the flag, we could replace the bottom leftmost coordinate of  $W_O$  by  $-1$  in both layers.

**Max of two numbers**

**3.2 Normalization**

**3.3 Boolean Operations**



## 4. MAIN RESULTS



## 5. CONCLUSIONS