



UNIVERSIDAD DE BUENOS AIRES
FACULTAD DE CIENCIAS EXACTAS Y NATURALES
DEPARTAMENTO DE COMPUTACIÓN

Understanding the Theoretical Limits of LLMs from the Perspective of Computational Complexity: The Role of Randomness in the Chain of Thought

Tesis de Licenciatura en Ciencias de la Computación

Rafael Romani

Director: Santiago Cifuentes

Codirector: Santiago Figueira

Buenos Aires, 2026

ENTENDIENDO LOS LÍMITES TEÓRICOS DE LOS LLMS DESDE EL PUNTO DE VISTA DE LA COMPLEJIDAD COMPUTACIONAL: EL ROL DE LA ALEATORIEDAD DENTRO DEL CHAIN OF THOUGHT

[Acá va el resumen]

Palabras claves: [acá van las palabras clave]

**ENTENDIENDO LOS LÍMITES TEÓRICOS DE LOS LLMS DESDE
EL PUNTO DE VISTA DE LA COMPLEJIDAD COMPUTACIONAL:
EL ROL DE LA ALEATORIEDAD DENTRO DEL CHAIN OF
THOUGHT**

[write your abstract here]

Keywords: [write your keywords here]

[acá van los agradecimientos]

CONTENTS

1. Introduction	1
2. Theoretical Framework	4
2.1 Transformers	4
2.2 Finite precision modeling	7
2.3 Complexity theory background	9
2.3.1 Boolean circuits	9
2.3.2 Boolean circuits and Turing machines	13
2.4 Transformers complexity classes	14
2.5 Known results on transformers power	14
2.5.1 Upper bound	15
2.5.2 Lower bound	16
3. General properties	17
3.1 Transformers Primitives	17
3.1.1 Flagging tokens and positions	17
3.1.2 Preserve and ignore coordinates through ATTN and FF	18
3.1.3 Add vector and constant functions	18
3.1.4 Make a vector ignore others in an ATTN layer	19
3.1.5 Linear transformations	20
3.1.6 Max coordinate	20
3.1.7 Repeat a token once it's printed	21
3.1.8 Force to halt after certain number of steps of CoT	21
3.1.9 Copy one vector into another with ATTN	22
3.2 Normalization	22
3.3 Boolean Operations	23
3.3.1 Complement of a language	23
3.3.2 Disjunction of two languages	23
4. Main results	25
4.1 Simulation of deterministic transformers	25
4.2 Normalization	25
4.3 Robustness of our Definition	26
4.4 Closed by Boolean Operations	27
4.5 Upper Bound	27
4.6 Lower Bound	29
4.7 Relations between models	29
4.7.1 Deterministic and probabilistic transformers	29
4.7.2 Uniform transformers and Turing machines	30
4.7.3 Uniform and non-uniform transformers	30
5. Conclusions	31

Appendix	32
A. Primitives	33
A.1 Construction for flagging	33
A.2 Ignore vectors in Attention layer	33
A.3 Linear transformations	34
A.4 Implementation of the special operation	35
A.5 Repeat tokens	36
B. Algorithm for normalizing	38
C. Boolean operations for deterministic transformers	40
C.1 Complement	40
C.2 Disjunction	40
D. Probability calculations	44
D.1 Number of runs for error reduction	44
D.2 Probabilities needed for disjunction	44
E. Proofs of probabilistic bounds	46
E.1 Sampling the distribution with a $\text{BPAC}^0[1]$ circuit	46
E.2 Simulation of randomized circuits by probabilistic transformers	46

1. INTRODUCTION

In 2017, a group of researchers working at Google published a paper introducing a technology that revolutionized society. In *Attention is All You Need*, they presented the transformers, today's core of large language models (LLMs) and generative AI.

This architecture was built as a text sequence-to-sequence model for machine translation but today finds applications in various fields, such as natural language processing, computer vision, and speech recognition.

The first model implementing transformers open for public use was ChatGPT by OpenAI. It made its debut in November 2022 and by December 2024 it reported 300 million global weekly users. Now there are many models competing for this market such as Claude, from Anthropic, or Gemini from Google.

These models can all be categorized as generative AI tools. This technology is arguably the most rapidly adopted in history. In the second half of 2025, roughly one in six people worldwide were using generative AI tools (Microsoft AI Economy Institute, 2026).

LLMs are being used not only as chatbots for asking small questions, but are shaping the industry and future. For example, Liang et al. (2024) found that more than 17% of the Computer Science's papers published between January 2020 and February 2024 on the *arXiv*, *bioRxiv*, and *Nature* portfolio journals had sentences heavily edited by ChatGPT. Naturally, the question of their computational power arises. Are these models genuinely capable of performing the tasks for which they are being employed? In this thesis, we aim to formally answer that question. We approach the problem from the perspective of formal languages.

First, we need to define what an LLM is. We consider LLMs just as instantiations of the transformer architecture. Thus, we need to understand what transformers are.

A transformer is a machine consisting of a tape from which it reads the input and onto which it appends the output character by character in an autoregressive manner. It reads the content of the tape, then computes the next character, appends it at the end, and repeats. Vaswani et al. (2017a) presented a concrete structure of how to do this.

First, the input text written by the user is split into tokens. This is a deterministic process and is done mainly as a resource optimization. Thus, in this thesis (and in all the papers of the area) it is assumed that the user talks in tokens, i.e., the transformers receive tokens in their input tape. The set of tokens can be interpreted as the language of the transformers.

Then, the first real component of a transformer is the mechanism that maps the tokens of the tape into vectors of numbers. Each occupied position of the tape is mapped into a vector. Each vector is computed depending on its position in the tape and the token at that position. Thus, this mapping from tape positions to vectors can be decomposed into two components: the position encoding and the token embedding. The dimension of the vectors is the dimension of the transformer.

Then, these vectors pass through a fixed sequence of layers consisting of an attention mechanism and a feed-forward network. These mechanisms encode information about how vectors relate to one another, i.e., how a token at a given position relates to a token at another position. The output of each layer is added to the original vectors and passed to

the next layer.

In the vector corresponding to the last token of the tape, after passing through all the layers, a probability distribution over the next token is encoded. Said probability is extracted and used to pick the new token to append in the tape. The token is selected, and the process is repeated. In this way, the LLMs like ChatGPT generate the text that they answer to our prompts.

The performance of LLMs improves when we allow them to generate more intermediate tokens. Intuitively, it is like giving them a scratch pad. This process is known as Chain of Thought, and it is known that increasing the number of intermediate tokens improves the quality of the answers (Wei et al., 2022). Notably, using heuristic phrases in prompts, such as 'let us think step by step', improves the answers without the need to retrain the model (Kojima et al., 2022). Also, prompting the first steps of the reasoning, even if they are wrong, gives better results (Wang et al., 2023; Madaan and Yazdanbakhsh, 2022). This is why modern models perform this reasoning in the background before responding to the user.

In this thesis, we analyze transformers as language recognizers. We define them as machines that receive an input string and determine whether it has a certain property, generating the token q_{accept} or q_{reject} respectively. When working with Turing machines as language recognizers, it is standard to study the languages recognizable when the machines are limited to a certain time or space budget as a function of the input size. When working with transformers, there are different resources of interest to bound. For example, Merrill and Sabharwal (2026) studied families of transformers that grow the number of layers with the input size. Merrill and Sabharwal (2026) prove that even with a logarithmic growth on the number of layers, there is a significant improvement in the languages that transformers can recognize.

In this work, we extend the results of Li et al. (2024). They studied the number of steps of Chain of Thought (CoT) as the limited resource. In their work, they give a formal argument to why transformers perform better with more steps of CoT. They prove that with $T(n)$ steps of CoT, a transformer can recognize languages recognizable with Boolean circuits of size $T(n)$. They also prove that a logarithmic number of steps of CoT does not grow the set of languages recognizable by transformers without CoT.

Li et al. (2024) give bounds on the computing power of transformers and define classes of languages depending on how many steps of CoT are required to recognize them. They do not study the properties of those classes. In this thesis, as a preliminary step toward our main objective, we prove that the classes of Li et al. (2024) are closed under Boolean operations.

Most of the formal work done to understand the computational power of transformers ignores one of their main characteristics. Li et al. (2024) and Merrill and Sabharwal (2026) work with deterministic abstractions, but transformers are intrinsically stochastic. As was mentioned before, in each step of CoT, transformers generate a probability distribution to decide next token. Formal works tend to not sample this distribution and just define the next token as the token with the largest probability in the distribution. In practice, the models behave in a nondeterministic way sampling the distribution for defining the next token. The stochasticity of the models is controlled through a parameter called temperature, and it is well known that a nonzero temperature is sometimes necessary for good performance (Demir and Dogan, 2025).

In this thesis, we adapt the abstraction of Li et al. (2024) to incorporate the stochastic behavior. We define the corresponding complexity classes and establish the analogues of the deterministic bounds in the probabilistic setting. We complete the work studying the properties of the classes of languages recognizable by probabilistic transformers with different number of steps of CoT. The classes defined in this work are closed under Boolean operations and behave well as probabilistic classes, i.e., regardless of how much error we allow in the probabilistic transformers, the recognizable languages remain the same.

The main contribution of this work is that we prove relations between the power of probabilistic and deterministic transformers. These relations provide a solid foundation for understanding when it suffices to study deterministic transformers and when a probabilistic abstraction is warranted.

2. THEORETICAL FRAMEWORK

2.1 Transformers

In this work we analyze transformers as language recognizers. A transformer is a machine that works over an alphabet Σ . It has a single tape from which it reads the input and onto which it appends additional tokens in an autoregressive manner. There are two distinguished tokens called **decision symbols**. Once the transformer writes one of those, it halts. We say that a transformer accepts a word if, when it halts, the decision symbol printed is the accepting one (q_{accept}). Analogously, a transformer rejects a word if the decision symbol printed is the rejecting one (q_{reject}).

In each step, the transformer maps the content of the tape to a probability distribution over Σ . Then, depending on this distribution, the transformer prints the next token. The transformer iterates this process until it halts. These iterations are called **chain of thought** (Wei et al., 2022).

We split the process of generating the next token into two separate steps. First, the computation of the probability distribution over the alphabet which we call distribution generation. Second, the selection of a token based on the probability distribution, which we call token selection. In this work we analyze and compare two different algorithms for token selection.

Most prior work only considers transformers that operate over an infinite tape. This is far from the real models, thus the transformers studied in this thesis work over a finite tape. The size of the tape is defined, for each input size n , as $\text{budget}(n) + n$. It is important to remark that the size of the tape bounds the number of steps of CoT that the transformer makes. For an input word of length n , no computation can take more than $\text{budget}(n)$ steps because every step writes a token in the tape.

The process of distribution generation is dependent on many parameters. In the real models, these parameters are the ones computed during the training.

We refer to them as θ and consider them as part of the definition of a transformer.

We denote by $\text{TF}(\alpha)$ the token outputted after one step of the computation of the transformer TF when the tape content is $\alpha \in \Sigma^*$. The autoregressive token generation is defined recursively. The token outputted after i steps is $\text{TF}^i(\alpha) := \text{TF}^{i-1}(\alpha \text{TF}(\alpha))$ with base case $\text{TF}^1(\alpha) := \text{TF}(\alpha)$. In other words, the output of the i th step is $\sigma_i = \text{TF}^i(\alpha) = \text{TF}(\alpha \sigma_1 \dots \sigma_{i-1})$ where $\sigma_j \in \Sigma$ for all j .

We denote by $\text{TF}^*(\alpha)$ the token written by the transformer when it halts, i.e., at the end of the computation, when the input is α . As said before, the transformer only halts when it prints a decision symbol (q_{accept} or q_{reject}). Therefore, $\text{TF}^*(\alpha) \in \{q_{accept}, q_{reject}\}$.

As stated before, the **distribution generation** is a mapping of the content of the tape to a probability distribution over Σ . There are different algorithms for this process in the literature. All transformers considered in this work employ the one presented in Li et al. (2024). This algorithm consists of four parts: a token embedding layer (TE), a position encoding layer (PE), an output linear layer (OUTPUT) and a stack of L identical decoder layers, where L is also called the depth of the model. Each decoder layer

has two sub-layers: a multi-head self-attention layer (ATTN) and a position-wise fully-connected feed-forward network (FF). Each part and layer has its own trainable parameters. Then, we can define the distribution generation algorithm by its parameters $\theta = \{\theta_{\text{TE}}, \theta_{\text{PE}}, \{\theta_{\text{ATTN}}^l, \theta_{\text{FF}}^l\}_{l=0}^{L-1}, \theta_{\text{OUTPUT}}\}$. The algorithm works over vectors of numbers of certain dimension d . The dimension is dependent on the size of the input word.¹ We use $d(n)$ to refer to the **embedding size** when the input word has size n .

We now describe each part of the transformer’s architecture on detail:

Token Embedding: It is a mapping from Σ to vectors in $\mathbb{R}^d$. This mapping is defined by $\theta_{\text{TE}} \in \mathbb{R}^{d \times |\Sigma|}$. For all $\sigma \in \Sigma$, we denote by $\theta_{\text{TE}}(\sigma) \in \mathbb{R}^d$ the mapping of token σ .

Position Encoding: It is a mapping from the positions of the tape to vectors in $\mathbb{R}^d$. This mapping is defined by $\theta_{\text{PE}} \in \mathbb{R}^{d \times (\text{budget}(n)+n)}$. For all positions of the tape $1 \leq i \leq \text{budget}(n) + n$ we denote by $\theta_{\text{PE}}(i) \in \mathbb{R}^d$ the mapping of position i .

Self Attention Layer: It is a mapping from $(\mathbb{R}^d)^a$ to $(\mathbb{R}^d)^a$. Given parameters $\theta_{\text{ATTN}} = (W_Q, W_K, W_V, W_O) \in \mathbb{R}^{d \times d} \times \mathbb{R}^{d \times d} \times \mathbb{R}^{d \times d} \times \mathbb{R}^{d \times d}$, the layer is defined by the algorithm 1. This layer is defined only as a single head attention layer because a multi-head attention layer with any constant number of heads can be simulated with multiple single-head attention layers (Yang et al., 2025).²

Algorithm 1 Causal Self-Attention, ATTN

Input: Parameter $\theta_{\text{ATTN}} = (W_Q, W_K, W_V, W_O)$ and embedding $h = (h_1, \dots, h_a) \in \mathbb{R}^d \times \dots \times \mathbb{R}^d$.

Output: Embedding $h' = (h'_1, \dots, h'_a) := \text{ATTN}_{\theta_{\text{ATTN}}}(h_1, \dots, h_a)$.

- 1: $q_i := W_Q h_i, k_i := W_K h_i, v_i := W_V h_i, \forall i \in [1, \dots, a]$
 - 2: $s_i := \text{softmax}(\langle q_i, k_1 \rangle, \dots, \langle q_i, k_i \rangle) \parallel (0, \dots, 0)$
 - 3: $h'_i := W_O \sum_{j=1}^a (s_i)_j v_j$
-

Feed Forward Network: It is a mapping from $\mathbb{R}^d$ to $\mathbb{R}^d$. Given $\theta_{\text{FF}} = (W_1, b_1, W_2, b_2) \in \mathbb{R}^{d \times d} \times \mathbb{R}^d \times \mathbb{R}^{d \times d} \times \mathbb{R}^d$, the layer is defined as $\text{FF}_{\theta_{\text{FF}}}(h) := W_2 \text{relu}(W_1 h + b_1) + b_2$, where $\text{relu} : \mathbb{R}^d \rightarrow \mathbb{R}^d$ replaces each position x_i by $\max(0, x_i)$.

Output Layer: It is a mapping from $\mathbb{R}^d$ to a probability distribution over Σ . Given $\theta_{\text{OUTPUT}} \in \mathbb{R}^{|\Sigma| \times d}$, we define it as $\text{OUTPUT}_{\theta_{\text{OUTPUT}}}(h) := \text{softmax}(\theta_{\text{OUTPUT}} h)$.

All these parts come together for the distribution generator process in the algorithm 2.

The OUTPUT layer and the ATTN layer use a softmax function $\mathbb{R}^n$ to $\mathbb{R}^n$ defined as $\text{softmax}(x)_i = \exp(x_i) / \sum_{j=1}^n \exp(x_j)$.

The second component of a transformer is the **token selector**. This algorithm consumes the probability distribution computed by the distribution generator and returns a token.

In this work we study two selectors, we call them **deterministic selector** and **probabilistic selector**. Let $p \in \Delta(\Sigma)$ be the distribution over Σ outputted by the distribution

¹ This is note the case in the real models since the same model is used for all input sizes.

² The matrices are named Q for query, K for key, V for value and O for output. It is called query because q_i represents a question of how much the i -th vector should attend to the others in a specific setting (determined by W_Q). The k_j vector answers said question such that $\langle q_i, k_j \rangle$ is large if the j -th vector is relevant to the i -th. Thus, the key matrix is used to compute the relevance. The value matrix has the role of extracting the information that the j -th vector should give to the i -th. Finally, the output matrix is just a post-processing.

Algorithm 2 Distribution generator

Input: Transformer parameter $\theta = (\theta_{\text{TE}}, \theta_{\text{PE}}, \{\theta_{\text{ATTN}}^l, \theta_{\text{FF}}^l\}_{l=0}^{L-1}, \theta_{\text{OUTPUT}})$ and tokens of the tape $\alpha \in \Sigma^*$.

Output: distribution $p_\theta(\cdot|\alpha)$.

- 1: $h_i^{(0)} \leftarrow \theta_{\text{TE}}(\alpha_i) + \theta_{\text{PE}}(i), \forall 1 \leq i \leq |\alpha|$
- 2: **for** $l = 0, \dots, L - 1$ **do**
- 3: $(h_1^{(l+0.5)}, \dots, h_{|\alpha|}^{(l+0.5)}) \leftarrow (h_1^{(l)}, \dots, h_{|\alpha|}^{(l)}) + \text{ATTN}_{\theta_{\text{ATTN}}^{(l)}}(h_1^{(l)}, \dots, h_{|\alpha|}^{(l)})$
- 4: $h_i^{(l+1)} \leftarrow h_i^{(l+0.5)} + \text{FF}_{\theta_{\text{FF}}^{(l)}}(h_i^{(l+0.5)}), \forall 1 \leq i \leq |\alpha|$
- 5: **end for**
- 6: $p_\theta(\cdot|\alpha) \leftarrow \text{OUTPUT}_{\theta_{\text{OUTPUT}}}(h_{|\alpha|}^{(L)})$

generator. Then, the deterministic selector always returns the token with the largest probability (argmax over p). In transformers with this selector, we do not allow for two or more tokens to be tied with the maximum probability, i.e., $\nexists \sigma$ s.t. $p(\text{argmax}(p)) = p(\sigma) \wedge \sigma \neq \text{argmax}(p)$. In contrast, the probabilistic selector operates stochastically. It returns a token $\sigma \in \Sigma$ with probability $p(\sigma)$. We allow for more than one maximum with this selector.

Figure 2.1 shows how the token selector connects with the distribution generator and its components.




Fig. 2.1: Transformer architecture

Although the presentation for transformers introduced in this section operates over real numbers, the transformers analyzed during this work utilize $\mathbb{F}_{e,s}$ for representing them.

The formal construction is the same, but all operations over the set of real numbers are replaced by their corresponding in the set of $\mathbb{F}_{e,s}$ numbers defined in Section 2.2.

Definition 2.1.1 (Deterministic transformer). We call a transformer deterministic when the token selector is the deterministic selector. We say that a family of deterministic transformers $\{\text{TF}_n\}_{n \in \mathbb{N}}$ recognizes a language $\mathcal{L}$ when:

$$\alpha \in \mathcal{L} \iff \text{TF}_{|\alpha|}^*(\alpha) = q_{\text{accept}}$$

As mentioned in Chapter 1, there is an unexplored area in the study of transformers expressiveness; a study of non-determinism is missing.

When TF is a deterministic transformer, $\text{TF}^i(\alpha)$ is an exact symbol for all $\alpha \in \Sigma^*$ and $i \in \mathbb{N}$. On the contrary, when the token selection process of the transformer samples the distribution, it becomes a random variable. Therefore, we introduce the following definition.

Definition 2.1.2 (Probabilistic transformer). We call probabilistic transformer to the transformers that use the probabilistic selector. We say that a family of probabilistic transformers $\{\text{TF}_n\}_{n \in \mathbb{N}}$ recognizes a language $\mathcal{L}$ when:

$$\alpha \in \mathcal{L} \iff \Pr[\text{TF}_{|\alpha|}^*(\alpha) = q_{\text{accept}}] > 2/3$$

2.2 Finite precision modeling

In real models, the transformers typically work with 16- or 32-bit floating-point numbers, and there is active work in reducing their size. For this reason, in this thesis we follow the decision of Li et al. (2024) to work with constant-precision transformers, where the output of each arithmetic operation is rounded to the closest floating-point number representable by a fixed number of digits following IEEE 754 standard. This analysis avoids the unrealistic infinite precision assumption made in prior works (Pérez et al., 2019). The definitions stated in this section are taken from Li et al. (2024).

Intuitively, we represent the numbers as a tuple $\langle \text{sign}, \text{significand}, \text{exponent} \rangle$ with a fixed number of bits for each coordinate. The tuple is interpreted as $\text{sign} \cdot \text{significand} \cdot 2^{\text{exponent}}$. For example, the number 8 is represented as $\langle +, 1, 3 \rangle$. When fixing the number of bits for the significand and the exponent there are numbers non-representable. For example, with only 2 bits of significand and 3 of exponent, the number 9 cannot be represented exactly. We can represent $8 = 1 \cdot 2^3$ and $12 = 3 \cdot 2^2$, thus when an arithmetic operation should output 9 we output 8.

Below we give a formal definition of the floating-point number and rounding operation. We use $\phi(a) = \sum_{i=1}^k 2^{k-i} a_i$ to denote the integer value of the binary string $a \in \{0, 1\}^k$ for any $k \in \mathbb{N}^+$.

Definition 2.2.1 (Floating-point Representation). Let e be the number of bits for exponents and s be the number of bits for significand. A $(e + 2s + 1)$ -bit binary string $a = (a_1, a_2, \dots, a_{e+2s+1}) \in \{0, 1\}^{e+2s+1}$ is a floating-point binary representation of number $\phi_{e,s}(a) \triangleq \text{sign}(a) \cdot 2^{\text{exponent}(a)} \cdot \text{significand}(a)$ with e -bit exponent and $2s$ -precision, where the sign is $\text{sign}(a) \triangleq 2a_1 - 1$, the significand is $\text{significand}(a) \triangleq 2^{-s} \phi(a_{2:2s+1})$, and the exponent is $\text{exponent}(a) \triangleq \phi(a_{2s+2:2s+e+1}) - 2^{\max(0, e-1)}$. We define the set of numbers expressible with e bits of exponent and s bits of significand as $\mathbb{F}_{e,s}$. We define $\infty_{e,s}$ as

RR: I'm missing good symbols for ∞ y $-\infty$.

the maximum number in $\mathbb{F}_{e,s}$. Similarly, we define $-\infty_{e,s}$ as the minimum. Since in this work we only analyze constant precision transformers, we will omit the subscripts of ∞ and $-\infty$.

Definition 2.2.2 (Correct Rounding). For any $x \in \mathbb{R}$, we define correct rounding $\text{round}(x, \mathbb{F}_{e,s})$ as the closest number to x in $\mathbb{F}_{e,s}$. We break the tie by picking the one with a smaller absolute value. In particular, we denote the rounding operation with e -bit exponent, $2s$ -bit precision by $\text{round}_{e,s}(\cdot) \triangleq \text{round}(\cdot, \mathbb{F}_{e,s})$, which is also denoted by $[\cdot]_{e,s}$ for convenience. We extend the definition of round and $\text{round}_{e,s}$ to vector inputs by applying rounding coordinate-wise.

Our notion of floating-point number simplifies the IEEE 754 Standard for Floating-point Arithmetic (IEEE, 2008) by removing ∞ and $-\infty$. When overflow happens, we always round the output to ∞ , and when underflow happens to $-\infty$. For unary functions like $\exp(\cdot)$ and binary functions including addition, subtraction, multiplication, and division, we simply define their rounded version by rounding their outputs. Whenever division by 0 happens, we define the output as 0.³

Next, we define finite-precision summation over more than two numbers by decomposing it as a chain of rounded binary addition in a fixed order.

Definition 2.2.3 (Summation with Iterative Rounding). For any $s, n \in \mathbb{N}^+$ and vector $x \in \mathbb{R}^n$, we define summation with iterative rounding to e bit exponent and $2s$ -bit precision as $\text{sum}_{e,s} : \cup_{n \in \mathbb{N}^+} (\mathbb{F}_{e,s})^n \rightarrow \mathbb{F}_{e,s}$, where for any $n \in \mathbb{N}^+$ and $x \in \mathbb{R}^n$,

$$\text{sum}_{e,s}(x) \triangleq \left[\left[\left[[x_1 + x_2]_{e,s} + x_3 \right]_{e,s} + \cdots + x_{n-1} \right]_{e,s} + x_n \right]_{e,s}.$$

We further define the following operations:

- Finite-precision inner product: $\langle x, y \rangle_{e,s} \triangleq \text{sum}_{e,s}(x \odot y)$;
- Finite-precision matrix product: $(A \times_{e,s} B)_{i,j} \triangleq \langle (A_{i,:})^\top, B_{:,j} \rangle_{e,s}$;
- Finite-precision softmax: $\text{softmax}_{e,s}(x) \triangleq \left[[\exp(x)]_{e,s} / \text{sum}_{e,s}([\exp(x)]_{e,s}) \right]_{e,s}$.

From the saturation and the iterative rounding presented by Li et al. (2024), some reasonable properties appear that are important to remark:

- $\exp(-\infty) = 0$
- $\exp(\infty) = \infty$
- $-\infty \cdot \infty = -\infty$

³ Notice that the only division is in the softmax. If it is the one in the attention mechanism, it is reasonable to give an attention score of 0 to every vector. If it is the one in the output layer it means that all symbols have an extremely small probability, therefore it is reasonable to assign the same probability to all of them.

SC Decías que esto se parcheaba en la práctica, no? Digo porque capaz es buen lugar para una refe.

RR En la práctica tienen trucos para que no ocurra la division por cero. Mi ref es hablar con Luciana Ferrer y un post de stack exchange xd

- $\infty \cdot \infty = \infty$
- $\forall x \in \mathbb{F}_{e,s}$ it holds that $x + 0 = x$
- $\forall x \in \mathbb{F}_{e,s}$ it holds that $x - x = 0$
- $\forall x \in \mathbb{F}_{e,s}$ it holds that $x + \infty + \infty + -\infty = 0$

2.3 Complexity theory background

In this section, we define the standard notions of complexity theory used throughout this thesis.

Definition 2.3.1 (Language). In this work we consider the following notion of problems: given a sequence of input tokens, output a token as the answer. Mathematically, given a vocabulary Σ , we call a mapping $\mathcal{L} : \cup_{k \in \mathbb{N}^+} \Sigma^k \rightarrow \Sigma$ a problem. If the correct answer is always q_{accept} or q_{reject} , we call $\mathcal{L}$ a decision problem. In circuit complexity, such $\mathcal{L}$ is also called a **language**.

2.3.1 Boolean circuits

The Boolean circuits are a model of computation generalizing Boolean formulas. It is a simplified model of the silicon chips used to make modern computers. It is a natural model for nonuniform computation, as is the case for transformers; thus, most of our results on the expressiveness of transformers are stated in the language of circuit theory.

A Boolean circuit is a procedural diagram showing how to derive an output from a binary input string by applying a sequence of basic Boolean operations OR ($\vee$), AND ($\wedge$), and NOT ($\neg$) on the input bits.

Though the standard definition of circuit complexity only deals with binary strings, given any finite vocabulary Σ , we can always replace each token in Σ by its binary representation, and the length of the input only blows up by a constant factor. Therefore, we can extend existing complexity classes to arbitrary finite vocabulary naturally.

Definition 2.3.2 (Boolean circuits). For every $n \in \mathbb{N}$, an n -input, single-output **Boolean circuit** is a directed acyclic graph with n sources (vertices with no incoming edges) and one sink (vertex with no outgoing edges). All nonsource vertices are called gates and are labeled with one of $\wedge$, $\vee$ or $\neg$ (i.e., the logical operations OR, AND, and NOT). The vertices labeled with $\vee$ and $\wedge$ have fan-in (i.e., number of incoming edges) equal to 2 and the vertices labeled with $\neg$ have fan-in 1. The size of C , denoted by $|C|$, is the number of vertices in it.

If C is a Boolean circuit, and $x \in \{0, 1\}^n$ is some input, then the output of C on x , denoted by $C(x)$, is defined in the natural way. More formally, for every vertex v of C , we give it a value $val(v)$ as follows: If v is the i -th input vertex then $val(v) = x_i$ and otherwise $val(v)$ is defined recursively by applying v 's logical operation on the values of the vertices connected to v . The output $C(x)$ is the value of the output vertex.

Definition 2.3.3 (Circuit families and language recognition). Let $T : \mathbb{N} \rightarrow \mathbb{N}$ be a function. A $T(n)$ -size **circuit family** is a sequence $\{C_n\}_{n \in \mathbb{N}}$ of Boolean circuits, where C_n has n inputs and a single output, and its size $|C_n| \leq T(n)$ for every n . We say that a




Fig. 2.2: A circuit C computing the XOR function (i.e., $C(x_1, x_2) = 1$ iff $x_1 \neq x_2$).

language $\mathcal{L}$ is in $\text{SIZE}(T(n))$ if there exists a $T(n)$ -size circuit family $\{C_n\}_{n \in \mathbb{N}}$ such that for every $x \in \{0, 1\}^n$, $x \in \mathcal{L} \iff C_n(x) = 1$.

Interesting complexity classes arise when we consider “small” circuits, such as the class of languages recognized by all polynomial-size families.

Definition 2.3.4 (P/Poly). P/Poly is the class of languages that are decidable by polynomial-sized circuit families. That is, $\text{P/Poly} = \cup_{c \geq 0} \text{SIZE}(n^c)$.

The class P/Poly fits rather awkwardly in the complexity world since it contains undecidable languages such as Unary halt⁴. The root of the problem is that for a language $\mathcal{L}$ to be in P/Poly, it suffices that a circuit family for $\mathcal{L}$ exists even if we have no way of actually constructing the circuits. Thus, it may be fruitful to try to restrict attention to circuits that can actually be built, say using a fairly efficient Turing machine.

Definition 2.3.5 (P-uniform circuit families). A circuit family $\{C_n\}_{n \in \mathbb{N}}$ is P-uniform if there is a polynomial-time Turing machine that on input 1^n outputs the description of the circuit C_n .

Boolean circuits are a good model to represent parallel computation. The **depth** of a circuit is the length of the longest directed path from an input node to the output node.

Definition 2.3.6 (NC). For every d , a language $\mathcal{L}$ is in NC^d if $\mathcal{L}$ can be decided by a family of circuits $\{C_n\}_{n \in \mathbb{N}}$ where C_n has $\text{poly}(n)$ size and depth $O(\log^d(n))$. The class NC is $\cup_{i \geq 1} \text{NC}^i$.

When studying the depth of a circuit, the fan-in of the gates becomes of interest. A polynomial fan-in can be simulated with a logarithmic blow-up in depth. This motivates the study of circuits with unbounded fan-in.

Definition 2.3.7 (AC). The class AC^i is defined similarly to NC^i except gates are allowed to have unbounded fan-in (i.e., the OR and AND gates can be applied to more than two bits). The class AC is $\cup_{i \geq 0} \text{AC}^i$.

⁴ UHALT = $\{1^n : n\text{'s binary expansion encodes a pair } \langle M, x \rangle \text{ such that } M \text{ halts on input } x\}$

AC^0 circuits allow for an interesting number of operations. For example, addition and multiplication of numbers of fixed precision can be done with AC^0 circuits.

SC No termino de entender a qué problema te referís acá con el fixed precision.

RR fixed precision = punto fijo

Also, the equality of two strings of bits can be done in AC^0 . From the ability to compute equality it follows that any function f with $O(\log(n))$ bits as input can be computed in AC^0 . The circuit computing f enumerates the result for each of the inputs, and outputs the one that corresponds to the branch equal to the input. If k denotes the number of input bits of f , then f can be represented by the formula $\bigwedge_{i=0}^{2^k} (f(i) \wedge i = x)$. Observe that the formula represents a circuit of constant depth. Also note that the formula has size $O(2^k)$; then, if $k \in O(\log(n))$, the circuit has polynomial size; and, therefore, it is in AC^0 .

One of the main results of circuit complexity is that the MAJORITY problem⁵ cannot be computed by AC^0 circuits (Furst et al., 1984).

Theorem 2.3.1. $MAJORITY \notin AC^0$

This theorem motivates the study of threshold circuits, circuits allowed to have MAJORITY gates.

Definition 2.3.8 (TC). The class TC^i is defined similarly to AC^i except the circuits can have gates OR, AND, NOT and MAJORITY of unbounded fan-in. The class TC is $\bigcup_{i \geq 0} TC^i$.

Following Theorem 2.3.1, a family of variations of MAJORITY began to be studied.

Definition 2.3.9 (MAJORITY with promise). For every $e \in (0, 1/2]$ we define the problem **Majority with promise** (MAJ_e) as MAJORITY but restricted to inputs that have more than $(1/2 + e)n$ 1s or more than $(1/2 + e)n$ 0s where n is the length of the input.

Observe that $MAJ_0 = MAJORITY$.

Theorem 2.3.2 (Viola (2009)). For all $e > 0$, there is a circuit in AC^0 recognizing MAJ_e .

Randomized Boolean circuits

In this thesis we also work with randomized Boolean circuits. These are Boolean circuits, but they are allowed to have random gates. These gates have fan-in 0, and their value is a uniformly distributed bit. This value is fixed once the evaluation of the circuit starts. Then, the evaluation of a randomized circuit C with input x can be thought as evaluating a Boolean circuit C' with input xr , where r are the values of the random gates.

Definition 2.3.10 (Randomized Boolean circuits). A **randomized Boolean circuit** is a directed acyclic graph with $n+r$ sources and one sink. All nonsource vertices are labeled as $\wedge$, $\vee$ or $\neg$ (i.e., the logical operations OR, AND, and NOT). The vertices labeled with $\vee$ and $\wedge$ have fan-in (i.e., number of incoming edges) equal to 2 and the vertices labeled with $\neg$ have fan-in 1. The size of C , denoted by $|C|$, is the number of vertices in it. The source vertices are either labeled as input or random.

If C is a Boolean circuit, and $x \in \{0,1\}^n$ is some input, then the output of C on x , denoted by $C(x)$, is defined as random variable in the following way. For every vertex v

⁵ MAJORITY = $\left\{ x : \frac{\sum_{i=1}^{|x|} x_i}{|x|} > \frac{1}{2} \right\}$, i.e., the words with more 1s than 0s

of C , we give it a value $val(v)$ as follows: If v is the i -th input vertex, then $val(v) = x_i$; if v is labeled as **random**, then $val(v)$ is a uniform random Boolean variable; and otherwise $val(v)$ is defined recursively by applying v 's logical operation on the values of the vertices connected to v . The output $C(x)$ is the value of the output vertex.

In general, we say that a family of random circuits recognizes a language $\mathcal{L}$ when $\Pr[C_{|x|}(x) = \mathcal{L}(x)]$ is high.

Definition 2.3.11 (Randomized circuit families and language recognition). Let $T : \mathbb{N} \rightarrow \mathbb{N}$ be a function. A $T(n)$ -size **randomized circuit family** is a sequence $\{C_n\}_{n \in \mathbb{N}}$ of randomized Boolean circuits, where C_n has n inputs and a single output, and its size $|C_n| \leq T(n)$ for every n . We say that a language $\mathcal{L}$ is in $\text{BPSIZE}(T(n))[f(n)]$ if there exists a $T(n)$ -size randomized circuit family $\{C_n\}_{n \in \mathbb{N}}$ with $O(f(n))$ random gates such that for every $x \in \{0, 1\}^n$, $\Pr[C_n(x) = 1] \geq 2/3 \iff \mathcal{L}(x)$.⁶

Since every random Boolean circuit can be interpreted as Boolean circuit with the random bits padded in the input, we can also define recognition in the following way:

Definition 2.3.12. We say that a language $\mathcal{L}$ is in $\text{BPSIZE}(T(n))[f(n)]$ if there exists a $T(n)$ -size deterministic circuit family $\{C'_n\}_{n \in \mathbb{N}}$ such that for every $x \in \{0, 1\}^n$, $\Pr_{r \in \{0, 1\}^{f(n)}}[C'_n(xr) = 1] \geq 2/3 \iff \mathcal{L}(x)$.




Fig. 2.3: Circuit C padded with extra random inputs $r_1, \dots, r_{f(n)}$ to form C'

In general, we drop the bound for the number of random bits when it is not of interest.

Definition 2.3.13. $\text{BPSIZE}(T(n)) = \text{BPSIZE}(T(n))[T(n)]$.

Definition 2.3.14. Similar to P/Poly, we define BPP/Poly as the class of languages that are decidable by polynomial-sized random circuit families. That is, $\text{BPP/Poly} = \bigcup_{c \geq 0} \text{BPSIZE}(n^c)$.

We extend the definition of NC, NC^i , AC, AC^i , TC, TC^i uniform and not uniform to randomized Boolean circuits in the natural way. We denote by $\mathcal{C}[f(n)]$ the class $\mathcal{C}$ restricted to $O(f(n))$ random gates.

Observe that a Boolean circuit is a random Boolean circuit without random gates.

Theorem 2.3.3. Every circuit class $\mathcal{C}$ is contained in its corresponding randomized circuit class. For example $\text{AC}^0 \subseteq \text{BPAC}^0$.

⁶ As in many probabilistic complexity classes, the constant $2/3$ is arbitrary and can be changed to any constant in the range $(1/2, 1)$.

SC Me parece bien este nombre, pero diría que es no estándar si es que no es estándar. Aparte, porque hay algo llamado $BP \oplus NP$ (o en general, podés poner $BP \oplus$ adelante de muchas clases) y puede mezclarse.

RR en complexity zoo hay muchas clases que se llaman $BP[C]$ que son la versión bounded probability de otra clase C . A qué te referís vos con $BP \oplus C$?

The other inclusion is not known in general. In the case of AC^0 and $BPAC^0$, restricting $BPAC^0$ to only $O(\log(n))$ random bits is enough to make it fall inside AC^0 .

Theorem 2.3.4. $BPAC^0[\log(n)] \subseteq AC^0$

Proof. In AC^0 we can enumerate all possible values for the random bits. Then, we compute the output of the random Boolean circuit for each case. There are $O(n)$ possible random values. We can run the circuit for each of them in parallel leading only to a linear blow-up in width but constant in depth. Then, by definition of the random circuit, at least $2/3$ of the cases have the same output. Therefore, we can output $MAJ_{2/3}$. $\square$

2.3.2 Boolean circuits and Turing machines

Most of the work in complexity theory is done over Turing machines. In this section, we present some known relations between circuits and Turing machines.

Definition 2.3.15 (DTIME). Let $T : \mathbb{N} \rightarrow \mathbb{N}$ be some function. A language $\mathcal{L}$ is in $DTIME(T(n))$ iff there is a Turing machine that runs in time $c \cdot T(n)$ for some constant $c > 0$ and decides $\mathcal{L}$.

Definition 2.3.16. $P = \cup_{c \geq 0} DTIME(n^c)$.

A probabilistic Turing machine is a Turing machine with access to a random coin.

Definition 2.3.17 (BPTIME). Let $T : \mathbb{N} \rightarrow \mathbb{N}$ be some function. A language $\mathcal{L}$ is in $BPTIME(T(n))$ iff there is a probabilistic Turing machine M that runs in time $c \cdot T(n)$ for some constant $c > 0$ and for all $x \in \{0, 1\}^n$, $x \in \mathcal{L} \iff \Pr[M(x) = 1] \geq 2/3$

Definition 2.3.18. $BPP = \cup_{c \geq 0} BPTIME(n^c)$.

When restricting $P/Poly$ and $BPP/Poly$ to only their uniform families, they coincide with P and BPP . To simulate a circuit by a Turing machine, one first computes the circuit description (this can be done since we are only analyzing uniform families) and then evaluates it, both in polynomial time. For the simulation of Turing machines by circuits, the construction of Cook-Levin works.

Theorem 2.3.5. $Uniform \cap P/Poly = P$

Theorem 2.3.6. $Uniform \cap BPP/Poly = BPP$

An important result relating machines with and without randomness is Adleman's theorem. Adleman proved that the non-uniformity of $P/Poly$ gives enough power to simulate the randomness of BPP (Charles and Gill, 1981).

Theorem 2.3.7 (Adleman's theorem). $BPP \subseteq P/Poly$

The machines of $P/Poly$ can be interpreted as polynomial deterministic Turing machines that receive a polynomial-size advice. Then $(P/Poly)/poly$ are the machines of $P/Poly$ with another polynomial advice. Clearly $(P/Poly)/poly = P/Poly$; thus, it follows from Adleman's theorem that

Corollary 2.3.1. $BPP/Poly = P/Poly$

RR no sé si me convence mi argumento tirado así no más. Creo que lo tendría que explicar un poco más. Vale la pena que lo haga? Siento que no va a ser lo suficientemente largo como para ameritar un apéndice pero si lo suficientemente largo como para ser molesto.

2.4 Transformers complexity classes

As discussed in the first chapter, there are different resources to bound in the transformers to limit their expressive power. In this thesis, we analyze the effect of bounding the number of steps in the chain of thought that the transformers are allowed to make before answering. With this in mind, we define the following complexity classes.

RR:
analizar
título

Given two functions $T(n)$ and $d(n)$ we define the complexity class $CoT[T(n), d(n)]$. Informally, a language $\mathcal{L}$ is in $CoT[T(n), d(n)]$ if and only if there is a deterministic family of transformers with constant depth, embedding size $d(n)$, and budget $T(n)$ that recognizes it.

Definition 2.4.1 ($CoT[T(n), d(n)]$). A language $\mathcal{L}$ is in $CoT[T(n), d(n)]$ if and only if there is an integer L and two functions $T'(n) = O(T(n))$, $d'(n) = O(d(n))$, such that for every positive integer n , there is an L -layer deterministic transformer, denoted by TF_n with embedding size $d'(n)$, that outputs $\mathcal{L}(\alpha)$ given any input α in Σ^n , using $T'(n)$ steps of chain of thought.

In the same way, we define the classes of languages recognized by probabilistic transformers. Given two functions $T(n)$ and $d(n)$ we define the complexity class $RCoT[T(n), d(n)]$. Informally, a language $\mathcal{L}$ is in $RCoT[T(n), d(n)]$ if and only if there is a probabilistic family of transformers with constant depth, embedding size $d(n)$, and budget $T(n)$ that recognizes it.

Definition 2.4.2 ($RCoT[T(n), d(n)]$). A language $\mathcal{L}$ is in $RCoT[T(n), d(n)]$ if and only if there is an integer L and two functions $T'(n) = O(T(n))$, $d'(n) = O(d(n))$, such that for every positive integer n , there is an L -layer probabilistic transformer, denoted by TF_n with embedding size $d'(n)$, that correctly outputs $\mathcal{L}(\alpha)$ for any input $\alpha \in \Sigma^n$ with probability greater than $2/3$, using $T'(n)$ steps of chain of thought.

2.5 Known results on transformers power

The main objective of this thesis is to expand on the state-of-the-art bounds for transformer expressiveness. We adapt Li et al. (2024)'s results from deterministic transformers to probabilistic transformers. Their main results imply that $CoT[poly(n), log(n)] = P/Poly$.

2.5.1 Upper bound

Li et al. (2024) expands on the work of Merrill and Sabharwal (2023). Merrill and Sabharwal (2023) have shown that one step of a constant-depth, polynomial-embedding-dimension and log-precision deterministic transformer can be simulated in a small parallel time, i.e., using TC^0 circuits. This result is built on the fact that multiplication and division of n -bit binary numbers are in TC^0 (Hesse, 2001), as well as the iterated addition over n different n -bit binary integers.

However, such TC^0 expressiveness upper bounds may be unrealistic for transformers operating with floating-point numbers. Merrill and Sabharwal (2023) implicitly assumed that, when adding more than one floating-point number, the algorithm first computes the exact answer without rounding, using arbitrarily higher precision, and performs rounding only at the end. However, in practice, rounding occurs after each addition of two numbers, and it remains open whether such TC^0 upper bounds still hold. Immediate rounding makes iterated addition over floating-point numbers no longer associative (Goldberg, 1991). The associativity of integer addition is essential for the result that iterated addition over n different n -bit binary integers is in TC^0 .

The abstraction for transformers presented by Li et al. (2024) uses the numeric representation presented in Section 2.2. Their model rounds after every numerical operation and represents numbers with a constant number of bits for the significand and the exponent.

Li et al. (2024) prove that $\text{sum}_{e,s} : \mathbb{F}_{e,s}^n \rightarrow \mathbb{F}_{e,s}$ has an AC^0 circuit when e, s are both constants. Then, composing circuits, Theorem 2.5.1 is obtained.

Theorem 2.5.1. *One step of a deterministic, constant-depth, polynomial-embedding-dimension transformer can be simulated with a circuit in AC^0 .*

Observe that the fact that a single forward pass of a transformer can be simulated by an AC^0 circuit immediately implies that a transformer with $O(\log(n))$ steps of CoT can also be simulated by AC^0 . Thus, we obtain the following theorem.

Theorem 2.5.2. $\text{CoT}[\log(n), \text{poly}(n)] \subseteq \text{AC}^0$

Proof. Theorem 2.5.2 is sustained on the fact that a transformer with T steps of CoT can be interpreted as an OR of $|\Sigma|^T$ disjoint subcircuits, where each of them enumerates all possible sequences of T chain-of-thought tokens and outputs the value of the token in the branch where all the intermediate token values are consistent.

We say a sequence of tokens $\sigma_1, \dots, \sigma_T$ is consistent with the input α when it satisfies $\text{TF}(\alpha) = \sigma_1$, $\text{TF}(\alpha\sigma_1) = \sigma_2$, $\dots$, $\text{TF}(\alpha\sigma_1 \dots \sigma_{T-1}) = \sigma_T$. Each of the conditions can be checked by an AC^0 circuit in parallel. Then, for each sequence of tokens there is an AC^0 circuit validating its consistency. Since the transformer is deterministic, there is only one consistent sequence. The circuit outputs the value of the last token in the consistent sequence.

The enumeration of all possible sequences can be done in parallel and thus only takes constant depth. When $T = O(\log(n))$, this only leads to a $\text{poly}(n)$ blow-up in circuit size. Therefore, it follows that $\text{CoT}[\log(n), \text{poly}(n)] \subseteq \text{AC}^0$. $\square$

Since each pass of the transformer can be simulated by a circuit in AC^0 , a transformer performing a polynomial number of CoT steps $T(n)$ can be simulated by a circuit of depth $T(n)$ (and polynomial size), by stacking the circuit for one pass. Therefore, a transformer performing a polynomial number of steps can be simulated by a circuit in P/Poly.

Corollary 2.5.1. $\text{CoT}[\text{poly}(n), \text{poly}(n)] \subseteq \text{P/Poly}$

2.5.2 Lower bound

Theorem 2.5.3. *Any family of circuits of polynomial size $T(n)$ can be simulated by a family of constant-depth, constant-precision deterministic transformers using $T(n)$ steps of CoT with $O(\log(n))$ embedding size*

Li et al. (2024) construct, for each circuit, a transformer such that at each CoT step the output of one gate of the circuit is written to the tape. Theorem 2.5.3 can also be phrased with complexity classes as:

Theorem 2.5.4 (Lower bound for deterministic transformers). *For any polynomial function $T : \mathbb{N}^+ \rightarrow \mathbb{N}^+$, $\text{SIZE}(T(n)) \subseteq \text{CoT}[T(n), \log(n)]$. In particular, $\text{P/Poly} \subseteq \text{CoT}[\text{poly}(n), \log(n)]$.*

Proof. Let $\{C_n\}_{n \in \mathbb{N}}$ be a family of circuits of polynomial size. Li et al. (2024) construct a family $\{\text{TF}_n\}_{n \in \mathbb{N}}$ such that each TF_n simulates C_n .

TF_n evaluates C_n gate by gate, processing one gate per CoT step. Let us assign an index to each gate of C_n in topological order, such that the input gates receive indices 1 through n . Then, in the i -th step of CoT, TF_n prints the value of the $(n+i)$ -th gate, i.e., $\text{TF}_n^i(\alpha)$ equals the value of the $(n+i)$ -th gate when the input is α . Thus, at the end of the computation, the i -th position of the tape contains the value of the i -th gate.

Their construction encodes the information of each gate of the circuit into the position encoding. This information is encoded in binary, i.e., each coordinate takes the value 0 or 1. They store in the i -th vector the current gate index (i), the type (AND, OR, or NOT) of the $(i+1)$ -th gate, and the indices of its two input gates. The requirement of $d(n) = \Theta(\log(n))$ embedding size comes from the fact that there are $\text{poly}(n)$ gates and all of their indices must be uniquely representable in binary.

Then, at the i -th step of the computation, the attention mechanism is used to retrieve the values of the two input gates of the $(n+i)$ -th gate. Next, a feed-forward network computes the value of the $(n+i)$ -th gate based on its type and the values of its input gates retrieved in the previous step. Finally, in the last CoT step, TF_n outputs the value of the output gate (sink) of C_n . $\square$

3. GENERAL PROPERTIES

In this chapter, we study the properties of $\text{CoT}[T(n), d(n)]$ classes and the primitive operations that transformers can perform. We define the notion of normalized transformer and prove that $\text{CoT}[T(n), d(n)]$ classes are closed under Boolean operations.

3.1 Transformers Primitives

In this section we introduce a set of primitive operations that are used through the thesis. The operations are applied as modifications to an already existing transformer. If the operations are applied to an entire family of transformers $\{\text{TF}_n\}_{n \in \mathbb{N}}$ such that the language it computes is in $\text{CoT}[T(n), d(n)]$, then the language of the modified family will be in the same class. This is because, although the operations add layers and coordinates to the size of the embedding, this is only a constant factor.

Most operations do not make sense on their own, but they are used in later proofs.

It is interesting to remark that these are not only a theoretical gadget. It has been observed that transformers during training develop these types of structures and can be identified by patterns in the attention matrices (Olsson et al., 2022).

3.1.1 Flagging tokens and positions

For complex operations inside the loop of the transformer we will need to identify when a vector comes from a certain token or position. In this section we present how to mark the vectors with flags during the embedding.

For flagging a token we use the token embedding, and in the same way, for flagging a position we use the position encoding. Either way, the strategy is exactly the same. The embedding dimension is extended to allocate the flags. A new coordinate is added per flag. Therefore, using a constant number of flags only adds a constant to the embedding dimension. Thus, if the flagging is made in a family of transformers its $\text{CoT}[T(n), d(n)]$ class does not change.

For example, let v_σ and v_τ be vectors of a transformer TF where v_σ comes from embedding the token σ and v_τ from the token τ (with $\sigma \neq \tau$). We can define a transformer TF' that behaves like TF but recognize when a vector comes from embedding the token σ . TF' has embedding dimension $d' = d + 1$ where d is the embedding dimension of TF .

$$v_\sigma = \begin{pmatrix} x \\ y \\ z \end{pmatrix} \rightarrow v'_\sigma = \begin{pmatrix} x \\ y \\ z \\ f_{\text{on}} \end{pmatrix} \qquad v_\tau = \begin{pmatrix} i \\ j \\ k \end{pmatrix} \rightarrow v'_\tau = \begin{pmatrix} i \\ j \\ k \\ f_{\text{off}} \end{pmatrix}$$

Fig. 3.1: Change in a flagged and unflagged vector

For the explicit construction read Appendix A.1.

This primitive works also as a mechanism to extend the embedding dimension without changing the behavior. If v_{on} and v_{off} are defined as 0, the resulting transformer is the same as the initial with an extra coordinate that is always 0.

3.1.2 Preserve and ignore coordinates through ATTN and FF

When defining new transformers over existing ones we will need to extend the embedding dimension to store new information. In order to preserve the behavior of the original transformers across the attention and feed-forward layers, we need to construct a mechanism that ignores the extra coordinates. For example, this mechanism is used for preserving flags during ATTN and FF layers that don't use them.

Let TF be a transformer of embedding dimension d and assume we are constructing a transformer TF' of embedding dimension $d' = d + 1$. Let W_Q, W_K, W_V, W_O be the matrices of an ATTN layer of TF and assume we want to perform the same layer in TF' but simulating that the i th coordinate does not exist. Notice that the embedding dimension of TF' has one more coordinate than TF 's. If we define the matrices of the layer W'_Q, W'_K, W'_V, W'_O as W_Q, W_K, W_V, W_O by inserting a row and column of zeros at the i th position, we obtain the desired ATTN layer.

Note that because the i th row of W'_O is zero, the output vectors of the ATTN layer are 0 in the i th coordinate. This is the point that makes it possible to not just ignore the value of the i th coordinate, but to not affect it.

To preserve the coordinate without ignoring it, it suffices to set to zero only the corresponding row of W'_O . This construction works for more than one coordinate. To extend it to a set of coordinates, it suffices to set more rows and columns to 0.

The same applies in the feed forward layer, making the i th column and row of the matrices and vectors of the layer zero, makes it ignore the value of the i th coordinate. Furthermore, the result of the FF layer will have a zero in the i th coordinate, thus preserving its value.

$$\begin{pmatrix} A^{i \times r} & 0^{i \times 1} & B^{i \times n-1-r} \\ 0^{1 \times r} & 0 & 0^{1 \times n-1-r} \\ C^{n-1-i \times r} & 0^{n-1-i \times 1} & D^{n-1-i \times n-1-r} \end{pmatrix} \begin{pmatrix} v^{i \times 1} \\ x^{1 \times 1} \\ w^{n-1-i \times 1} \end{pmatrix} = \begin{pmatrix} Av + Bw \\ x \\ Cv + Dw \end{pmatrix} \sim \begin{pmatrix} A & B \\ C & D \end{pmatrix} \begin{pmatrix} v \\ w \end{pmatrix}$$

Fig. 3.2: Matrices for ignoring/preserving the $(i + 1)$ -th coordinate

SC diagrama?

RR how? pongo las matrices? Siento que a veces queda muy cargado. Quizás ahora se entiende mejor?

SC Mmm todavía no es super claro. Una duda: ese 0 nuevo que ponés, no le llega al soft max? Me refiero al attention.

RR No, no afecta al attn. Esto funciona cómo diagrama?

3.1.3 Add vector and constant functions

In this section, we develop a mechanism for adding a fixed vector to a previously flagged one. This technique works even when there is more than one flagged vector.

Assume there is a flag with values $f_{\text{on}} = 1$ and $f_{\text{off}} = 0$. For clarity, suppose the flag is in the last coordinate. Let v be the vector we want to add. We define a mechanism that replaces every h such that $h_d = f_{\text{on}}$ with $h + v$.

The operation is performed with only one FF layer. We construct a layer such that $FF(h) = \mathbb{I}(h_d = f_{\text{on}}) \cdot v$. Said layer is achieved with the following parameters.

$$W_1 = id \quad W_2 = \begin{pmatrix} 0 & \dots & 0 & v_1 \\ \vdots & \ddots & \vdots & \vdots \\ 0 & \dots & 0 & v_d \end{pmatrix} \quad b_1 = b_2 = \begin{pmatrix} 0 \\ \vdots \\ 0 \end{pmatrix}$$

Notice that if the coordinate for the flag had been the k -th instead of the last one, then v would have appeared in the k -th column of W_2 instead of the last.

This technique can be used to apply a constant function to the flagged vectors. This is achieved by exploiting the finite representation system. We can saturate the vector to the ∞ and then make it 0 to finally add our objective. This holds since for all $x, c \in \mathbb{F}_{e,s}$ it happens that $x + \infty + \infty + -\infty + c = c$. Because as mentioned in Section 2.2, $x + \infty + \infty = \infty$ and $\infty + -\infty = 0$.

The constant function $f(h) = v$ is implemented by adding four vectors that perform this trick. First, it is necessary to add the vector consisting entirely of ∞ twice; then, a vector consisting entirely of $-\infty$; and finally, v ; as shown in Equation 3.1.

$$h \leftarrow h + \begin{pmatrix} \infty \\ \vdots \\ \infty \end{pmatrix} + \begin{pmatrix} \infty \\ \vdots \\ \infty \end{pmatrix} + \begin{pmatrix} -\infty \\ \vdots \\ -\infty \end{pmatrix} + v \quad (3.1)$$

3.1.4 Make a vector ignore others in an ATTN layer

We show how to make a vector ignore others in the ATTN layers, a technique that will prove useful in later constructions. A vector h_i ignores vector h_j when the attention score is 0, which can be achieved by making $\langle q_i, k_j \rangle = -\infty$.

We extend the embedding dimension with three more coordinates for preserving the attention score of the other vectors previous to taking softmax. We construct TF' based on TF such that the l -th vector ignores a set of vectors M in an ATTN layer; specifically, we require said layer to satisfy:

$$\langle q'_i, k'_j \rangle = \begin{cases} -\infty & \text{if } i = l \wedge j \in M \\ \langle q_i, k_j \rangle & \text{if } i \neq l \vee j \notin M \end{cases}$$

This can be achieved by saturating the finite representation since $\langle q'_l, k'_i \rangle = \sum_{x=1}^{d+3} (q'_l)_x (k'_i)_x$. If the last two terms of the summation are $-\infty$, then because of the iterative rounding, $\langle q'_l, k'_i \rangle = -\infty$. Since we do not want to change the attention scores of vectors not in M , when $i \notin M$ the last three terms should be 0. The explicit construction achieving this is explained in Appendix A.2

3.1.5 Linear transformations

Although it may seem that transformers are a model that should be able to perform linear transformations in a native way, it is not the case. Applying a linear transformation to a vector is not straightforward, since in every step of ATTN and FF the result of the layer is added to the original vector instead of replacing it.

SC Viste que la gente cheta a esto lo llama “residual vector” o state algo asi. Capaz podés buscarlo y sonar cheto xD.

RR lo llaman residual stream pero para usarlo acá siento que lo debería introducir antes y creo que no es worth.

An intuitive idea would be to apply the transformation $M - I$ when we want to transform the vector x to Mx , since $x + (M - I)x = x + Mx - Ix = Mx$. This does not hold in our model since the addition and multiplication are not commutative in $\mathbb{F}_{e,s}$. To avoid the representation error, we will extend the embedding dimension to get more memory per vector such that we can store in new coordinates the value of Mx and then replace it in the original ones.

We give a construction that applies the linear transformation M to a vector previously flagged. Let h_f be the flagged vector. As mentioned, we need to double the embedding dimension (from d to $2d$) so we can have coordinates to use as a temporary memory. Our construction works with two ATTN layers operating as shown in Figure 3.3 (see Appendix A.3 for the exact construction). The first ATTN layer computes the transformation saving its result in the coordinates $d + 1$ to $2d$. Then, the second layer we move the result to the right place.

$$\widetilde{h}_f = \begin{pmatrix} (h_f)_1 \\ \vdots \\ (h_f)_d \\ 0 \\ \vdots \\ 0 \end{pmatrix} \xrightarrow{\text{first ATTN layer}} \widetilde{h}_f = \begin{pmatrix} 0 \\ \vdots \\ 0 \\ (Mh_f)_1 \\ \vdots \\ (Mh_f)_d \end{pmatrix} \xrightarrow{\text{second ATTN layer}} \widetilde{h}_f = \begin{pmatrix} (Mh_f)_1 \\ \vdots \\ (Mh_f)_d \\ 0 \\ \vdots \\ 0 \end{pmatrix}$$

Fig. 3.3: Visual steps for linear transformations

3.1.6 Max coordinate

We show how to take the maximum of a set of coordinates. This operation will be needed for later constructions.

We give a construction for taking the maximum of two coordinates. This construction can be easily extended to more coordinates. The operation works properly when the set of coordinates has only one maximum. All the coordinates that are not the maximum are replaced by 0.

$$\begin{pmatrix} a \\ b \\ \vdots \end{pmatrix} \rightarrow \begin{cases} \begin{pmatrix} a \\ 0 \\ \vdots \end{pmatrix} & \text{if } a > b \\ \begin{pmatrix} 0 \\ b \\ \vdots \end{pmatrix} & \text{if } b > a \end{cases}$$

The maximum of two coordinates can be computed with the following primitives:

$$\begin{aligned} \begin{pmatrix} a \\ b \\ 0 \\ 0 \\ \vdots \end{pmatrix} &\xrightarrow{\text{FF layer}} \begin{pmatrix} a \\ b \\ \text{relu}(a-b) \\ \text{relu}(b-a) \\ \vdots \end{pmatrix} \xrightarrow[\text{the coordinates with the relu}]{\text{Multiply twice by } \infty} \begin{pmatrix} a \\ b \\ \infty \cdot \mathbb{I}(a > b) \\ \infty \cdot \mathbb{I}(b > a) \\ \vdots \end{pmatrix} \rightarrow \\ &\xrightarrow[\text{coordinates with the comparison}]{\text{Divide by } \infty} \begin{pmatrix} a \\ b \\ \mathbb{I}(a > b) \\ \mathbb{I}(b > a) \\ \vdots \end{pmatrix} \xrightarrow[\text{operation}]{\text{Special}} \begin{pmatrix} a \cdot \mathbb{I}(a > b) \\ b \cdot \mathbb{I}(b > a) \\ 0 \\ 0 \\ \vdots \end{pmatrix} = \begin{cases} \begin{pmatrix} a \\ 0 \\ \vdots \end{pmatrix} & \text{if } a > b \\ \begin{pmatrix} 0 \\ b \\ \vdots \end{pmatrix} & \text{if } b > a \end{cases} \end{aligned}$$

For computing the maximum of a set of coordinates, it suffices to compute a comparison for each pair of the set and multiply each coordinate x by $\mathbb{I}(x > y)$ for all $y \neq x$ in the set.

See appendix A.4 for the implementation of the special operation

3.1.7 Repeat a token once it's printed

For a later proof we will need to assume that a transformer consumes all of its budget. Thus, we will argue that for every transformer TF there is a transformer TF' of the same $\text{CoT}[T(n), d(n)]$ class that computes the same language and consumes all of its budget.

We construct TF' as TF but once it prints a token σ , it repeats it until the budget runs out (or another construction makes it stop). This is achieved by flagging the token σ and applying a constant function at the end of the body of the transformer to all the vectors flagged. The constant function force the probability distribution to have all the probability concentrated in the token σ .

We give the exact construction in appendix A.5

3.1.8 Force to halt after certain number of steps of CoT

Using a strategy similar to the one in Section 3.1.7, we use the PE to flag when a certain position is reached. Then, if the transformer is normalized, modifying the last vector

we decide the outputted distribution. Particularly, we can force it to concentrate the probability in q_{accept} or q_{reject} . We change the last vector with a constant function. The flag inserted by the PE is to decide if the function has to be applied or not.

Notice that this primitive, together with the ability to repeat a token, proves that for every transformer TF in $\text{CoT}[T(n), d(n)]$ there is a transformer TF' in the same class that always consumes all of its budget.

We extend the alphabet with two dummy versions of the decision tokens and construct TF' based on TF. We make TF' work as TF, but we make it print the dummy versions of the decision tokens when TF would have printed the real ones¹. Then, we force TF' to repeat the dummy decision tokens until it reaches the last position of the tape (uses all the budget). At that point, we force TF' to print the real decision token corresponding to the one it was repeating.

3.1.9 Copy one vector into another with ATTN

For the disjunction of transformers we will need to move values from the i th vector to the j th vector. In this section we show how to do it. This operation can only be done if $i < j$, since the attention mechanism only allows for a vector to attend to vectors to its left, i.e., h_k can attend to h_l if and only if $k \leq l$.

Let j be the number of the receiving vector and i the number of the copied, i.e., we copy from h_i to h_j . The operation is done in two steps.

First, we make 0 the coordinates of h_j that will receive the values. This is done with a constant function.

Second, we make h_j only attend to h_i in an attention layer. In said attention layer, we use $W_V = id$ and W_O such that it outputs the desired coordinates of h_i in the target coordinates. Formally, W_O has only 0 in the row k if h_j is not receiving anything from h_i in the k -th coordinate. If h_j is receiving the coordinate l of h_i in the position k , then the k th row of W_O has 0 in every position except the l th.

For example, if we want to copy the 2nd coordinate of h_i to the 3rd coordinate of h_j , W_O would look like 3.2. Then, the vector coming out of the ATTN layer will have 0 in every coordinate, except in the 3rd, where it will be $(h_i)_2$.

$$W_O = \begin{pmatrix} 0 & 0 & 0 \\ 0 & 0 & 0 \\ 0 & 1 & 0 \end{pmatrix} \quad (3.2)$$

3.2 Normalization

For an easier handling of transformers, we define a notion of normalized transformer. This notion will specify properties of the OUTPUT matrix and of the vector consumed by this step.

Definition 3.2.1 (Normalized deterministic transformer). We say a deterministic transformer is normalized if it satisfies two conditions.

- N1: The matrix of the OUTPUT has to be a projection, i.e., it is a diagonal matrix with ones on its diagonal.

¹ This is achieved giving the probability of each decision token to its dummy version.

- N2: The last vector after the last iteration must consist entirely of $-\infty$ in every projected coordinate except one, which must equal ∞ .

The motivation for this definition is to simplify the OUTPUT layer (N1) and to have more control over the probability distribution generated (N2). The exact values of the last vectors are such that after projecting and applying softmax in the OUTPUT layer we get a deterministic distribution with 100% of the probability concentrated in only one token.

For every family of transformers $\text{TF}_{n \in \mathbb{N}} \in \text{CoT}[T(n), d(n)]$ computing the language $\mathcal{L}$, there is a family $\text{TF}'_{n \in \mathbb{N}}$ in the same class, computing the same language, where for all n , TF'_n is normalized. In Appendix B we give an algorithm for normalizing transformers.

3.3 Boolean Operations

In this section we prove that $\text{CoT}[T(n), d(n)]$ classes are closed under Boolean operations. We show that the complement of a language in $\text{CoT}[T(n), d(n)]$ is in $\text{CoT}[T(n), d(n)]$ and that the disjunction of two languages of the same class, stays in the same class.

3.3.1 Complement of a language

Theorem 3.3.1. $\mathcal{L} \in \text{CoT}[T(n), d(n)] \implies \mathcal{L}^c \in \text{CoT}[T(n), d(n)]$

Let $\mathcal{L} \in \text{CoT}[T(n), d(n)]$. Then there is a family of transformers $\{\text{TF}_n\}_{n \in \mathbb{N}}$ of embedding size $d(n)$ that after $T(n)$ steps of CoT prints q_{accept} or q_{reject} depending on if the input word is in the language. For each transformer in the family, we construct a corresponding transformer that behaves identically (internally and externally) but prints the flipped decision token as represented in Figure 3.4. See Appendix C.1




Fig. 3.4: Visual representation of the construction for complement

3.3.2 Disjunction of two languages

Theorem 3.3.2. $\mathcal{L}_1, \mathcal{L}_2 \in \text{CoT}[T(n), d(n)] \implies \mathcal{L}_1 \cup \mathcal{L}_2 \in \text{CoT}[T(n), d(n)]$

Let $\mathcal{L}_1$ and $\mathcal{L}_2$ be languages in $\text{CoT}[T(n), d(n)]$, then there is a family of transformers $\{(\text{TF}_1)_n\}_{n \in \mathbb{N}}$ (resp. $\{(\text{TF}_2)_n\}_{n \in \mathbb{N}}$) of embedding size $d_1(n)$ (resp. $d_2(n)$) $\in O(d(n))$ that after $T_1(n)$ (resp. $T_2(n)$) $\in T(n)$ steps of CoT computes $\mathcal{L}_1$ (resp. $\mathcal{L}_2$).

We construct a family $\{\text{TF}_n\}_{n \in \mathbb{N}}$ that computes $\mathcal{L}_1 \cup \mathcal{L}_2$. It does it in $T_1(n) + T_2(n) + 1 \in O(T(n))$ steps of CoT and it has embedding size $O(d(n))$.

The idea behind the construction is to run first one of the transformers, then run the other and finally compute the disjunction of both results as represented in Figure 3.5. The exact construction for TF is shown in appendix C.2.




Fig. 3.5: Visual representation of the construction for disjunction

4. MAIN RESULTS

In this chapter, we present our results on probabilistic transformers. We show that probabilistic transformers are as strong as the deterministic ones, since probabilistic transformers are capable of simulating deterministic ones. We prove that the probabilistic classes support error reduction by means of a Monte Carlo simulation. Furthermore, we also demonstrate that RCoT classes are closed under Boolean operations.

Finally, we prove upper and lower bounds on the expressive power of probabilistic transformers. These theorems are extensions of Li et al. (2024)'s results for deterministic transformers.

4.1 Simulation of deterministic transformers

In this section we show how to simulate a deterministic transformer with a probabilistic one.

Theorem 4.1.1. *A constant-depth, constant-precision deterministic transformer with embedding dimension $d(n)$ and using $T(n)$ steps of CoT can be simulated by a constant-depth, constant-precision probabilistic transformer with the same embedding dimension and using the same number of CoT steps.*

Proof. Let TF be a deterministic transformer. We construct a probabilistic transformer TF' that accepts the same words in the same number of steps of CoT and with the same embedding size.

As per Section 3.2, we assume TF to be normalized. We define TF' exactly as TF, except for the token selector which by definition is the probabilistic. Now, TF and TF' behave exactly the same. Since TF is normalized, the distribution p generated in every step satisfies the property that there exists $\sigma \in \Sigma$ such that:

$$p(\tau) = \begin{cases} 1 & \text{if } \tau = \sigma \\ 0 & \text{otherwise} \end{cases}$$

Then, sampling this distribution p always returns the token σ , the same token that argmax would return. Next, the deterministic token selector and the probabilistic token selector behave the same. Thus, for all i and α , $\text{TF}^i(\alpha) = \text{TF}'^i(\alpha)$. In particular, this means that $\text{TF}^*(\alpha) = \text{TF}'^*(\alpha)$. Therefore, TF and TF' accept the same words (and in the same number of steps). By construction, the embedding size of TF and TF' is the same. $\square$

From Theorem 4.1.1, we immediately obtain:

Theorem 4.1.2. $\text{CoT}[T(n), d(n)] \subseteq \text{RCoT}[T(n), d(n)]$ for all $T(\cdot)$ and $d(\cdot)$

4.2 Normalization

Recall the definition of normalization for deterministic transformers. We define a deterministic transformer as normalized when it satisfies the conditions N1 and N2.

If we make a probabilistic transformer satisfy N2 it becomes deterministic. This changes the behavior of the transformer, so we cannot claim that for every probabilistic transformer there exists another with the same characteristics that computes the same language and satisfies N2. Thus, for probabilistic transformers, we define normalization solely as the condition N1.

Definition 4.2.1 (Normalized probabilistic transformer). We say a probabilistic transformer is normalized if it satisfies the following condition.

- **N1:** The matrix of the OUTPUT has to be a projection, i.e., it is a diagonal matrix with ones on its diagonal.

Observe that the algorithm presented in Appendix B for constructing equivalent transformers satisfying N1 in the deterministic setting also applies to probabilistic transformers.

4.3 Robustness of our Definition

Recall the definition of language recognition by probabilistic transformers. We say that a family of probabilistic transformers $\{\text{TF}_n\}_{n \in \mathbb{N}}$ recognizes a language $\mathcal{L}$ when

$$\alpha \in \mathcal{L} \iff \Pr[\text{TF}_{|\alpha|}^*(\alpha) = q_{\text{accept}}] > 2/3$$

The choice of the constant $2/3$ may seem somewhat arbitrary. We now show that we can replace $2/3$ with any constant in the interval $[1/2, 1)$.

SC ojo que es estricto, tmb hay que arrelgarlo en la demo

RR está bien así me parece. Aceptar con proba $\geq 1/2$ es lo que queremos

Let $\{\text{TF}_n\}_{n \in \mathbb{N}}$ be a family of probabilistic transformers such that $\alpha \in \mathcal{L} \iff \Pr[\text{TF}_{|\alpha|}^*(\alpha) = q_{\text{accept}}] > c$ for some fixed c in the interval $[1/2, 1)$ and show that it is equivalent to a family $\{\text{TF}'_n\}_{n \in \mathbb{N}}$ that answers correctly with probability $d < 1$.

For each TF_n we construct TF'_n such that for all words of length n , $\alpha \in \mathcal{L} \iff \Pr[\text{TF}_n^*(\alpha) = q_{\text{accept}}] > d$. First, if $d \leq c$, then $\text{TF}_n = \text{TF}'_n$. Let us explore the case where $c < d$.

We construct TF'_n based on TF_n using a Monte Carlo simulation. The idea is similar to the disjunction of languages. TF'_n runs the transformer TF_n sequentially r times and then analyzes whether the majority of the results accept.

The construction presented in Section 3.3.2 for the disjunction of two deterministic transformers can be easily extended to run a larger number of transformers sequentially. This construction also applies for probabilistic transformers.

To compute the disjunction, we analyzed the coordinate f_c at the last step of the computation. It worked as a counter of how many transformers accepted the word. We replicate this strategy to check whether the majority of the r runs of TF_n accepted the word, i.e., $f_c > r/2$. If that is the case, we force TF'_n to accept; otherwise, we force it to reject.

For fixed c and d , only a constant number of runs (depending on c and d) are needed for the construction to work properly. The exact number is computed using Chernoff's bound. The calculations are shown in Appendix D.1.

For $|\alpha| = n$, let $T(n)$ be the number of steps the family $\{\text{TF}_n\}_{n \in \mathbb{N}}$ needs to compute $\mathcal{L}(\alpha)$ and $d(n)$ its embedding dimension. Then, the family $\{\text{TF}'_n\}_{n \in \mathbb{N}}$ takes $O(r \cdot T(n))$ steps to compute $\mathcal{L}(\alpha)$ and uses embedding size $O(r \cdot d(n))$. Thus, the $\text{RCoT}[T(n), d(n)]$ classes are invariant to the choice of the constant in the definition of language recognition by probabilistic transformers.

4.4 Closed by Boolean Operations

We showed in Section 3.3 that deterministic transformers can perform Boolean operations with only a constant overhead; thus, $\text{CoT}[T(n), d(n)]$ classes are closed under Boolean operations. Similar constructions apply to probabilistic transformers. Therefore, we show that $\text{RCoT}[T(n), d(n)]$ classes are closed under complement and disjunction, which together imply closure under Boolean operations.

Theorem 4.4.1 ($\text{RCoT}[T(n), d(n)]$ is closed under complement). $\mathcal{L} \in \text{RCoT}[T(n), d(n)] \implies \mathcal{L}^c \in \text{RCoT}[T(n), d(n)]$

Let $\mathcal{L} \in \text{RCoT}[T(n), d(n)]$ and let $\{\text{TF}_n\}_{n \in \mathbb{N}}$ be the family of probabilistic transformers that computes $\mathcal{L}$. Let us assume all the transformers in the family are normalized. Then, similarly to the deterministic case, we define the family $\{\text{TF}'_n\}_{n \in \mathbb{N}}$ as $\{\text{TF}_n\}_{n \in \mathbb{N}}$, except the OUTPUT matrices have the rows corresponding to q_{accept} and q_{reject} swapped.

Next, for all $\alpha \in \Sigma^*$, if $\Pr[\text{TF}_{|\alpha|}^*(\alpha) = q_{\text{accept}}] > 2/3$, then $\Pr[\text{TF}'_{|\alpha|}(\alpha) = q_{\text{reject}}] > 2/3$, since every time $\text{TF}_{|\alpha|}$ prints q_{accept} , $\text{TF}'_{|\alpha|}$ prints q_{reject} . Therefore, $\{\text{TF}'_n\}_{n \in \mathbb{N}}$ computes $\mathcal{L}^c$ using the same number of CoT steps and the same embedding size $O(d(n))$ as $\{\text{TF}_n\}_{n \in \mathbb{N}}$ to compute $\mathcal{L}$. Thus, $\mathcal{L} \in \text{RCoT}[T(n), d(n)]$ implies $\mathcal{L}^c \in \text{RCoT}[T(n), d(n)]$.

Theorem 4.4.2 ($\text{RCoT}[T(n), d(n)]$ is closed under disjunction). $\mathcal{L}_1, \mathcal{L}_2 \in \text{RCoT}[T(n), d(n)] \implies \mathcal{L}_1 \cup \mathcal{L}_2 \in \text{RCoT}[T(n), d(n)]$

Let $\{\text{TF}_n^1\}_{n \in \mathbb{N}}$ be the family of probabilistic transformers recognizing $\mathcal{L}_1$ and $\{\text{TF}_n^2\}_{n \in \mathbb{N}}$ the one recognizing $\mathcal{L}_2$. We construct the family $\{\text{TF}_n\}_{n \in \mathbb{N}}$ following the same construction as for the disjunction of deterministic transformers (Section 3.3.2). By the result of Section 4.3, we can assume that the probability of error of each family is arbitrary small. We require for both families to answer correctly with probability $4/5$. Then, the probability of $\{\text{TF}_n\}_{n \in \mathbb{N}}$ answering $\mathcal{L}_1 \cup \mathcal{L}_2$ correctly is at least $1/2$ (see Appendix D.2). Therefore, by the construction in Section 4.3, there is a family of transformers in $\text{RCoT}[T(n), d(n)]$ that recognizes $\mathcal{L}_1 \cup \mathcal{L}_2$ in $\text{RCoT}[T(n), d(n)]$.

4.5 Upper Bound

In this section, we show that one forward pass of constant-depth, polynomial-embedding-dimension, and constant-precision probabilistic transformers can be simulated by shallow Boolean circuits.

SC ste resultado se puede probar de forma más directa, no? Como que solo tenés que simular el único paso del transformer y fijarte cuál item es el más probable.

Theorem 4.5.1. *One step of a probabilistic, constant-depth, polynomial-embedding-dimension transformer can be simulated with a circuit in $\text{BPAC}^0[1]$.*

For this proof, we build upon Li et al. (2024)'s proof of Theorem 2.5.1. They already showed that the distribution generator can be simulated by AC^0 circuits. Thus, we only show that the probabilistic token selector can be simulated by $\text{BPAC}^0[1]$ circuits.

Lemma 4.5.1. There is a $\text{BPAC}^0[1]$ circuit that simulates the sampling of a probability distribution for the tokens defined in $\mathbb{F}_{e,s}$.

*Proof Sketch.*¹ We define $|\Sigma|$ ranges in $[0, 1)$, one per token, each of length equal to the token's probability. These ranges can be computed by an AC^0 circuit. Then, we use a seed of constant size to index the ranges. The seed is generated by the random gates and the ranges are indexable with an AC^0 circuit. Thus, the circuit outputs the token σ with probability $p(\sigma)$. $\square$

By Theorem 2.3.4, one step of a probabilistic, constant-depth, polynomial-embedding-dimension transformer can also be simulated with a circuit in AC^0 . And as in Theorem 2.5.2, we now prove that a logarithmic number of CoT steps is not sufficient to go beyond AC^0 , even in a probabilistic transformer.

Theorem 4.5.2. $\text{RCoT}[\log(n), \text{poly}(n)] \subseteq \text{BPAC}^0[\log(n)] \subset \text{AC}^0$

Proof. Since we need a constant number of bits from the seed to simulate each CoT step, we need $O(\log(n))$ random bits to simulate $\log(n)$ steps. Thus, we can enumerate all possible chain-of-thought token sequences ($|\Sigma|^{O(\log(n))} = O(n)$) and seeds $2^{O(\log(n))=O(n)}$ and check when they are consistent. Observe that we need to check each seed with each sequence; thus, we check consistency of $O(n^2)$ pairs (sequence, seed).

In the deterministic transformers it was enough to perform an OR of all the consistent sequences. In the probabilistic setting, that idea does not work, since there could be false positives. Instead, our circuit computes MAJORITY of the consistent runs and outputs its result.

The circuit is clearly in AC^0 , except for the MAJORITY. The MAJORITY problem can be solved in AC^0 only when there is a promise that more than a fraction c of the bits are 1s (or 0s), for a fixed $c \in (1/2, 1]$ (Viola (2009)). Since the probabilistic transformer outputs the correct answer $2/3$ of the time, i.e., for $2/3$ of the seeds, we know that at least $2/3$ of the consistent runs agree. Thus, instead of MAJORITY, we compute $\text{MAJORITY}_{2/3}$, which can be done in AC^0 .

Finally, since there is an AC^0 circuit simulating one forward pass of constant-depth, polynomial-embedding-dimension, and constant-precision probabilistic transformers, there is an AC^0 circuit simulating $O(\log(n))$ steps of CoT of the same set-up. $\square$

Since one forward pass of constant-depth, polynomial-embedding-dimension, and constant-precision probabilistic transformers can be simulated by a circuit in $\text{BPAC}^0[1]$, stacking the circuits of $\text{BPAC}^0[1]$ for each CoT step gives a simulation for $\text{poly}(n)$ steps (as in Corollary 2.5.1). This circuit has polynomial size and uses a polynomial number of random bits; thus, we obtain Corollary 4.5.1.

Corollary 4.5.1. $\text{RCoT}[\text{poly}(n), \text{poly}(n)] \subseteq \text{BPP}/\text{Poly}$

¹ See Appendix E.1 for a complete proof.

4.6 Lower Bound

Building upon Theorem 2.5.3, we prove that constant-depth, constant-precision probabilistic transformers using T steps of CoT with $O(\log(n))$ embedding size can simulate a randomized circuit of size T .

Theorem 4.6.1. *Any family of randomized Boolean circuits of polynomial size $T(n)$ can be simulated by a family of constant-depth, constant-precision probabilistic transformers using $T(n)$ steps of CoT with $O(\log(n))$ embedding size*

Proof Sketch. ² Let $\{C_n\}_{n \in \mathbb{N}}$ be a family of randomized circuits of polynomial size. We construct a family of probabilistic transformers $\{TF_n\}_{n \in \mathbb{N}}$ such that each TF_n simulates C_n . As Definition 2.3.12 states, a randomized circuit C_n accepts α when $\Pr_{r_1 \dots r_{f(n)}}[C'_n(\alpha, r) = 1] > 2/3$, that is, the deterministic circuit C'_n accepts (α, r) for $2/3$ of the possible values of r . Our construction first generates r at random with one step of CoT per bit of r , and then simulates the deterministic circuit C' with input αr (Theorem 2.5.3 and Theorem 4.1.1). $\square$

Theorem 4.6.1 can also be phrased with complexity classes as:

Theorem 4.6.2. *For any polynomial function $T : \mathbb{N}^+ \rightarrow \mathbb{N}^+$, $\text{BPSIZE}(T(n)) \subseteq \text{RCoT}[T(n), \log(n)]$. In particular, $\text{BPP}/\text{Poly} \subseteq \text{RCoT}[\text{poly}(n), \log(n)]$.*

4.7 Relations between models

4.7.1 Deterministic and probabilistic transformers

From Corollary 2.5.1 and Theorem 2.5.4, we deduce the following theorem:

Theorem 4.7.1. $\text{CoT}[\text{poly}(n), \text{poly}(n)] = \text{P}/\text{Poly}$.

In the same way, from Corollary 4.5.1 and Theorem 4.6.2, it follows Theorem 4.7.2.

Theorem 4.7.2. $\text{RCoT}[\text{poly}(n), \text{poly}(n)] = \text{BPP}/\text{Poly}$.

Then, from Corollary 2.3.1 and Theorems 4.7.1 and 4.7.2 we deduce the following:

Theorem 4.7.3. $\text{CoT}[\text{poly}(n), \text{poly}(n)] = \text{RCoT}[\text{poly}(n), \text{poly}(n)]^3$.

This theorem implies that the power of deterministic and probabilistic transformers is the same when analyzing if a polynomial number of steps suffices to compute a language. Thus, unless we are working in fine-grained complexity, we can work either with deterministic or probabilistic transformers.

² See Appendix E.2 for the complete construction.

³ $\text{CoT}[\text{poly}(n), \text{poly}(n)] = \text{P}/\text{PolyBPP}/\text{Poly} = \text{RCoT}[\text{poly}(n), \text{poly}(n)]$

4.7.2 Uniform transformers and Turing machines

A more interesting question arises when we require the transformers to be efficiently constructible.

Observe that the proof of Theorem 2.5.3 yields a polynomial algorithm that, for each circuit, constructs its corresponding deterministic transformer. This motivates the definition of $(\text{Uniform})\text{CoT}[T(n), d(n)]$, which is $\text{CoT}[T(n), d(n)]$ restricted to families of transformers such that there exists a deterministic Turing machine that runs in polynomial time and, given the input length 1^n , outputs the transformer TF_n . Since P/Poly restricted to uniform circuits coincides with P (Theorem 2.3.5), we obtain the following corollary.

Corollary 4.7.1. $\text{P} \subseteq (\text{Uniform})\text{CoT}[\text{poly}(n), \text{poly}(n)]$

Observe that Corollary 4.7.1 and Corollary 2.5.1 together imply Theorem 4.7.4.

Theorem 4.7.4. $\text{P} = (\text{Uniform})\text{CoT}[\text{poly}(n), \text{poly}(n)]$

These results have their analogous in the probabilistic setting. In the same way as in the proof of Theorem 2.5.3, the proof of Theorem 4.6.1 yields a polynomial algorithm that, for each randomized circuit, constructs its corresponding probabilistic transformer. Next, we define $(\text{Uniform})\text{RCoT}[T(n), d(n)]$, which is $\text{RCoT}[T(n), d(n)]$ restricted to families of transformers such that there exists a deterministic Turing machine that runs in polynomial time and, given the input length 1^n , outputs the transformer TF_n . Since BPP/Poly restricted to uniform circuits coincides with BPP (Theorem 2.3.6), we obtain the following corollary.

Corollary 4.7.2. $\text{BPP} \subseteq (\text{Uniform})\text{RCoT}[\text{poly}(n), \text{poly}(n)]$

Observe that Corollaries 4.7.2 and 4.5.1 together imply the following theorem.

Theorem 4.7.5. $\text{BPP} = (\text{Uniform})\text{RCoT}[\text{poly}(n), \text{poly}(n)]$

Since it is not known whether $\text{P} = \text{BPP}$, we cannot deduce the uniform version of Theorem 4.7.3, we can only prove:

Theorem 4.7.6. $\text{P} = \text{BPP} \iff (\text{Uniform})\text{CoT}[\text{poly}(n), \text{poly}(n)] = (\text{Uniform})\text{RCoT}[\text{poly}(n), \text{poly}(n)]$.

Although $\text{P} = \text{BPP}$ is an open question in complexity theory, this theorem is an interesting result. It implies that, when restricted to uniform families of transformers, the constrained access to randomness that transformers have is equivalent in expressive power to the unrestricted access allowed to probabilistic polynomial-time Turing machines. This result suggests that studying probabilistic transformers is particularly relevant in fine-grained complexity, given the widely held (though unproven) belief that randomness yields a polynomial advantage.

4.7.3 Uniform and non-uniform transformers

We show that non-uniformity in deterministic transformers can simulate randomness in uniform probabilistic transformers. Theorems 4.7.5 (a), 2.3.7 (b) and 4.7.1 (c) give the following relation:

$$(\text{Uniform})\text{RCoT}[\text{poly}(n), \text{poly}(n)] \stackrel{(a)}{=} \text{BPP} \stackrel{(b)}{\subseteq} \text{P}/\text{Poly} \stackrel{(c)}{=} (\text{Uniform})\text{CoT}[\text{poly}(n), \text{poly}(n)]$$

Thus, we have Theorem 4.7.7:

Theorem 4.7.7. $(\text{Uniform})\text{RCoT}[\text{poly}(n), \text{poly}(n)] \subseteq \text{CoT}[\text{poly}(n), \text{poly}(n)]$

5. CONCLUSIONS

In this work, we studied the expressive power of constant-depth, constant-precision deterministic and probabilistic transformers. We proved that Boolean operations between languages can be computed in both settings without a blow-up in resources (Sections 3.3 and 4.4); and we were able to extend known expressiveness results from a deterministic setting (Section 2.5) to a randomized one (Sections 4.5 and 4.6).

One of the main results of this work relates the expressive power of deterministic and probabilistic transformers, between them and with other models of computation (Section 4.7). Theorem 4.7.3 says that in the nonuniform setting, the languages recognized by deterministic and probabilistic transformers are the same if we only worry whether a polynomial number of CoT steps are enough. Theorem 4.7.6 implies the same in the uniform setting if $P = BPP$. There is no consensus on said open question, but it is widely believed that access to randomness gives a polynomial advantage; thus studying probabilistic transformers is relevant when working in fine-grained complexity

There is no consensus in the community whether the best abstraction of real-life transformers is given by uniform or non-uniform families. We agree with Fortnow (2025), that it is reasonable to allow non-uniform families. The uniform constraint was artificially introduced to Boolean circuits to make the power of polynomial-size circuits match the power of polynomial-time deterministic Turing machines. This constraint semantically says that there has to be some relation between the circuits of the same family, i.e., there has to be a polynomial-time deterministic Turing machine that computes the family.

In real life, models undergo computationally intensive training in which they learn their weights. This extensive training is analogous to the power conferred by non-uniformity on polynomial-size circuits, which is generally believed to be slightly greater than that of P . Recall that it is widely believed that $P/Poly \not\subseteq NP$.

Consistent with the accepted belief that non-uniformity gives just enough power to simulate randomness, we proved that allowing non-uniformity in deterministic transformers is sufficient to simulate uniform probabilistic transformers (Theorem 4.7.7). It remains an open question whether deterministic transformers can simulate probabilistic ones without the help of non-uniformity. Determining whether such a simulation is possible, and characterizing its overhead, would provide insight into the relationship between the expressive power of the two models and the practical implications of Theorem 4.7.3.

Theorems 2.5.2 and 4.5.2 show that $O(\log(n))$ steps of CoT are insufficient to separate the expressive power of probabilistic and deterministic transformers, both in the uniform and non-uniform cases. We are confident that their power is the same, since AC^0 should be a tight upper bound for one pass of a transformer. We do not believe the ATTN mechanism can be simulated in $NC^0 \subsetneq AC^0$, since it involves a linear number of variables across a constant number of operations.

APPENDIX

A. PRIMITIVES

A.1 Construction for flagging

Let TF be a transformer, we define TF' that behaves like TF but recognize when a vector comes from embedding the token σ . TF' has embedding dimension $d' = d + 1$ where d is the embedding dimension of TF .

Essentially, the components of TF' are the same as TF . The main change is the token embedding function $\text{TE}' : \Sigma \rightarrow \mathbb{F}_{e,s}^{d+1}$ which is defined as

$$(\theta'_{TE}(\tau))_j = \begin{cases} (\theta_{TE}(\tau))_j & \text{if } j \neq d+1 \\ f_{\text{on}} & \text{if } j = d+1 \text{ and } \tau = \sigma \\ f_{\text{off}} & \text{if } j = d+1 \text{ and } \tau \neq \sigma \end{cases}$$

where f_{on} is the value that marks the vector as flagged and f_{off} as unflagged. Those values can be defined chosen arbitrarily for each flag. In this case, the coordinate inserted for the flag is the last one, but it can be added anywhere.

The other components of TF' ignore the new $d+1$ -th coordinate. The position encoding of TF' is equal to the position encoding of TF , except it puts a 0 in the $d+1$ coordinate (not present in TF) independently of the position. The matrices of the ATTN , FF and OUTPUT layers of TF' are the ones of TF extended with 0 in the $d+1$ rows and columns to match the change in embedding dimension.

TF' behaves exactly as TF except it has one more coordinate in the embedding. Said coordinate represents the flag and takes only the values v_{on} or v_{off} in every vector.

A.2 Ignore vectors in Attention layer

Let M be a set of indices. We construct an attention layer of TF' such that h_l ignore all the vectors h_i such that $i \in M$. As mentioned in the main text, we require the attention scores (previous to the softmax) to satisfy the following:

$$\langle q'_l, k'_j \rangle = \begin{cases} -\infty & \text{if } i = l \wedge j \in M \\ \langle q_i, k_j \rangle & \text{if } i \neq l \vee j \notin M \end{cases}$$

We achieved this by saturating the finite representation since $\langle q'_l, k'_i \rangle = \sum_{x=1}^{d+3} (q'_l)_x (k'_i)_x$.

We force $(k'_i)_{d+2}$ and $(k'_i)_{d+3}$ to be $-\infty$ when $i \in M$, and $(q'_i)_{d+2}$ and $(q'_i)_{d+3}$ to be 1 when $i = l$ and 0 otherwise. We also make $(k'_i)_{d+1}$ and $(q'_i)_{d+1}$ to be 0 for all i so the third to last term of the inner product is always 0. With these values, the only attention scores modified are the ones of the l th vector since every other will have 0 in the last three coordinates of the query thus making the last three terms of the inner product 0 and therefore not affecting the score.

This can be achieved with the following flags:

$$(h_i)_{d+1} = \begin{cases} 1 & \text{if } i = l \\ 0 & \text{otherwise} \end{cases} \quad (h_i)_{d+2} = (h_i)_{d+3} = \begin{cases} 1 & \text{if } i \in M \\ 0 & \text{otherwise} \end{cases}$$

Now the query and key matrices are extended as:

$$W'_Q = \begin{pmatrix} W_Q & 0^{\mathbb{R}^{d \times 1}}_{e,s} & 0^{\mathbb{R}^{d \times 1}}_{e,s} & 0^{\mathbb{R}^{d \times 1}}_{e,s} \\ 0^{\mathbb{R}^{d \times 1}}_{e,s} & 0 & 0 & 0 \\ 0^{\mathbb{R}^{d \times 1}}_{e,s} & 1 & 0 & 0 \\ 0^{\mathbb{R}^{d \times 1}}_{e,s} & 1 & 0 & 0 \end{pmatrix} \quad W'_K = \begin{pmatrix} W_K & 0^{\mathbb{R}^{d \times 1}}_{e,s} & 0^{\mathbb{R}^{d \times 1}}_{e,s} & 0^{\mathbb{R}^{d \times 1}}_{e,s} \\ 0^{\mathbb{R}^{d \times 1}}_{e,s} & 0 & 0 & 0 \\ 0^{\mathbb{R}^{d \times 1}}_{e,s} & 0 & -\infty & 0 \\ 0^{\mathbb{R}^{d \times 1}}_{e,s} & 0 & 0 & -\infty \end{pmatrix}$$

Then,

$$\begin{aligned} \langle q'_i, k'_j \rangle &= \sum_{x=1}^{d+3} (q'_l)_x (k'_i)_x = \langle q_i, k_j \rangle + (q'_l)_{d+1} (k'_i)_{d+1} + (q'_l)_{d+2} (k'_i)_{d+2} + (q'_l)_{d+3} (k'_i)_{d+3} \\ &= \langle q_i, k_j \rangle + 0 + (q'_l)_{d+2} (k'_i)_{d+2} + (q'_l)_{d+3} (k'_i)_{d+3} \\ &= \langle q_i, k_j \rangle + (q'_l)_{d+2} (k'_i)_{d+2} + (q'_l)_{d+3} (k'_i)_{d+3} \\ &= \begin{cases} \langle q_i, k_j \rangle + 1 \cdot -\infty + 1 \cdot -\infty & \text{if } i = l \wedge j \in M \\ \langle q_i, k_j \rangle + 1 \cdot 0 + 1 \cdot 0 & \text{if } i = l \wedge j \notin M \\ \langle q_i, k_j \rangle + 0 \cdot -\infty + 0 \cdot -\infty & \text{if } i \neq l \wedge j \in M \\ \langle q_i, k_j \rangle + 0 \cdot 0 + 0 \cdot 0 & \text{if } i \neq l \wedge j \notin M \end{cases} \\ &= \begin{cases} \langle q_i, k_j \rangle + -\infty + -\infty & \text{if } i = l \wedge j \in M \\ \langle q_i, k_j \rangle + 0 + 0 & \text{if } i = l \wedge j \notin M \\ \langle q_i, k_j \rangle + 0 + 0 & \text{if } i \neq l \wedge j \in M \\ \langle q_i, k_j \rangle + 0 + 0 & \text{if } i \neq l \wedge j \notin M \end{cases} \\ &= \begin{cases} -\infty & \text{if } i = l \wedge j \in M \\ \langle q_i, k_j \rangle & \text{otherwise} \end{cases} \end{aligned}$$

Namely, $\langle q'_i, k'_j \rangle = -\infty$ if $i = l$ and $j \in M$, but $\langle q'_i, k'_j \rangle = \langle q_i, k_j \rangle$ otherwise.

A.3 Linear transformations

We show a mechanism to replace a flagged vector with the result of applying the linear transformation M to it. Let f denote the index of the flagged vector. We perform the operation with two consecutive **ATTN** layers, in which h_f only attends to itself. Therefore, after these two layers h_f will be replaced with Mh_f .

We need to double the embedding size from d to $2d$ (plus some flags) since we need to store the result of the transformation in some coordinates (the added $d + 1$ to $2d$) as a middle step. In the first **ATTN** layer, we compute the transformation saving its result in the coordinates $d + 1$ to $2d$. Then, in the second layer we move the result to the right place.

$$\widetilde{h}_f = \begin{pmatrix} (h_f)_1 \\ \vdots \\ (h_f)_d \\ 0 \\ \vdots \\ 0 \end{pmatrix} \xrightarrow{\text{first ATTN layer}} \widetilde{h}_f = \begin{pmatrix} 0 \\ \vdots \\ 0 \\ (Mh_f)_1 \\ \vdots \\ (Mh_f)_d \end{pmatrix} \xrightarrow{\text{second ATTN layer}} \widetilde{h}_f = \begin{pmatrix} (Mh_f)_1 \\ \vdots \\ (Mh_f)_d \\ 0 \\ \vdots \\ 0 \end{pmatrix}$$

We assume that the $d + 1$ to $2d$ coordinates (used as memory) are set to 0 by the layers preceding this mechanism. This can be achieved by maintaining them as 0 since the beginning or transformed to 0 with the primitive described in Section 3.1.3.

Let us construct each of the ATTN layers. As mentioned before, in both layers we will need for $\widetilde{h}_f$ to only attend to itself. Therefore, we define W_O and W_K using the strategy of Section 3.1.4.

Recall from the definition of the ATTN mechanism (1) that the output of the ATTN layer is $h'_i = W_O \sum_{j=1}^a (s_i)_j v_j$. By this point, h'_f is looking like

$$h'_f = W_O \sum_{j=1}^a (s_f)_j v_j = W_O v_f = W_O (W_V \widetilde{h}_f)$$

- For the first layer, defining W_O and W_V as follows

$$W_V = id \quad \quad \quad W_O = \begin{pmatrix} -id^{\mathbb{R}^{d \times d}} & 0^{\mathbb{R}^{d \times d}} \\ M & 0^{\mathbb{R}^{d \times d}} \end{pmatrix}$$

makes us get $h'_f = (-(h_f)_1, \dots, -(h_f)_d, (Mh_f)_1, \dots, (Mh_f)_d)$, which added to $\widetilde{h}_f$ gives us what we were looking for after the first layer.

- For the second ATTN layer, we use the same mechanism for making $\widetilde{h}_f$ only attend to itself. And defining W_V and W_O as follows

$$W_V = id \quad \quad \quad W_O = \begin{pmatrix} 0^{\mathbb{R}^{d \times d}} & id^{\mathbb{R}^{d \times d}} \\ 0^{\mathbb{R}^{d \times d}} & -id^{\mathbb{R}^{d \times d}} \end{pmatrix}$$

we get $h'_f = ((Mh_f)_1, \dots, -(Mh_f)_d, -(Mh_f)_1, \dots, -(Mh_f)_d)$, which added to $\widetilde{h}_f$ gives us what we were looking for after the second layer.

Note that, after these two layers, the values of h_i for all $i \neq f$ have changed in ways that have not been studied in this paper, but this is not relevant to the way we use the primitive.

A.4 Implementation of the special operation

This operation is not a linear transformation; therefore, we need to define a new technique. Observe that one of $\mathbb{I}_a$ and $\mathbb{I}_b$ equals 1 and the other equals 0. Then, $\mathbb{I}_b = 1 - \mathbb{I}_a$ and $\mathbb{I}_a = 1 - \mathbb{I}_b$.

$$\begin{pmatrix} a \\ b \\ \mathbb{I}_a \\ \mathbb{I}_b \end{pmatrix} \xrightarrow[\text{operation}]{\text{Special}} \begin{pmatrix} a \cdot \mathbb{I}_a \\ b \cdot \mathbb{I}_b \\ 0 \\ 0 \end{pmatrix}$$

This operation is performed with three layers of FF. Naturally, after each layer a new vector is added to the original. The vectors we add are the following:

$$\begin{pmatrix} a \\ b \\ \mathbb{I}_a \\ \mathbb{I}_b \end{pmatrix} \xrightarrow[\text{operation}]{\text{Special}} \begin{pmatrix} a + \infty \cdot \mathbb{I}_b + \infty \cdot \mathbb{I}_b + -\infty \cdot \mathbb{I}_b \\ b + \infty \cdot \mathbb{I}_a + \infty \cdot \mathbb{I}_a + -\infty \cdot \mathbb{I}_a \\ \mathbb{I}_a + 0 + 0 + -\mathbb{I}_a \\ \mathbb{I}_b + 0 + 0 + -\mathbb{I}_b \end{pmatrix}$$

Note that if $\mathbb{I}_a = 1$, then $\mathbb{I}_b = 0$, so the first coordinate isn't modified and the second is changed to 0 because of the finite representation properties. In the same way, if $\mathbb{I}_a = 0$, then $\mathbb{I}_b = 1$, so the second coordinate isn't modified and the first is changed to 0.

Let us show how to instantiate the FF layers. The first two have parameters:

$$W_1 = id \quad W_2 = \begin{pmatrix} 0 & 0 & 0 & \infty \\ 0 & 0 & \infty & 0 \\ 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 \end{pmatrix} \quad b_1 = b_2 = \begin{pmatrix} 0 \\ 0 \\ 0 \\ 0 \end{pmatrix}$$

Notice that the output of these FF layer when the input is $(x, y, \mathbb{I}_a, \mathbb{I}_b)$ will be $(\infty \cdot \mathbb{I}_b, \infty \cdot \mathbb{I}_a, 0, 0)$.

The third FF layer has parameters:

$$W_1 = id \quad W_2 = \begin{pmatrix} 0 & 0 & 0 & -\infty \\ 0 & 0 & -\infty & 0 \\ 0 & 0 & -1 & 0 \\ 0 & 0 & 0 & -1 \end{pmatrix} \quad b_1 = b_2 = \begin{pmatrix} 0 \\ 0 \\ 0 \\ 0 \end{pmatrix}$$

which, when the input is $(x, y, \mathbb{I}_a, \mathbb{I}_b)$, outputs $(-\infty \cdot \mathbb{I}_b, -\infty \cdot \mathbb{I}_a, -\mathbb{I}_a, -\mathbb{I}_b)$

RR intenté escribir cómo quedaban las FF instanciadas pero queda medio ilegible. No sé si vale la pena

A.5 Repeat tokens

We show how to make a transformer repeat a token until we decide it to make it stop. This construction is needed to argue that for every transformer TF there is a transformer TF' of the same $\text{CoT}[T(n), d(n)]$ class that computes the same language and consumes all of its budget.

Assume that TF is normalized and let us define TF'. TF' is identical to TF except for the following modifications. Let σ be the token to repeat. We flag the vectors coming from σ with the token embedding. Since we are replicating a token that was printed, i.e., it is not part of the input, its corresponding vector is the last one. This is because we begin repeating it after the first time it is printed.

Since we flag the vectors coming from the token σ we can apply a constant function only to them. This function is applied in the last layer of the body of TF'. Since we assume TF was normalized, the OUTPUT matrix is just a projection. Then, the value of the constant function is what is outputted as the probability distribution (after softmax). If the constant function has value $-\infty$ in every coordinate except in the one associated with the probability of σ , where it has ∞ , the distribution outputted by the distribution generator will have all the probability concentrated in σ . Thus, TF' prints σ .

Observe that this only works if the token does not appear in the input word. This case can be handled by exploiting the non-uniformity of the model. Since the input word has size n , we can use the PE to flag the first n vectors. Then, we can use that flag to set to 0 the coordinate used by the TE for flagging the token to be repeated, i.e., turning off the flag. We make them 0 using a constant function as described in the Section 3.1.3.

B. ALGORITHM FOR NORMALIZING

For normalizing transformers, first we make them meet N1 and then N2.

Algorithm for N1

Notice that every linear transformation $M : \mathbb{R}^d \rightarrow \mathbb{F}_{e,s}^{|\Sigma|}$ with $d \geq |\Sigma|$ can be decomposed into two steps.¹

$$\begin{pmatrix} x_1 \\ \vdots \\ x_d \end{pmatrix} \xrightarrow{M'} \begin{pmatrix} (Mx)_1 \\ \vdots \\ (Mx)_{|\Sigma|} \\ v \end{pmatrix} \xrightarrow{\text{projection}} \begin{pmatrix} (Mx)_1 \\ \vdots \\ (Mx)_{|\Sigma|} \end{pmatrix}$$

where $v \in \mathbb{F}_{e,s}^{d-|\Sigma|}$ can have any values.

Let TF be a transformer and M be the matrix of the OUTPUT layer. Using the decomposition just presented, and the strategy defined in Section 3.1.5 for applying linear transformations, we can change the matrix of the OUTPUT layer to be only a projection. The linear transformation M has to be applied to the last vector, since it is the one going from the body of TF to the OUTPUT layer.

Algorithm for N2

Let TF be a transformer satisfying N1, we give an algorithm for making it also meet N2.

We use the maximum operation described in Section 3.1.6 to make all the positions of the last vector 0 except the one that was the maximum before. Then, if the maximum is positive, multiplying by ∞ , then subtracting 1 from every coordinate and multiplying by ∞ again makes TF meet N2.

$$\begin{pmatrix} 0 \\ \vdots \\ 0 \\ x_{max} \\ 0 \\ \vdots \\ 0 \end{pmatrix} \xrightarrow[\infty]{\text{multiplying by}} \begin{pmatrix} 0 \\ \vdots \\ 0 \\ \infty \\ 0 \\ \vdots \\ 0 \end{pmatrix} \xrightarrow{\text{subtracting 1}} \begin{pmatrix} -1 \\ \vdots \\ -1 \\ \infty - 1 \\ -1 \\ \vdots \\ -1 \end{pmatrix} \xrightarrow[\infty]{\text{multiplying by}} \begin{pmatrix} -\infty \\ \vdots \\ -\infty \\ \infty \\ -\infty \\ \vdots \\ -\infty \end{pmatrix}$$

It could be the case that the maximum was negative, in that case the resulting vector is not what we presented. We show a mechanism for making the maximum strictly positive so we can insert it in the algorithm before subtracting 1.

¹ We can always add more coordinates to the embedding if $d < |\Sigma|$. Since the language is fixed, it will be a constant number of coordinates, thus not changing the class of the transformer.

At this point our vector looks like $(0, \dots, 0, x_{max}, 0, \dots, 0)$, we are going to transform it to $(0, \dots, 0, |x_{max}|, 0, \dots, 0)$. Recall that the coordinates we are interested in are only the first $|\Sigma|$. Let us call v to the block containing these coordinates. Since $\text{relu}(x) + \text{relu}(-x) = |x|$ for all x , our algorithm does the following:

$$\begin{pmatrix} v \\ \vdots \end{pmatrix} \rightarrow \begin{pmatrix} v \\ -v \\ \vdots \end{pmatrix} \rightarrow \begin{pmatrix} \text{relu}(v) \\ \text{relu}(-v) \\ \vdots \end{pmatrix} \rightarrow \begin{pmatrix} \text{relu}(v) + \text{relu}(-v) \\ \vdots \end{pmatrix}$$

where relu is applied position wise.

The first step is done in the same way as the first half of the linear transformations. For the second step, we take max-coordinate (Section 3.1.6) in v and $-v$ independently. This works since all the coordinates of v are 0 except one thus there are only two scenarios. The non-zero coordinate of v is positive, then it will be the maximum (so it won't be modified) and all the others will remain 0; or the non-zero coordinate of v is negative, then the maximum will be 0 and all the coordinates will become 0. Whichever happens in v , the other happens in $-v$. Then, for adding both vectors we can use one **FF** layer with the following parameters. Observe that this works since all the coordinates of $\text{relu}(-v)$ are non-negative.

$$W_1 = id^{d \times d} \quad W_2 = \begin{pmatrix} 0^{|\Sigma| \times |\Sigma|} & id^{|\Sigma| \times |\Sigma|} & 0^{|\Sigma| \times (d-2|\Sigma|)} \\ 0^{d-|\Sigma| \times |\Sigma|} & 0^{d-|\Sigma| \times |\Sigma|} & 0^{d-|\Sigma| \times d-2|\Sigma|} \end{pmatrix} \quad b_1 = b_2 = (0^{d \times d})$$

SC hacer step by step esta **FF**

C. BOOLEAN OPERATIONS FOR DETERMINISTIC TRANSFORMERS

C.1 Complement

Let TF be a transformer, we construct TF' that flips the decision token of TF , i.e., if $\text{TF}^*(\alpha) = q_{\text{accept}}$, then $\text{TF}'^*(\alpha) = q_{\text{reject}}$ and if $\text{TF}^*(\alpha) = q_{\text{reject}}$, then $\text{TF}'^*(\alpha) = q_{\text{accept}}$. This is the only change made to TF , thus $\text{TF}'^i(\sigma) = \text{TF}^i(\sigma)$ for every i when σ is not a decision symbol.

This is done by swapping the rows associated with q_{accept} and q_{reject} in the matrix of the OUTPUT layer of TF . Thus, when the largest probability is associated with q_{accept} in TF , that same number will be associated with q_{reject} in TF' .

Formally, TF' is defined exactly the same as TF except in the OUTPUT layer, where the matrix is different. Let M be the matrix of the OUTPUT layer of TF , let i be the position of the token q_{accept} and j the position of q_{reject} , then the matrix of the OUTPUT layer of TF' will be $M' = PM$ where P is the matrix that flips the rows i and j .

$$M'_{kl} = \begin{cases} 1 & \text{if } k = l, k \neq i \text{ and } k \neq j \\ 1 & \text{if } k = j \text{ and } l = i \\ 1 & \text{if } k = i \text{ and } l = j \\ 0 & \text{otherwise} \end{cases}$$

Since the OUTPUT matrix is the only difference between TF and TF' , they have the same budget and embedding size as TF , therefore, TF' belongs to the same class as TF .

We can define the family $\{\text{TF}'_n\}_{n \in \mathbb{N}}$ where TF'_i is the negation of TF_i constructed as explained in this section. Then $\{\text{TF}'_n\}_{n \in \mathbb{N}}$ computes the complement of $\{\text{TF}_n\}_{n \in \mathbb{N}}$. Since for every $i \in \mathbb{N}$, TF'_i has the same budget and embedding size as TF_i , $\tilde{\mathcal{L}}$ is in $\text{CoT}[T(n), d(n)]$.

C.2 Disjunction

Let $\mathcal{L}_1$ and $\mathcal{L}_2$ be languages in $\text{CoT}[T(n), d(n)]$, then there is a family of transformers $\{(\text{TF}_1)_n\}_{n \in \mathbb{N}}$ (resp. $\{(\text{TF}_2)_n\}_{n \in \mathbb{N}}$) of embedding size $d_1(n)$ (resp. $d_2(n)$) $\in O(d(n))$ that after $T_1(n)$ (resp. $T_2(n)$) $\in T(n)$ steps of CoT computes $\mathcal{L}_1$ (resp. $\mathcal{L}_2$).

In this appendix, we construct a family $\{\text{TF}_n\}_{n \in \mathbb{N}}$ that computes $\mathcal{L}_1 \cup \mathcal{L}_2$. It does it in $T_1(n) + T_2(n) + 1 \in O(T(n))$ steps of CoT and it has embedding size $O(d(n))$.

As mentioned in the main text, TF first runs TF_1 , then TF_2 , and lastly, computes the disjunction of the outputs.

First, observe that for every input α , the tape of a transformer, at the end of the computation, ends up looking as $\alpha\beta q$, where q is the decision token printed by the transformer and β ($\in \Sigma^*$) is the sequence of intermediate tokens printed during the chain of thought. We denote by $\alpha\beta q_1$ the content of the tape of TF_1 at the end of the computation when the input word is α . Similarly, we denote by $\alpha\gamma q_2$ the content of the tape of TF_2 at the end of the computation when the input word is α . The tape of TF , computing the disjunction

of TF_1 and TF_2 , at the end of the computation will have content $\alpha\beta q_1\gamma q_2q$, where q is the decision token representing the disjunction of q_1 and q_2 .

Let us construct TF from TF_1 and TF_2 . Since we cannot change the body of the transformer during the iterations, we are required to use some tricks.

First, we grow the dimension of the embedding for running the body of both transformers at the same time over the same vectors but independently. In the final layers we choose which probability distribution will go to the output layer, this makes TF print the token generated by TF_1 or TF_2 . As always, we assume that the transformers are normalized.

We use the first half of the embedding for the transformer TF_1 and the second for the transformer TF_2 . We have some extra coordinates for flags at the end, but we will ignore them for now.

$$\text{TE}(\sigma)_i = \begin{cases} \text{TE}_1(\sigma)_i & \text{if } i \leq d_1(|\alpha|) \\ \text{TE}_2(\sigma)_{i-d_1(|\alpha|)} & \text{if } d_1(|\alpha|) < i \leq d_2(|\alpha|) \\ 0 & \text{if } d_2(|\alpha|) < i \end{cases}$$

$$\text{PE}(n)_i = \begin{cases} \text{PE}_1(n)_i & \text{if } i \leq d_1(|\alpha|) \\ \text{shiftedPE}_2(n)_{i-d_1(|\alpha|)}^1 & \text{if } d_1(|\alpha|) < i \leq d_2(|\alpha|) \\ \text{flags} & \text{if } d_2(|\alpha|) < i \end{cases}$$

The transformer TF has one layer for each layer of TF_1 and each of TF_2 . The first layers correspond to the body of TF_1 . They are constructed ignoring the half of the embedding storing the coordinates used by TF_2 . The second half of layers are the other way around, they compute the layers of TF_2 and ignore the TF_1 coordinates. Then, all the matrices in the layers of the first part of the body of TF look like V_1 and the ones in the second layers will look like V_2 where W_1 and W_2 are the matrices of the corresponding layers in TF_1 and TF_2 .

$$V_1 = \begin{pmatrix} W_1 & 0 \\ 0 & 0 \end{pmatrix} \quad V_2 = \begin{pmatrix} 0 & 0 \\ 0 & W_2 \end{pmatrix}$$

It is important to notice that the values of the coordinates corresponding to one transformer do not affect the values of the other. Thus, if the tape of the transformer has content γ , then at the end of the body of TF , the last vector would look like (x, y, flags) where x is the last vector at the end of the body of TF_1 when its tape content is γ , and y the one at the end of TF_2 when its tape content is γ . Therefore, if we always choose to transform said vector into $(x, 0, \dots, 0)$, then TF will produce the same tape as TF_1 . This works for generating the first $T_1(|\alpha|)$ tokens which are βq_1 . Thus, we would like to apply this transformation only to the vectors generating them. Therefore, we need to flag them with the position encoding. Since $h_{|\alpha|+i}$ carries the distribution to output in the $(i+1)$ -th step, i.e., the $(|\alpha|+i+1)$ -th token of the tape, the vectors that we need to flag will be $h_{|\alpha|}, h_{|\alpha|+1}, \dots, h_{|\alpha|+T_1(|\alpha|)-1}$.

After running $T_1(|\alpha|)$ steps of this construction, the tape will be $\alpha\beta q_1$. At this point, we want to start generating γ , i.e., the tape of TF_2 . Just changing what we choose before the **OUTPUT** layer does not work, since what the construction will produce would be $\text{TF}_2(\alpha\beta q_1)$, and we need $\text{TF}_2(\alpha)$.

Now we show how to run TF_2 in this construction ignoring the tokens produced by TF_1 . We make TF_2 not attend to the tokens in positions $|\alpha| + 1$ to $T_1(|\alpha|)$ (this can be done using the primitive presented in the Section 3.1.4). We also need for them to be invisible to the portion of the position encoding corresponding to TF_2 , thus we "shift" it in the following way.

$$\text{PE}(n)_i = \begin{cases} \text{PE}_1(n)_i & \text{if } i \leq d_1(|\alpha|) \\ \text{shiftedPE}_2(n)_{i-d_1(|\alpha|)} & \text{if } d_1(|\alpha|) < i \leq d_2(|\alpha|) \\ \text{flags} & \text{if } d_2(|\alpha|) < i \end{cases}$$

$$\text{shiftedPE}_2(n) = \begin{cases} \text{PE}_2(n) & \text{if } n \leq |\alpha| \\ 0 & \text{if } |\alpha| < n < T_1(|\alpha|) \\ \text{PE}_2(n - T_1(|\alpha|)) & \text{if } T_1(|\alpha|) < n \end{cases}$$

The first case is for the positions occupied by the input, we do not want to change their encoding. The second case is for the positions occupied by the tokens generated by TF_1 , we do not care about them. The third case is for the tokens that TF_2 will generate. It should believe that they are in their natural positions, i.e., shifted $T_1(|\alpha|)$ positions left (the positions occupied by the tokens generated by TF_1). We also need to change the flag used for choosing the distribution in the vectors at positions after $|\alpha| + T_1(|\alpha|)$ so TF chooses the distribution corresponding to TF_2 instead of TF_1 .

At this point, the construction almost produces the tape of TF_1 followed by the tape of TF_2 except for one token. We need to do something special for the first token generated by TF_2 . The distribution corresponding to it is in the vector $h_{|\alpha|}$ but the last vector of the first run computing TF_2 tokens is $h_{|\alpha|+T_1(|\alpha|)}$, the one coming from the last token of TF_1 . Therefore, before choosing the distribution, we need to copy the vector $h_{|\alpha|}$ to $h_{|\alpha|+T_1(|\alpha|)}$. Notice that nothing more is needed, since the attention mechanism only attends to vectors to the right, so $h_{|\alpha|}$ is not affected by the vectors coming from the tokens produced by TF_1 . This is done with the operation presented in Section 3.1.9.

With this last fix, TF generates the tape of TF_1 followed by the tape of TF_2 . We now show how to compute q , i.e., the disjunction of q_1 and q_2 . This is easily done making $h_{|\alpha|+T_1(|\alpha|)+T_2(\alpha)}$ only attend to itself and to $h_{|\alpha|+T_1(|\alpha|)}$, i.e., to the vectors coming from q_1 and q_2 .

We add a flag in the token embedding of TF to recognize the decision tokens of TF_1 and TF_2 . The flag is in the coordinate f_c .

$$\text{TE}(\sigma)_{f_c} = \begin{cases} 1 & \text{if } \sigma = q_{\text{accept}} \\ 0 & \text{otherwise} \end{cases}$$

Then, adding $h_{|\alpha|+T_1(|\alpha|)}$ to $h_{|\alpha|+T_1(|\alpha|)+T_2(\alpha)}$, we have 1 or 2 in the coordinate f_c iff TF_1 or TF_2 accepted α . Now we make all of $h_{|\alpha|+T_1(|\alpha|)+T_2(\alpha)}$ 0, except the f_c coordinate (using Section 3.1.3).

Before going to the **OUTPUT** layer we move the value of f_c to the coordinate that will project to the probability of q_{accept} and add some number larger than 0 but smaller than 1 to the coordinate that will define the probability of q_{reject} . Therefore, the token with the maximum probability will be q_{accept} iff TF_1 or TF_2 accepted α , and it will be q_{reject} in the other case.

Then, TF computes the union of the languages of TF_1 and TF_2 . The number of steps of CoT that it needs is exactly the ones made by TF_1 plus the ones made by TF_2 plus 1 extra for taking the disjunction, thus it has the same asymptotic cost as the maximum of TF_1 and TF_2 . The embedding dimension of TF is the same as the dimension of TF_1 plus the dimension of TF_2 plus some constant number of flags, thus in $O(\cdot)$ notation it is the same as the maximum of TF_1 and TF_2 .

Therefore, the language computed by the family of transformers $\{\text{TF}_n\}_{n \in \mathbb{N}}$ is in $\text{CoT}[T(n), d(n)]$. Thus, $\mathcal{L}1 \cup \mathcal{L}2 \in \text{CoT}[T(n), d(n)]$.

D. PROBABILITY CALCULATIONS

D.1 Number of runs for error reduction

The question is how many runs of TF_n are needed, given that it answers correctly with probability $c > 1/2$, in order to obtain the correct answer with probability at least d . We estimate the necessary r using Chernoff's bound.

Theorem D.1.1 (Chernoff's bound). *Let $y_1, \dots, y_r$ be independent Boolean random variables with $\Pr[y_i = 1] = c$ for every $1 \leq i \leq r$. Let $Y = \sum_{i=1}^r y_i$ and let $\delta \in (0, 1)$. Then $\Pr[Y \leq (1 - \delta)rc] \leq e^{-\frac{\delta^2}{4}rc}$ and $\Pr[Y \geq (1 + \delta)rc] \leq e^{-\frac{\delta^2}{4}rc}$.*

We need to find δ and large enough r such that $\Pr[Y > r/2] > d$. Let us analyze how the bound applies to our case. Each y_i represents a run of TF_n . $y_i = 1$ indicates that the i th run of TF_n answered correctly, and $y_i = 0$ that it answered incorrectly. Since TF_n answers correctly with probability higher than c , we have $\Pr[y_i = 1] > c$.

Applying Chernoff's bound, we require δ and r such that $(1 + \delta)rc > r/2$ and $e^{-\delta^2 rc/4} \geq d$.

$$\begin{aligned} (1 + \delta)rc &> \frac{r}{2} & e^{-\frac{\delta^2}{4}rc} &\geq d \\ (1 + \delta)c &> \frac{1}{2} & -\frac{\delta^2}{4}rc &\geq \log_e(d) \\ c + \delta c &> \frac{1}{2} & r &\geq -4 \frac{\log_e(d)}{c\delta^2} \\ \delta &> \frac{1/2 - c}{c} \end{aligned}$$

Since $c \geq 1/2$, we can bound $(1/2 - c)/c$ by 0; thus, choosing any $\delta \in (0, 1)$ satisfies the first equation. Then, with a fixed δ , choosing $r = -4 \log_e(d)/(c\delta^2)$ makes $\Pr[Y > r/2] > d$. Therefore, TF'_n runs TF_n $\lceil -4 \log_e(d)/(c\delta^2) \rceil$ times and accepts the word if more than half of the runs accepted; otherwise, it rejects. Finally, TF'_n answers correctly with probability at least d .

D.2 Probabilities needed for disjunction

Let $\{\text{TF}_n^1\}_{n \in \mathbb{N}}$ be the family of probabilistic transformers recognizing $\mathcal{L}^1$ and $\{\text{TF}_n^2\}_{n \in \mathbb{N}}$ the family recognizing $\mathcal{L}^2$. Let p_1 be the probability of $\{\text{TF}_n^1\}_{n \in \mathbb{N}}$ answering correctly and p_2 the probability of $\{\text{TF}_n^2\}_{n \in \mathbb{N}}$ answering correctly.

We construct the family $\{\text{TF}_n\}_{n \in \mathbb{N}}$ that runs $\{\text{TF}_n^1\}_{n \in \mathbb{N}}$ and $\{\text{TF}_n^2\}_{n \in \mathbb{N}}$ and outputs its disjunction.

What is the probability of $\{\text{TF}_n\}_{n \in \mathbb{N}}$ answering correctly $\mathcal{L}^1(\alpha) \vee \mathcal{L}^2(\alpha)$? For each word α , there are four possible cases ($\alpha \in \mathcal{L}^1 \wedge \alpha \in \mathcal{L}^2, \alpha \in \mathcal{L}^1 \wedge \alpha \notin \mathcal{L}^2, \alpha \notin \mathcal{L}^1 \wedge \alpha \in \mathcal{L}^2, \alpha \notin \mathcal{L}^1 \wedge \alpha \notin \mathcal{L}^2$). In Figure D.1, we show the probability of TF^1 and TF^2 accepting or rejecting α in each of the cases.

TF accepts whenever TF^1 or TF^2 accepts. Let us analyze case by case the probability of TF answering correctly.

RR: ya sé que estas tablas se ven mal. Tengo que re-visarles el estilo

	TF ¹ accepts	TF ¹ rejects		TF ¹ accepts	TF ¹ rejects
TF ² accepts	$p_1 p_2$	$(1 - p_1) p_2$	TF ² accepts	$p_1 (1 - p_2)$	$(1 - p_1) (1 - p_2)$
TF ² rejects	$p_1 (1 - p_2)$	$(1 - p_1) (1 - p_2)$	TF ² rejects	$p_1 p_2$	$(1 - p_1) p_2$

Case 1: $\alpha \in \mathcal{L}_1 \wedge \alpha \in \mathcal{L}_2$ Case 2: $\alpha \in \mathcal{L}_1 \wedge \alpha \notin \mathcal{L}_2$

	TF ¹ accepts	TF ¹ rejects		TF ¹ accepts	TF ¹ rejects
TF ² accepts	$(1 - p_1) p_2$	$p_1 p_2$	TF ² accepts	$(1 - p_1) (1 - p_2)$	$p_1 (1 - p_2)$
TF ² rejects	$(1 - p_1) (1 - p_2)$	$p_1 (1 - p_2)$	TF ² rejects	$(1 - p_1) p_2$	$p_1 p_2$

Case 3: $\alpha \notin \mathcal{L}_1 \wedge \alpha \in \mathcal{L}_2$ Case 4: $\alpha \notin \mathcal{L}_1 \wedge \alpha \notin \mathcal{L}_2$ Fig. D.1: Probabilities of all possible behaviors of TF¹ and TF²

- Case 1: $p_1 + p_2 - p_1 p_2 = p_1 p_2 + (1 - p_1) p_2 + p_1 (1 - p_2)$
- Case 2: $1 - p_2 + p_1 p_2 = p_1 (1 - p_2) + (1 - p_1) (1 - p_2) + p_1 p_2$
- Case 3: $1 - p_1 + p_1 p_2 = (1 - p_1) p_2 + p_1 p_2 + (1 - p_1) (1 - p_2)$
- Case 4: $p_1 p_2$

Then, if $p_1 = p_2 = 4/5$, the probability of TF answering correctly in each case is larger than $1/2$.

- Case 1: $24/25 > 1/2$
- Case 2: $21/25 > 1/2$
- Case 3: $21/25 > 1/2$
- Case 4: $16/25 > 1/2$

Then, if $\{\text{TF}_n^1\}_{n \in \mathbb{N}}$ and $\{\text{TF}_n^2\}_{n \in \mathbb{N}}$ answer correctly with probability $4/5$, the probability of $\{\text{TF}_n\}_{n \in \mathbb{N}}$ answering correctly $\mathcal{L}_1 \cup \mathcal{L}_2$ is at least $1/2$. Therefore, by the construction in Section 4.3, there is a family of transformers in $\text{RCoT}[T(n), d(n)]$ that recognizes $\mathcal{L}_1 \cup \mathcal{L}_2$ in $\text{RCoT}[T(n), d(n)]$.

E. PROOFS OF PROBABILISTIC BOUNDS

E.1 Sampling the distribution with a $\text{BPAC}^0[1]$ circuit

Let p be the probability distribution computed by the distribution generator algorithm of a probabilistic, constant-depth, polynomial-embedding-dimension transformer. In this section we show that sampling p can be done with a constant-depth and polynomial-size random circuit using only a constant number of random gates.

We can sample from the distribution p by drawing a uniform random variable on the interval $[0, 1)$. We refer to the value of this variable as the seed. This is because we can partition $[0, 1)$ into disjoint intervals, one per token, each of length equal to the token's probability. Let p be the probability computed by the distribution generator. Let $\Sigma = \{\sigma_1, \dots, \sigma_{|\Sigma|}\}$. Then the ranges are defined as $\text{range}(\sigma_i) = [\sum_{j=1}^{i-1} p(\sigma_j), \sum_{j=1}^i p(\sigma_j))$.

For example, if $\Sigma = \{t_1, t_2, t_3\}$ and p is such that $p(t_1) = 0.33$, $p(t_2) = 0.22$, $p(t_3) = 0.45$, then the ranges would be $[0, 0.33)$ for t_1 , $[0.33, 0.55)$ for t_2 , and $[0.55, 1)$ for t_3 .

Observe that although the numbers $p(\sigma_j)$ are representable in $\mathbb{F}_{e,s}$, the summations are not performed with iterative rounding. They are computed with the maximum precision needed to avoid rounding error; this precision can be bounded as a function of $|\Sigma|$, e and s alone, independently of the values of $p(\sigma_i)$.

Since the precision for the numbers is fixed and the number of terms of the summations are bounded by the size of the alphabet, the ranges can be computed by circuits in AC^0 .

Let us analyze the seed. It is just a random number in the range $[0, 1)$. We can interpret it as a binary string. How many bits of the seed are needed to determine which range it belongs to? It suffices to examine as many bits of the seed as the logarithm of the smallest probability.

Since p comes from the distribution generator, for all tokens, $p(\sigma)$ has to be representable in $\mathbb{F}_{e,s}$. Then, the smallest non-empty range has size at least ϵ , where ϵ is the smallest positive number representable in $\mathbb{F}_{e,s}$. Let r be the number of bits needed for representing ϵ in binary. Therefore, knowing the first r bits of the binary representation of the seed is enough to sample p . Thus, we only need r random bits for sampling p .

Note that r only depends on the sizes (e and s) of the significand and exponent of $\mathbb{F}_{e,s}$. Thus, r is constant for constant-precision transformers.

Observe that once the ranges are computed, they act as a lookup table with one entry for each number in $[0, 1)$ whose binary representation has r bits. Since r is a constant, the ranges can be interpreted as a constant-size lookup table.

Therefore, since the ranges can be computed by AC^0 circuits and the lookup table has constant size, sampling from p can be performed by a circuit in $\text{BPAC}^0[1]$.

E.2 Simulation of randomized circuits by probabilistic transformers

Let $\{C_n\}_{n \in \mathbb{N}}$ be a family of randomized circuits of polynomial size. We construct a family of probabilistic transformers $\{\text{TF}_n\}_{n \in \mathbb{N}}$ such that each TF_n simulates C_n . As mentioned in Section 2.3, a randomized circuit C_n accepts α when $\Pr_{r_1 \dots r_{f(n)}}[C'_n(\alpha, r) = 1] > 2/3$, that is, the deterministic circuit C'_n accepts (α, r) for $2/3$ of the possible values of r . Our

construction first generates r at random and then simulates the deterministic circuit using the construction of Li et al. (2024) for Theorem 2.5.3.

Let TF'_n be the deterministic transformer simulating C'_n in $|C'_n|$ steps of CoT. Then, $\text{TF}'_n(\alpha r) = C'_n(\alpha, r)$.

We assume TF'_n to be normalized. Then, its OUTPUT matrix is a projection.

TF'_n computes the i -th gate when the tape contains $i - 1$ elements, i.e., when the values of all gates with lower index have been computed. It computes C'_n starting by the gate with index $n + f(n) + 1$. By our assignment of the indices, the first Boolean gate is the one with index $n + f(n) + 1$. Thus, TF'_n begins meaningful computation once its tape contains $n + f(n)$ printed tokens.

We construct TF_n that simulates C_n by modifying TF'_n . Observe that since TF'_n is normalized, the change in token selector from a deterministic to a probabilistic selector does not affect its behavior. We define TF_n exactly as the probabilistic version of TF'_n , with one modification: it ignores the distributions generated by TF'_n in the first $f(n)$ steps of CoT.

The value in each position in the tape of TF_n corresponds to the value of a gate in C_n . The positions may hold the value 1, the value 0, or remain empty if their value has not yet been computed. The first n positions are the bits of α . The following $f(n)$ are the bits of r . TF'_n can generate the tokens for the positions after $f(n)$, i.e., compute the values of the Boolean gates. The probability distributions TF'_n generates in the first $f(n)$ steps, i.e., the distributions corresponding to the tokens representing r , do not correspond to what we want (since from TF'_n 's perspective, they are part of the input). Thus, using the PE, we flag positions n through $n + f(n) - 1$ so that a constant function can be applied to them at the end of the body, forcing the probability distribution for tokens at positions $n + 1$ through $n + f(n)$ to be the uniform distribution over $\{0, 1\}$.

Then, after $f(n)$ steps of CoT, in the tape of TF_n there is written α followed by $f(n)$ random bits. From this point onward, TF'_n simulates $C'_n(\alpha, r)$.

TF_n accepts α if it accepts for at least $2/3$ of the possible runs. TF_n uses randomness only to generate the first $f(n)$ tokens. It generates them following a uniform distribution. Then, $\Pr[\text{TF}_n(\alpha) = 1] = \Pr_r[\text{TF}'_n(\alpha r) = 1] = \Pr_r[C'_n(\alpha, r) = 1] = \Pr[C_n(\alpha) = 1]$. Therefore, $\Pr[\text{TF}_n(\alpha) = 1] > 2/3$ if and only if $\Pr[C_n(\alpha) = 1] > 2/3$.

BIBLIOGRAPHY

- Arora, S. and Barak, B. (2009). Computational Complexity A Modern Approach. Cambridge University Press.
- Bick, A., Blandin, A., and Deming, D. J. (2026). The rapid adoption of generative AI. Management Science.
- Charles, H. B. and Gill, J. (1981). Relative to a random oracle a , $p^a \neq np^a \neq conp^a$ with probability 1*. SIAM Journal on Computing, 10(1):96–113.
- Demir, S. and Dogan, Z. (2025). Optimal attention temperature enhances in-context learning under distribution shift. arXiv preprint arXiv:2511.01292.
- Fortnow, L. (2025). AI and the Power of Nonuniform Circuits. <https://blog.computationalcomplexity.org/2025/10/ai-and-power-of-nonuniform-circuits.html>. Computational Complexity Weblog, accessed July 6, 2026.
- Furst, M., Saxe, J. B., and Sipser, M. (1984). Parity, circuits, and the polynomial-time hierarchy. Mathematical systems theory, 17(1):13–27.
- Goldberg, D. (1991). What every computer scientist should know about floating-point arithmetic. ACM computing surveys (CSUR), 23(1):5–48.
- Hesse, W. (2001). Division is in uniform tc0. In International Colloquium on Automata, Languages, and Programming, pages 104–114. Springer.
- IEEE (2008). Ieee standard for floating-point arithmetic. IEEE Std 754-2008, pages 1–70.
- Kojima, T., Gu, S. S., Reid, M., Matsuo, Y., and Iwasawa, Y. (2022). Large language models are zero-shot reasoners. Advances in neural information processing systems, 35:22199–22213.
- Li, Z., Liu, H., Zhou, D., and Ma, T. (2024). Chain of thought empowers transformers to solve inherently serial problems. In The Twelfth International Conference on Learning Representations.
- Liang, W., Zhang, Y., Wu, Z., Lepp, H., Ji, W., Zhao, X., Cao, H., Liu, S., He, S., Huang, Z., Yang, D., Potts, C., Manning, C. D., and Zou, J. Y. (2024). Mapping the increasing use of llms in scientific papers.
- Lin, T., Wang, Y., Liu, X., and Qiu, X. (2022). A survey of transformers. AI Open, 3:111–132.
- Madaan, A. and Yazdanbakhsh, A. (2022). Text and patterns: For effective chain of thought, it takes two to tango. arXiv preprint arXiv:2209.07686.
- Merrill, W. and Sabharwal, A. (2023). The parallelism tradeoff: Limitations of log-precision transformers. Transactions of the Association for Computational Linguistics, 11:531–545.

- Merrill, W. and Sabharwal, A. (2026). A little depth goes a long way: The expressive power of log-depth transformers. Advances in Neural Information Processing Systems, 38:95315–95339.
- Microsoft AI Economy Institute (2026). Global ai adoption in 2025: A widening digital divide. Report, Microsoft. AI Diffusion Report, 2025 H2.
- Olsson, C., Elhage, N., Nanda, N., Joseph, N., DasSarma, N., Henighan, T., Mann, B., Askell, A., Bai, Y., Chen, A., et al. (2022). In-context learning and induction heads. arXiv preprint arXiv:2209.11895.
- Pérez, J., Marinković, J., and Barceló, P. (2019). On the turing completeness of modern neural network architectures. arXiv preprint arXiv:1901.03429.
- Stryker, C. and Bergmann, D. (2026). What is a transformer model? IBM Think.
- Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, L., and Polosukhin, I. (2017a). Attention is all you need. Advances in neural information processing systems, 30.
- Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, L. u., and Polosukhin, I. (2017b). Attention is all you need. In Guyon, I., Luxburg, U. V., Bengio, S., Wallach, H., Fergus, R., Vishwanathan, S., and Garnett, R., editors, Advances in Neural Information Processing Systems, volume 30. Curran Associates, Inc.
- Viola, E. (2009). On approximate majority and probabilistic time. Computational Complexity, 18(3):337.
- Wang, B., Min, S., Deng, X., Shen, J., Wu, Y., Zettlemoyer, L., and Sun, H. (2023). Towards understanding chain-of-thought prompting: An empirical study of what matters. In Rogers, A., Boyd-Graber, J., and Okazaki, N., editors, Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers), pages 2717–2739, Toronto, Canada. Association for Computational Linguistics.
- Wei, J., Wang, X., Schuurmans, D., Bosma, M., Xia, F., Chi, E., Le, Q. V., Zhou, D., et al. (2022). Chain-of-thought prompting elicits reasoning in large language models. Advances in neural information processing systems, 35:24824–24837.
- Yang, A., Watson, C., Xue, A., Bhattamishra, S., Llarena, J., Merrill, W., Ferreira, E. D. S., Svete, A., and Chiang, D. (2025). The transformer cookbook.