



UNIVERSIDAD DE BUENOS AIRES  
FACULTAD DE CIENCIAS EXACTAS Y NATURALES  
DEPARTAMENTO DE COMPUTACIÓN

[your title here]

Tesis de Licenciatura en Ciencias de la Computación

Rafael Romani

Director: Santiago Cifuentes

Codirector: Santiago Figueira

Buenos Aires, 2026

[YOUR TITLE HERE]

[Acá va el resumen]

**Palabras claves:** [acá van las palabras clave]

[YOUR TITLE HERE]

[write your abstract here]

**Keywords:** [write your keywords here]

*[acá van los agradecimientos]*

# CONTENTS

|                                                                                             |    |
|---------------------------------------------------------------------------------------------|----|
| 1. Introduction . . . . .                                                                   | 1  |
| 1.1 Uso y poder de los LLMS . . . . .                                                       | 1  |
| 1.2 Pero qué es un LLM? . . . . .                                                           | 2  |
| 1.3 Formalizaciones previas y resultados que ignoran los aspectos probabilísticos . . . . . | 2  |
| 1.4 Objetivos y contribuciones . . . . .                                                    | 2  |
| 2. Theoretical Framework . . . . .                                                          | 3  |
| 2.1 Transformers . . . . .                                                                  | 3  |
| 2.2 Finite precision modeling . . . . .                                                     | 5  |
| 2.3 Transformer Complexity classes . . . . .                                                | 7  |
| 2.4 Circuit Complexity . . . . .                                                            | 8  |
| 2.5 Known results . . . . .                                                                 | 8  |
| 3. General properties . . . . .                                                             | 9  |
| 3.1 Transformers Primitives . . . . .                                                       | 9  |
| 3.1.1 Flagging Tokens and Positions . . . . .                                               | 9  |
| 3.1.2 Preserve and ignore coordinates through ATTN and FF . . . . .                         | 10 |
| 3.1.3 Add vector and constant functions . . . . .                                           | 11 |
| 3.1.4 Make a vector ignore others in an ATTN layer . . . . .                                | 11 |
| 3.1.5 Linear transformations . . . . .                                                      | 12 |
| 3.1.6 Max coordinate . . . . .                                                              | 14 |
| 3.1.7 Repeat a token once it's printed . . . . .                                            | 15 |
| 3.1.8 Force to halt after certain number of steps of CoT . . . . .                          | 16 |
| 3.1.9 Copy one vector into another with ATTN . . . . .                                      | 16 |
| 3.2 Normalization . . . . .                                                                 | 17 |
| 3.2.1 Algorithm for normalizing . . . . .                                                   | 17 |
| 3.3 Boolean Operations . . . . .                                                            | 19 |
| 3.3.1 Complement of a language . . . . .                                                    | 19 |
| 3.3.2 Disjunction of two languages . . . . .                                                | 20 |
| 4. Main results . . . . .                                                                   | 23 |
| 4.1 Probabilistic can simulate det . . . . .                                                | 23 |
| 4.2 Normalization . . . . .                                                                 | 23 |
| 4.3 Robustness of our Definition . . . . .                                                  | 24 |
| 4.4 Closed by Boolean Operations . . . . .                                                  | 25 |
| 4.5 Upper Bound . . . . .                                                                   | 26 |
| 4.6 Lower Bound . . . . .                                                                   | 26 |
| 5. Conclusions . . . . .                                                                    | 27 |

# 1. INTRODUCTION

## 1.1 Uso y poder de los LLMs

In 2017, a group of researchers working at Google published a paper introducing a technology that revolutionized society. In *Attention is All You Need*, they presented the transformers, today's core of large language models (LLMs) and generative AI.

This architecture was build intended as text sequence-to-sequence model for machine translation but today has applications in various fields as natural language processing, computer vision or speech recognition.

The first model implementing transformers open for public use was ChatGPT by OpenAI. It made its debut in November 2022 and by December 2024 it reported 300 million global weekly users. Now there are many models competing for this market as Claude, from Anthropic, or Gemini from Google.

These models can be all be grouped under the generative AI tools. This is a technology that is arguably the most rapid technology adopted by society. In the second half of 2025, roughly one in six people worldwide were using generative AI tools. not sure how to cite this <https://www.microsoft.com/en-us/research/wp-content/uploads/2026/01/Microsoft-AI-Diffusion-Report-2025-H2.pdf>

LLMs are being used not only as chatbots for asking small questions, but are shaping the industry and future. For example, [?] found that more than 17% of the Computer Science's papers published between January 2020 and February 2024 on the *arXiv*, *bioRxiv*, and *Nature* portfolio journals had sentences heavily edited by ChatGPT.

### Missing usage in industry

So naturally, the question about their computing power pops up. Are this kind of models really capable of doing the stuff for what they are being used?

- well-suited to parallelism enables transformer models to take full advantage of the power and speed offered by GPUs during both training and inference. This possibility, in turn, unlocked the opportunity to train transformer models on unprecedentedly massive datasets through self-supervised learning.
- chatgpt, claude, gemini,
- every day users
- number of users?
- use in the industry?
- other uses besides the chatbots?
- no usamos normalizacion porque andie lo usa, es complicado. Si se achican los bitsize empieza a tener menos sentido

SF: Lo citás como miscellaneous

cite this <https://www.ibm.com> model

1.2 Pero qué es un LLM?

1.3 Formalizaciones previas y resultados que ignoran los aspectos probabilísticos

1.4 Objetivos y contribuciones

## 2. THEORETICAL FRAMEWORK

### 2.1 Transformers

In this work we are going to analyze transformers as language recognizers. A transformer is a machine that works over an alphabet  $\Sigma$ . It has a single tape from which it reads the input and onto which it appends additional tokens in an autoregressive manner. There are two distinguished tokens called **decision symbols**. Once the transformer writes one of those, it halts. We say that a transformer accepts a word if, when it halts, the decision symbol printed is the accepting one ( $q_{accept}$ ). Analogously, a transformer rejects a word if the decision symbol printed is the rejecting one ( $q_{reject}$ ).

In each step, the transformer maps the content of the tape to a probability distribution over  $\Sigma$ . Then, depending on this distribution, the transformer prints the next token. The transformer iterates this process until it halts. These iterations are called **chain of thought**.

We split the process of generating the next token into two separate steps. First, the computation of the probability distribution over the alphabet which we call distribution generation.

**SF** así como hiciste con cot, definí nuevos términos con un estilo especial. Yo uso defstlye como comando y después decido si es itálica, negrita, etc

**RR** defino distribution generation y token selection más abajo y los puse en negrita. Debería marcarlos también acá?

**RR:**  
Agregar una cita

Second, the selection of a token based on the probability distribution, which we call token selection. In this work we analyze and compare two different algorithms for token selection.

Most prior work only considers transformers that operate over an infinite tape. This is far from the real models, thus the transformers studied in this thesis work over a finite tape. The size of the tape is defined, for each input size  $n$ , as  $\text{budget}(n) + n$ . It is important to remark that the size of the tape bounds the number of steps of CoT that the transformer makes. For an input word of length  $n$ , no computation can take more than  $\text{budget}(n)$  steps because every step writes a token in the tape.

The process of distribution generation is dependent on many parameters. In the real models, these parameters are the ones calculated during the training.

We refer to them as  $\theta$  and consider them as part of the definition of a transformer.

We denote by  $\text{TF}(\alpha)$  the token outputted after one step of the computation of the transformer  $\text{TF}$  when the tape content is  $\alpha \in \Sigma^*$ . The autoregressive token generation is defined recursively. The token outputted after  $i$  steps is  $\text{TF}^i(\alpha) := \text{TF}^{i-1}(\alpha \text{TF}(\alpha))$  with base case  $\text{TF}^1(\alpha) := \text{TF}(\alpha)$ . In other words, the output of the  $i$ th step is  $\sigma_i = \text{TF}^i(\alpha) = \text{TF}(\alpha \sigma_1 \dots \sigma_{i-1})$  where  $\sigma_j \in \Sigma$  for all  $j$ .

We denote by  $\text{TF}^*(\alpha)$  the token written by the transformer when it halts, i.e., at the end of the computation, when the input is  $\alpha$ . As said before, the transformer only halts when it prints a decision symbol ( $q_{accept}$  or  $q_{reject}$ ). Therefore,  $\text{TF}^*(\alpha) \in \{q_{accept}, q_{reject}\}$ .

As stated before, the **distribution generation** is a mapping of the content of the tape to a probability distribution over  $\Sigma$ . There are different algorithms for this process in the



literature. All transformers considered in this work employ the one presented in [?]. This algorithm consists of four parts: a token embedding layer (TE), a position encoding layer (PE), an output linear layer (OUTPUT) and a stack of  $L$  identical decoder layers, where  $L$  is also called the depth of the model. Each decoder layer has two sub-layers: a multi-head self-attention layer (ATTN) and a position-wise fully-connected feed-forward network (FF). Each part and layer has its own trainable parameters. Then, we can define the distribution generation algorithm by its parameters  $\theta = \{\theta_{\text{TE}}, \theta_{\text{PE}}, \{\theta_{\text{ATTN}}^l, \theta_{\text{FF}}^l\}_{l=0}^{L-1}, \theta_{\text{OUTPUT}}\}$ . The algorithm works over vectors of numbers of certain dimension  $d$ . The dimension is dependent on the size of the input word. We call  $d(n)$  to the **embedding size** when the input word has size  $n$ .

**Token Embedding:** It is a mapping from  $\Sigma$  to vectors in  $\mathbb{R}^d$ . This mapping is defined by  $\theta_{\text{TE}} \in \mathbb{R}^{d \times |\Sigma|}$ . For all  $\sigma \in \Sigma$ , we note  $\theta_{\text{TE}}(\sigma) \in \mathbb{R}^d$  to the mapping of token  $\sigma$ .

**Position Encoding:** It is a mapping from the positions of the tape to vectors in  $\mathbb{R}^d$ . This mapping is defined by  $\theta_{\text{PE}} \in \mathbb{R}^{d \times \text{budget}(n)+n}$ . For all position of the tape  $1 \leq i \leq \text{budget}(n) + n$  we note  $\theta_{\text{PE}}(i) \in \mathbb{R}^d$  to the mapping of position  $i$ .

**SC** Intuición de todas estas capas?

**RR** Lo hago en la intro? Me parece el mejor lugar para hacerlo

**Self Attention Layer:** It is a mapping from  $(\mathbb{R}^d)^a$  to  $(\mathbb{R}^d)^a$ . Given parameters  $\theta_{\text{ATTN}} = (W_Q, W_K, W_V, W_O) \in \mathbb{R}^{d \times d} \times \mathbb{R}^{d \times d} \times \mathbb{R}^{d \times d} \times \mathbb{R}^{d \times d}$ , the layer is defined by the algorithm 1. This layer is defined only as a single head attention layer because a multi-head attention layer with any constant number of heads can be simulated with multiple single-head attention layers [?].

---

**Algorithm 1** Causal Self-Attention, ATTN

---

**Input:** Parameter  $\theta_{\text{ATTN}} = (W_Q, W_K, W_V, W_O)$  and embedding  $h = (h_1, \dots, h_a) \in \mathbb{R}^d \times \dots \times \mathbb{R}^d$ .

**Output:** Embedding  $h' = (h'_1, \dots, h'_a) := \text{ATTN}_{\theta_{\text{ATTN}}}(h_1, \dots, h_a)$ .

- 1:  $q_i := W_Q h_i, k_i := W_K h_i, v_i := W_V h_i, \forall i \in [1, \dots, a]$
  - 2:  $s_i := \text{softmax}(\langle q_i, k_1 \rangle, \dots, \langle q_i, k_i \rangle) \parallel (0, \dots, 0)$
  - 3:  $h'_i := W_O \sum_{j=1}^a (s_i)_j v_j$
- 

**Feed Forward Network:** It is a mapping from  $\mathbb{R}^d$  to  $\mathbb{R}^d$ . Given  $\theta_{\text{FF}} = (W_1, b_1, W_2, b_2) \in \mathbb{R}^{d \times d} \times \mathbb{R}^d \times \mathbb{R}^{d \times d} \times \mathbb{R}^d$ , the layer is defined as  $\text{FF}_{\theta_{\text{FF}}}(h) := W_2 \text{relu}(W_1 h + b_1) + b_2$ , where  $\text{relu} : \mathbb{R}^d \rightarrow \mathbb{R}^d$  replaces each position  $x_i$  by  $\max(0, x_i)$ .

**Output Layer:** It is a mapping from  $\mathbb{R}^d$  to a probability distribution over  $\Sigma$ . Given  $\theta_{\text{OUTPUT}} \in \mathbb{R}^{|\Sigma| \times d}$ , we define it as  $\text{OUTPUT}_{\theta_{\text{OUTPUT}}}(h) := \text{softmax}(\theta_{\text{OUTPUT}} h)$ .

All these parts come together for the distribution generator process in the algorithm 2.

The OUTPUT layer and the ATTN layer use a softmax function  $\mathbb{R}^n$  to  $\mathbb{R}^n$  defined as  $\text{softmax}(x)_i = \exp(x_i) / \sum_{j=1}^n \exp(x_j)$ .

The second component of a transformer is the **token selector**. This algorithm consumes the probability distribution computed by the distribution generator and returns a token.

In this work we study two selectors, we call them **deterministic selector** and **probabilistic selector**. Let  $p \in \Delta(\Sigma)$  be the distribution over  $\Sigma$  outputted by the distribution

**Algorithm 2** Distribution generator

**Input:** Transformer parameter  $\theta = (\theta_{\text{TE}}, \theta_{\text{PE}}, \{\theta_{\text{ATTN}}^l, \theta_{\text{FF}}^l\}_{l=0}^{L-1}, \theta_{\text{OUTPUT}})$  and tokens of the tape  $\alpha \in \Sigma^*$ .

**Output:** distribution  $p_\theta(\cdot|\alpha)$ .

- 1:  $h_i^{(0)} \leftarrow \theta_{\text{TE}}(\alpha_i) + \theta_{\text{PE}}(i), \forall 1 \leq i \leq |\alpha|$
- 2: **for**  $l = 0, \dots, L - 1$  **do**
- 3:    $(h_1^{(l+0.5)}, \dots, h_{|\alpha|}^{(l+0.5)}) \leftarrow (h_1^{(l)}, \dots, h_{|\alpha|}^{(l)}) + \text{ATTN}_{\theta_{\text{ATTN}}^{(l)}}(h_1^{(l)}, \dots, h_{|\alpha|}^{(l)})$
- 4:    $h_i^{(l+1)} \leftarrow h_i^{(l+0.5)} + \text{FF}_{\theta_{\text{FF}}^{(l)}}(h_i^{(l+0.5)}), \forall 1 \leq i \leq |\alpha|$
- 5: **end for**
- 6:  $p_\theta(\cdot|\alpha) \leftarrow \text{OUTPUT}_{\theta_{\text{OUTPUT}}}(h_{|\alpha|}^{(L)})$

generator. Then, the deterministic selector always returns the token with the largest probability (argmax over  $p$ ). In transformers with this selector, we do not allow for two or more tokens to be tied with the maximum probability, i.e.,  $\nexists \sigma$  s.t.  $p(\text{argmax}(p)) = p(\sigma) \wedge \sigma \neq \text{argmax}(p)$ . In contrast, the probabilistic selector works non-deterministically. It returns a token  $\sigma \in \Sigma$  with probability  $p(\sigma)$ . We allow for more than one maximum with this selector.

**RR** Add an image representing the transformer architecture

**Definition 2.1.1** (Deterministic transformer). We call a transformer deterministic when the token selector is the deterministic selector. We say that a family of deterministic transformers  $\{\text{TF}_n\}_{n \in \mathbb{N}}$  recognizes a language  $\mathcal{L}$  when:

$$\alpha \in \mathcal{L} \iff \text{TF}_{|\alpha|}^*(\alpha) = q_{\text{accept}}$$

As mentioned in Chapter 1, there is an unexplored area in the study of transformers expressiveness; a study of non-determinism is missing.

When  $\text{TF}$  is a deterministic transformer,  $\text{TF}^i(\alpha)$  is an exact symbol for all  $\alpha \in \Sigma^*$  and  $i \in \mathbb{N}$ . On the contrary, when the token selection process of the transformer samples the distribution, it becomes a random variable. Therefore, we introduce the following definition.

**Definition 2.1.2** (Probabilistic transformer). We call probabilistic transformer to the transformers that use the probabilistic selector. We say that a family of probabilistic transformers  $\{\text{TF}_n\}_{n \in \mathbb{N}}$  recognizes a language  $\mathcal{L}$  when:

$$\alpha \in \mathcal{L} \iff \Pr[\text{TF}_{|\alpha|}^*(\alpha) = q_{\text{accept}}] > 2/3$$

## 2.2 Finite precision modeling

I'm missing good symbols for  $\infty$  y  $-\infty$ .

In real models, the transformers typically work with 16- or 32-bit floating point numbers, and there is active work in reducing their size. For this reason, in this thesis we follow the decision of [?] to work with constant-precision transformers, where the output of each arithmetic operation is rounded to the closest floating point number representable

by a fixed number of digits following IEEE 754 standard. This analysis avoids the unrealistic infinite precision assumption made in prior works [?]. The definitions stated in this section are taken from [?].

Intuitively, we represent the numbers as a tuple  $\langle \text{sign}, \text{significand}, \text{exponent} \rangle$  with a fixed number of bits for each coordinate. The tuple is interpreted as  $\text{sign} \cdot \text{significand} \cdot 2^{\text{exponent}}$ . For example, the number 8 is represented as  $\langle +, 1, 3 \rangle$ . When fixing the number of bits for the significand and the exponent there are numbers non-representable. For example, with only 2 bits of significand and 3 of exponent, the number 9 can not be represented exactly. We can represent  $8 = 1 \cdot 2^3$  and  $12 = 3 \cdot 2^2$ , thus when an arithmetic operation should output 9 we output 8.

Below we give a formal definition of the floating-point number and rounding operation. We use  $\phi(a) = \sum_{i=1}^k 2^{k-i} a_i$  to denote the value of the binary number represented by  $a \in \{0, 1\}^k$  for any  $k \in \mathbb{N}^+$ .

**Definition 2.2.1** (Floating-point Representation). Let  $e$  be the number of bits for exponents and  $s$  be the number of bits for significand. A  $(e + 2s + 1)$ -bit binary string  $a = (a_1, a_2, \dots, a_{e+2s+1}) \in \{0, 1\}^{e+2s+1}$  is a floating-point binary representation of number  $\phi_{e,s}(a) \triangleq \text{sign}(a) \cdot 2^{\text{exponent}(a)} \cdot \text{significand}(a)$  with  $e$ -bit exponent and  $2s$ -precision, where the sign is  $\text{sign}(a) \triangleq 2a_1 - 1$ , the significand is  $\text{significand}(a) \triangleq 2^{-s} \phi(a_{2:2s+1})$ , and the exponent is  $\text{exponent}(a) \triangleq \phi(a_{2s+2:2s+e+1}) - 2^{\max(0, e-1)}$ . We define the set of numbers expressible with  $e$  bits of exponent and  $s$  bits of significand as  $\mathbb{F}_{e,s}$ . We define  $\infty_{e,s}$  as the maximum number in  $\mathbb{F}_{e,s}$ . Similarly, we define  $-\infty_{e,s}$  as the minimum. Since in this work we only analyze constant precision transformers, we will omit the subscripts of  $\infty$  and  $-\infty$ .

**Definition 2.2.2** (Correct Rounding). For any  $x \in \mathbb{R}$ , we define correct rounding  $\text{round}(x, \mathbb{F}_{e,s})$  as the closest number to  $x$  in  $\mathbb{F}_{e,s}$ . We break the tie by picking the one with a smaller absolute value. In particular, we denote the rounding operation with  $e$ -bit exponent,  $2s$ -bit precision by  $\text{round}_{e,s}(\cdot) \triangleq \text{round}(\cdot, \mathbb{F}_{e,s})$ , which is also denoted by  $[\cdot]_{e,s}$  for convenience. We extend the definition of  $\text{round}$  and  $\text{round}_{e,s}$  to vector inputs by rounding coordinate-wisely.

Our notion of floating-point number simplifies the IEEE 754 Standard for Floating-point Arithmetic [?] by removing  $\infty$  and  $-\infty$ . When overflow happens, we always round the output to  $\infty$ , and when underflow happens to  $-\infty$ . For unary functions like  $\exp(\cdot)$  and binary functions including addition, subtraction, multiplication, and division, we simply define their rounded version by rounding their outputs. Whenever division by 0 happens, we define the output as 0.<sup>1</sup>

Next, we define finite-precision summation over more than two numbers by decomposing it as a chain of rounded binary addition in a fixed order.

**Definition 2.2.3** (Summation with Iterative Rounding). For any  $s, n \in \mathbb{N}^+$  and vector  $x \in \mathbb{R}^n$ , we define summation with iterative rounding to  $e$  bit exponent and  $2s$ -bit precision as  $\text{sum}_{e,s} : \cup_{n \in \mathbb{N}^+} (\mathbb{F}_{e,s})^n \rightarrow \mathbb{F}_{e,s}$ , where for any  $n \in \mathbb{N}^+$  and  $x \in \mathbb{R}^n$ ,

$$\text{sum}_{e,s}(x) \triangleq \left[ \left[ \left[ [x_1 + x_2]_{e,s} + x_3 \right]_{e,s} + \dots + x_{n-1} \right]_{e,s} + x_n \right]_{e,s}.$$

<sup>1</sup> Notice that the only division is in the softmax. If it's the one in the attention mechanism, it's reasonable to give an attention score of 0 to every vector. If it's the one in the output layer it means that all symbols have an extremely small probability, therefore is reasonable to assign the same probability to all of them.

We further define the following operations:

- Finite-precision inner product:  $\langle x, y \rangle_{e,s} \triangleq \text{sum}_{e,s}(x \odot y)$ ;
- Finite-precision matrix product:  $(A \times_{e,s} B)_{i,j} \triangleq \langle (A_{i,:})^\top, B_{:,j} \rangle_{e,s}$  ;
- Finite-precision softmax:  $\text{softmax}_{e,s}(x) \triangleq \left[ \frac{[\exp(x)]_{e,s}}{\text{sum}_{e,s}([\exp(x)]_{e,s})} \right]_{e,s}$ .

From the saturation and the iterative rounding presented by [?], some properties appear that are important to remark:

- $\exp(-\infty) = 0$
- $\exp(\infty) = \infty$
- $-\infty \cdot \infty = -\infty$
- $\infty \cdot \infty = \infty$
- $\forall x \in \mathbb{F}_{e,s}$  it holds that  $x + 0 = x$
- $\forall x \in \mathbb{F}_{e,s}$  it holds that  $x - x = 0$
- $\forall x \in \mathbb{F}_{e,s}$  it holds that  $x + \infty + \infty + -\infty = 0$

### 2.3 Transformer Complexity classes

Maybe rethink this title?

As discussed in the first chapter, there are different resources to bound in the transformers to limit their expressive power. In this thesis, we analyze the effect of bounding the number of steps in the chain of thought that the transformers are allowed to make before answering. With this in mind, we define the following complexity classes.

Given two functions  $T(n)$  and  $d(n)$  we define the complexity class  $\text{CoT}[T(n), d(n)]$ . Informally, a language  $\mathcal{L}$  is in  $\text{CoT}[T(n), d(n)]$  if and only if there is a deterministic family of transformers with constant depth, embedding size  $d(n)$  and budget  $T(n)$  that recognizes it.

**Definition 2.3.1** ( $\text{CoT}[T(n), d(n)]$ ). A language  $\mathcal{L}$  is in  $\text{CoT}[T(n), d(n)]$  if and only if there is an integer  $L$  and two functions  $T'(n) = O(T(n))$ ,  $d'(n) = O(d(n))$ , such that for every positive integer  $n$ , there is an  $L$ -layer deterministic transformer, denoted by  $\text{TF}_n$  with embedding size  $d'(n)$ , that outputs  $\mathcal{L}(\alpha)$  given any input  $\alpha$  in  $\Sigma^n$ , using  $T'(n)$  steps of chain of thought.

**SF** Explicar qué es  $\mathcal{L}(\alpha)$ .

**RR** No es notación estándar?

In the same way, we define the classes of languages recognized by probabilistic transformers. Given two functions  $T(n)$  and  $d(n)$  we define the complexity class  $\text{RCoT}[T(n), d(n)]$ . Informally, a language  $\mathcal{L}$  is in  $\text{RCoT}[T(n), d(n)]$  if and only if there is a probabilistic family of transformers with constant depth, embedding size  $d(n)$  and budget  $T(n)$  that recognizes it.

---

**Definition 2.3.2** ( $\text{RCoT}[T(n), d(n)]$ ). A language  $\mathcal{L}$  is in  $\text{RCoT}[T(n), d(n)]$  if and only if there is an integer  $L$  and two functions  $T'(n) = O(T(n))$ ,  $d'(n) = O(d(n))$ , such that for every positive integer  $n$ , there is an  $L$ -layer probabilistic transformer, denoted by  $\text{TF}_n$  with embedding size  $d'(n)$ , that outputs  $\mathcal{L}(\alpha)$  given any input  $\alpha$  in  $\Sigma^n$ , using  $T'(n)$  steps in the chain of thought with probability higher than  $2/3$ .

## 2.4 Circuit Complexity

## 2.5 Known results

### 3. GENERAL PROPERTIES

In this chapter we study the properties of  $\text{CoT}[T(n), d(n)]$  classes and the primitive operations that transformers can perform. We define the notion of normalized transformer and prove that  $\text{CoT}[T(n), d(n)]$  classes are closed under boolean operations.

#### 3.1 Transformers Primitives

In this section we introduce a set of primitive operations that are used through the thesis. The operations are applied as modifications to an already existing transformer. If the operations are applied to an entire family of transformers  $\{\text{TF}_n\}_{n \in \mathbb{N}}$  such that the language it computes is in  $\text{CoT}[T(n), d(n)]$ , then the language of the modified family will be in the same class. This is because, although the operations add layers and coordinates to the size of the embedding, this is only a constant factor.

Most operations don't make sense on their own, but they are used in later demonstrations.

##### 3.1.1 Flagging Tokens and Positions

For complex operations inside the loop of the transformer we will need to identify when a vector comes from a certain token or position. In this section we present how to mark the vectors with flags during the embedding.

For flagging a token we use the token embedding, and in the same way, for flagging a position we use the position encoding. Either way, the strategy is exactly the same. The embedding dimension is extended to allocate the flags. A new coordinate is added per flag. Therefore, using a constant number of flags only adds a constant to the embedding dimension. Thus, if the flagging is made in a family of transformers its  $\text{CoT}[T(n), d(n)]$  class does not change.

For example, let  $v_\sigma$  and  $v_\tau$  be vectors of a transformer  $\text{TF}$  where  $v_\sigma$  comes from embedding the token  $\sigma$  and  $v_\tau$  from the token  $\tau$  (with  $\sigma \neq \tau$ ). We can define a transformer  $\text{TF}'$  that behaves like  $\text{TF}$  but recognize when a vector comes from embedding the token  $\sigma$ .  $\text{TF}'$  has embedding dimension  $d' = d + 1$  where  $d$  is the embedding dimension of  $\text{TF}$ .

$$v_\sigma = \begin{pmatrix} x \\ y \\ z \end{pmatrix} \rightarrow v'_\sigma = \begin{pmatrix} x \\ y \\ z \\ f_{\text{on}} \end{pmatrix} \qquad v_\tau = \begin{pmatrix} i \\ j \\ k \end{pmatrix} \rightarrow v'_\tau = \begin{pmatrix} i \\ j \\ k \\ f_{\text{off}} \end{pmatrix}$$

Fig. 3.1: Change in a flagged and unflagged vector

Formally, for flagging the token  $\sigma$  we define the token embedding function of  $\text{TF}'$  :  $\Sigma \rightarrow \mathbb{R}^{d+1}$  as

$$(\theta'_{TE}(\tau))_j = \begin{cases} (\theta_{TE}(\tau))_j & \text{if } j \neq d+1 \\ f_{\text{on}} & \text{if } j = d+1 \text{ and } \tau = \sigma \\ f_{\text{off}} & \text{if } j = d+1 \text{ and } \tau \neq \sigma \end{cases}$$

where  $f_{\text{on}}$  is the value that marks the vector as flagged and  $f_{\text{off}}$  as unflagged. Those values can be defined chosen arbitrarily for each flag. In this case, the coordinate inserted for the flag is the last one, but it can be added anywhere.

The other components of  $\text{TF}'$  ignore  $d+1$  the coordinate. The position encoding of  $\text{TF}'$  is equal to the position encoding of  $\text{TF}$ , except it puts a 0 in the  $d+1$  coordinate (not present in  $\text{TF}$ ) independently of the position. The matrices of the  $\text{ATTN}$ ,  $\text{FF}$  and  $\text{OUTPUT}$  layers of  $\text{TF}'$  are the ones of  $\text{TF}$  extended with 0 in the  $d+1$  rows and columns to match the change in embedding dimension.

$\text{TF}'$  behaves exactly as  $\text{TF}$  except it has one more coordinate in the embedding. Said coordinate represents the flag and takes only the values  $v_{\text{on}}$  or  $v_{\text{off}}$  in every vector.

This primitive works also as a mechanism to extend the embedding dimension without changing the behavior. If  $v_{\text{on}}$  and  $v_{\text{off}}$  are defined as 0, the resulting transformer is the same as the initial with an extra coordinate that is always 0.

### 3.1.2 Preserve and ignore coordinates through $\text{ATTN}$ and $\text{FF}$

When defining new transformers over existing ones we will need to extend the embedding dimension to store new information. In order to preserve the behavior of the original transformers across the attention and feed-forward layers, we need to construct a mechanism that ignores the extra coordinates. For example, this mechanism is used for preserving flags during  $\text{ATTN}$  and  $\text{FF}$  layers that don't use them.

Let  $\text{TF}$  be a transformer of embedding dimension  $d$  and assume we are constructing a transformer  $\text{TF}'$  of embedding dimension  $d' = d+1$ . Let  $W_Q, W_K, W_V, W_O$  be the matrices of an  $\text{ATTN}$  layer of  $\text{TF}$  and assume we want to perform the same layer in  $\text{TF}'$  but simulating that the  $i$ th coordinate does not exist. Notice that the embedding dimension of  $\text{TF}'$  has one more coordinate than  $\text{TF}$ 's. If we define the matrices of the layer  $W'_Q, W'_K, W'_V, W'_O$  as  $W_Q, W_K, W_V, W_O$  by inserting a row and column of zeros at the  $i$ th position, we obtain the desired  $\text{ATTN}$  layer.

Note that because the  $i$ th row of  $W'_O$  is zero, the output vectors of the  $\text{ATTN}$  layer are 0 in the  $i$ th coordinate. This is the point that makes it possible to not just ignore the value of the  $i$ th coordinate, but to not affect it.

For preserving the coordinate but not ignore it, it is enough to have zeros in only its corresponding row of  $W'_O$ . This construction works for more than one coordinate. To extend it to a set of coordinates, it suffices to set more rows and columns to 0.

The same applies in the feed forward layer, making the  $i$ th column and row of the matrices and vectors of the layer zero, makes it ignore the value of the  $i$ th coordinate. Furthermore, the result of the  $\text{FF}$  layer will have a zero in the  $i$ th coordinate, thus preserving its value.

**SC** diagrama?

**RR** how? pongo las matrices? Siento que a veces queda muy cargado. Quizás ahora se entiende mejor?

### 3.1.3 Add vector and constant functions

In this section we develop a mechanism for adding a fixed vector to a previously flagged one. This technique works even when there is more than one flagged vector.

Assume there is a flag with values  $f_{\text{on}} = 1$  and  $f_{\text{off}} = 0$ . For clarity, suppose the flag is in the last coordinate. Let  $v$  be the vector we want to add. We define a mechanism that replaces every  $h$  such that  $h_d = f_{\text{on}}$  with  $h + v$ .

The operation is performed with only one FF layer. We construct a layer such that  $FF(h) = \mathbb{I}(h_d = f_{\text{on}}) \cdot v$ . Said layer is achieved with the following parameters.

$$W_1 = id \quad W_2 = \begin{pmatrix} 0 & \dots & 0 & v_1 \\ \vdots & \ddots & \vdots & \vdots \\ 0 & \dots & 0 & v_d \end{pmatrix} \quad b_1 = b_2 = \begin{pmatrix} 0 \\ \vdots \\ 0 \end{pmatrix}$$

Notice that if the coordinate for the flag had been the  $k$ th instead of the last one, then  $v$  would have appeared in the  $k$ th column of  $W_2$  instead of the last.

This technique can be used to apply a constant function to the flagged vectors. This is achieved by exploiting the finite representation system. We can saturate the vector to the  $\infty$  and then make it 0 to finally add our objective. This holds since for all  $x, c \in \mathbb{F}_{e,s}$  it happens that  $x + \infty + \infty + -\infty + c = c$ . Because as mentioned in Section 2.2,  $x + \infty + \infty = \infty$  and  $\infty + -\infty = 0$ .

The constant function  $f(h) = v$  is implemented by adding four vectors that perform this trick. First, it is necessary to add the vector consisting entirely of  $\infty$  twice; then, a vector consisting entirely of  $-\infty$ ; and finally,  $v$ ; as shown in equation 3.1.

$$h \leftarrow h + \begin{pmatrix} \infty \\ \vdots \\ \infty \end{pmatrix} + \begin{pmatrix} \infty \\ \vdots \\ \infty \end{pmatrix} + \begin{pmatrix} -\infty \\ \vdots \\ -\infty \end{pmatrix} + v \quad (3.1)$$

### 3.1.4 Make a vector ignore others in an ATTN layer

We show how to make a vector ignore others in the ATTN layers, a technique that will prove useful in later constructions. A vector  $h_i$  ignores vector  $h_j$  when the attention score is 0, which can be achieved by making  $\langle q_i, k_j \rangle = -\infty$ .

We extend the embedding dimension with three more coordinates for preserving the attention score of the other vectors previous to taking softmax. We construct  $\text{TF}'$  over  $\text{TF}$  such that the  $l$ th vector ignores a set of vectors  $M$  in an ATTN layer; specifically, we require said layer to satisfy:

$$\langle q'_i, k'_j \rangle = \begin{cases} -\infty & \text{if } i = l \wedge j \in M \\ \langle q_i, k_j \rangle & \text{if } i \neq l \vee j \notin M \end{cases}$$

This can be achieved by saturating the finite representation since  $\langle q'_l, k'_i \rangle = \sum_{x=1}^{d+3} (q'_l)_x (k'_i)_x$ . If the last two terms of the summation are  $-\infty$ , then because of the iterative rounding,  $\langle q'_l, k'_i \rangle = -\infty$ . Since we do not want to change the attention scores of vectors not in  $M$ , when  $i \notin M$  the last three terms should be 0.



We force  $(k'_i)_{d+2}$  and  $(k'_i)_{d+3}$  to be  $-\infty$  when  $i \in M$ , and  $(q'_i)_{d+2}$  and  $(q'_i)_{d+3}$  to be 1 when  $i = l$  and 0 otherwise. We also make  $(k'_i)_{d+1}$  and  $(q'_i)_{d+1}$  to be 0 for all  $i$  so the third to last term of the inner product is always 0. With these values, the only attention scores modified are the ones of the  $l$ th vector since every other will have 0 in the last three coordinates of the query thus making the last three terms of the inner product 0 and therefore not affecting the score.

This can be achieved with the following flags:

$$(h_i)_{d+1} = \begin{cases} 1 & \text{if } i = l \\ 0 & \text{otherwise} \end{cases} \quad (h_i)_{d+2} = (h_i)_{d+3} = \begin{cases} 1 & \text{if } i \in M \\ 0 & \text{otherwise} \end{cases}$$

Now the query and key matrices are extended as:

$$W'_Q = \begin{pmatrix} W_Q & 0^{\mathbb{R}^{d \times 1}} & 0^{\mathbb{R}^{d \times 1}} & 0^{\mathbb{R}^{d \times 1}} \\ 0^{\mathbb{R}^{d \times 1}} & 0^{\mathbb{R}^{1 \times 1}} & 0^{\mathbb{R}^{1 \times 1}} & 0^{\mathbb{R}^{1 \times 1}} \\ 0^{\mathbb{R}^{d \times 1}} & 1^{\mathbb{R}^{1 \times 1}} & 0^{\mathbb{R}^{1 \times 1}} & 0^{\mathbb{R}^{1 \times 1}} \\ 0^{\mathbb{R}^{d \times 1}} & 1^{\mathbb{R}^{1 \times 1}} & 0^{\mathbb{R}^{1 \times 1}} & 0^{\mathbb{R}^{1 \times 1}} \end{pmatrix} \quad W'_K = \begin{pmatrix} W_K & 0^{\mathbb{R}^{d \times 1}} & 0^{\mathbb{R}^{d \times 1}} & 0^{\mathbb{R}^{d \times 1}} \\ 0^{\mathbb{R}^{d \times 1}} & 0^{\mathbb{R}^{1 \times 1}} & 0^{\mathbb{R}^{1 \times 1}} & 0^{\mathbb{R}^{1 \times 1}} \\ 0^{\mathbb{R}^{d \times 1}} & 0^{\mathbb{R}^{1 \times 1}} & -\infty^{\mathbb{R}^{1 \times 1}} & 0^{\mathbb{R}^{1 \times 1}} \\ 0^{\mathbb{R}^{d \times 1}} & 0^{\mathbb{R}^{1 \times 1}} & 0^{\mathbb{R}^{1 \times 1}} & -\infty^{\mathbb{R}^{1 \times 1}} \end{pmatrix}$$

Then,

$$\begin{aligned} \langle q'_i, k'_j \rangle &= \sum_{x=1}^{d+3} (q'_i)_x (k'_j)_x = \langle q_i, k_j \rangle + (q'_l)_{d+1} (k'_i)_{d+1} + (q'_l)_{d+2} (k'_i)_{d+2} + (q'_l)_{d+3} (k'_i)_{d+3} \\ &= \langle q_i, k_j \rangle + 0 + (q'_l)_{d+2} (k'_i)_{d+2} + (q'_l)_{d+3} (k'_i)_{d+3} \\ &= \langle q_i, k_j \rangle + (q'_l)_{d+2} (k'_i)_{d+2} + (q'_l)_{d+3} (k'_i)_{d+3} \\ &= \begin{cases} \langle q_i, k_j \rangle + 1 \cdot -\infty + 1 \cdot -\infty & \text{if } i = l \wedge j \in M \\ \langle q_i, k_j \rangle + 1 \cdot 0 + 1 \cdot 0 & \text{if } i = l \wedge j \notin M \\ \langle q_i, k_j \rangle + 0 \cdot -\infty + 0 \cdot -\infty & \text{if } i \neq l \wedge j \in M \\ \langle q_i, k_j \rangle + 0 \cdot 0 + 0 \cdot 0 & \text{if } i \neq l \wedge j \notin M \end{cases} \\ &= \begin{cases} \langle q_i, k_j \rangle + -\infty + -\infty & \text{if } i = l \wedge j \in M \\ \langle q_i, k_j \rangle + 0 + 0 & \text{if } i = l \wedge j \notin M \\ \langle q_i, k_j \rangle + 0 + 0 & \text{if } i \neq l \wedge j \in M \\ \langle q_i, k_j \rangle + 0 + 0 & \text{if } i \neq l \wedge j \notin M \end{cases} \\ &= \begin{cases} -\infty & \text{if } i = l \wedge j \in M \\ \langle q_i, k_j \rangle & \text{otherwise} \end{cases} \end{aligned}$$

Namely,  $\langle q'_i, k'_j \rangle = -\infty$  if  $i = l$  and  $j \in M$ , but  $\langle q'_i, k'_j \rangle = \langle q_i, k_j \rangle$  otherwise.

### 3.1.5 Linear transformations

Although it may seem that transformers are a model that should be able to perform linear transformations in a native way, it is not the case. Applying a linear transformation to a vector is not straightforward, since in every step of ATTN and FF the result of the layer is added to the original vector instead of replacing it. An intuitive idea would be to apply the transformation  $M - I$  when we want to transform the vector  $x$  to  $Mx$ , since  $x + (M - I)x = x + Mx - Ix = Mx$ . This does not hold in our model since the addition

and multiplication are not commutative in  $\mathbb{F}_{e,s}$ . For avoiding the representation error, we will extend the embedding dimension to get more memory per vector such that we can store in new coordinates the value of  $Mx$  and then replace it in the original ones.

Let  $f$  denote the index of the flagged vector. We show a mechanism to replace it with the result of applying the linear transformation  $M$  to it. We perform the operation with two consecutive **ATTN** layers, in which  $h_f$  only attends to itself. Therefore, after these two layers  $h_f$  will be replaced with  $Mh_f$ .

We need to double the embedding size from  $d$  to  $2d$  (plus some flags) since we need to store the result of the transformation in some coordinates (the added  $d + 1$  to  $2d$ ) as a middle step. In the first **ATTN** layer, we compute the transformation saving its result in the coordinates  $d + 1$  to  $2d$ . Then, in the second layer we move the result to the right place.

$$\widetilde{h}_f = \begin{pmatrix} (h_f)_1 \\ \vdots \\ (h_f)_d \\ 0 \\ \vdots \\ 0 \end{pmatrix} \xrightarrow{\text{first ATTN layer}} \widetilde{h}_f = \begin{pmatrix} 0 \\ \vdots \\ 0 \\ (Mh_f)_1 \\ \vdots \\ (Mh_f)_d \end{pmatrix} \xrightarrow{\text{second ATTN layer}} \widetilde{h}_f = \begin{pmatrix} (Mh_f)_1 \\ \vdots \\ (Mh_f)_d \\ 0 \\ \vdots \\ 0 \end{pmatrix}$$

We assume that the  $d + 1$  to  $2d$  coordinates (used as memory) are set to 0 by the layers preceding this mechanism. This can be achieved maintaining them as 0 since the beginning or transformed to 0 with the primitive described in Section 3.1.3.

Let us construct each of the **ATTN** layers. As said before, in both layers we will need for  $\widetilde{h}_f$  to only attend to itself. Therefore, we define  $W_Q$  and  $W_K$  using the strategy of Section 3.1.4.

Recall from the definition of the **ATTN** mechanism (1) that the output of the **ATTN** layer is  $h'_i = W_O \sum_{j=1}^a (s_i)_j v_j$ . By this point,  $h'_f$  is looking like

$$h'_f = W_O \sum_{j=1}^a (s_f)_j v_j = W_O v_f = W_O (W_V \widetilde{h}_f)$$

- For the first layer, defining  $W_O$  and  $W_V$  as follows

$$W_V = id \quad W_O = \begin{pmatrix} -id^{\mathbb{R}^{d \times d}} & 0^{\mathbb{R}^{d \times d}} \\ M & 0^{\mathbb{R}^{d \times d}} \end{pmatrix}$$

makes us get  $h'_f = (-(h_f)_1, \dots, -(h_f)_d, (Mh_f)_1, \dots, (Mh_f)_d)$ , which added to  $\widetilde{h}_f$  gives us what we were looking for after the first layer.

- For the second **ATTN** layer, we use the same mechanism for making  $\widetilde{h}_f$  only attend to itself. And defining  $W_V$  and  $W_O$  as follows

$$W_V = id \quad W_O = \begin{pmatrix} 0^{\mathbb{R}^{d \times d}} & id^{\mathbb{R}^{d \times d}} \\ 0^{\mathbb{R}^{d \times d}} & -id^{\mathbb{R}^{d \times d}} \end{pmatrix}$$

we get  $h'_f = ((Mh_f)_1, \dots, (Mh_f)_d, -(Mh_f)_1, \dots, -(Mh_f)_d)$ , which added to  $\widetilde{h}_f$  gives us what we were looking for after the second layer.

Note that, after these two layers, the values of  $h_i$  for all  $i \neq f$  have changed in ways that have not been studied in this paper, but this is not relevant to the way we use the primitive.

### 3.1.6 Max coordinate

We show how to take the maximum of a set of coordinates. This operation will be needed for later constructions.

We give a construction for taking the maximum of two coordinates. This construction can be easily extended to more coordinates. The operation works properly when the set of coordinates has only one maximum. All the coordinates that are not the maximum are replaced by 0.

$$\begin{pmatrix} a \\ b \\ \vdots \end{pmatrix} \rightarrow \begin{cases} \begin{pmatrix} a \\ 0 \\ \vdots \end{pmatrix} & \text{if } a > b \\ \begin{pmatrix} 0 \\ b \\ \vdots \end{pmatrix} & \text{if } b > a \end{cases}$$

The maximum of two coordinates can be computed with the following primitives:

$$\begin{aligned} \begin{pmatrix} a \\ b \\ 0 \\ 0 \\ \vdots \end{pmatrix} &\xrightarrow{\text{FF layer}} \begin{pmatrix} a \\ b \\ \text{relu}(a-b) \\ \text{relu}(b-a) \\ \vdots \end{pmatrix} \xrightarrow[\text{the coordinates with the relu}]{\text{Multiply twice by } \infty} \begin{pmatrix} a \\ b \\ \infty \cdot \mathbb{I}(a > b) \\ \infty \cdot \mathbb{I}(b > a) \\ \vdots \end{pmatrix} \rightarrow \\ &\xrightarrow[\text{the coordinates with the comparison}]{\text{Divide by } \infty} \begin{pmatrix} a \\ b \\ \mathbb{I}(a > b) \\ \mathbb{I}(b > a) \\ \vdots \end{pmatrix} \xrightarrow[\text{operation}]{\text{Special}} \begin{pmatrix} a \cdot \mathbb{I}(a > b) \\ b \cdot \mathbb{I}(b > a) \\ 0 \\ 0 \\ \vdots \end{pmatrix} = \begin{cases} \begin{pmatrix} a \\ 0 \\ \vdots \end{pmatrix} & \text{if } a > b \\ \begin{pmatrix} 0 \\ b \\ \vdots \end{pmatrix} & \text{if } b > a \end{cases} \end{aligned}$$

For computing the maximum of a set of coordinates it is only needed to compute a comparison for each pair of the set and multiply each coordinate  $x$  by  $\mathbb{I}(x > y)$  for all  $y \neq x$  in the set.

#### Implementation of the special operation

This operation is not a linear transformation; therefore, we need to define a new technique. Observe that one of  $\mathbb{I}_a$  and  $\mathbb{I}_b$  equals 1 and the other equals 0. Then,  $\mathbb{I}_b = 1 - \mathbb{I}_a$  and

$$\mathbb{I}_a = 1 - \mathbb{I}_b.$$

$$\begin{pmatrix} a \\ b \\ \mathbb{I}_a \\ \mathbb{I}_b \end{pmatrix} \xrightarrow[\text{operation}]{\text{Special}} \begin{pmatrix} a \cdot \mathbb{I}_a \\ b \cdot \mathbb{I}_b \\ 0 \\ 0 \end{pmatrix}$$

This operation is performed with three layers of **FF**. Naturally, after each layer a new vector is added to the original. The vectors we add are the following:

$$\begin{pmatrix} a \\ b \\ \mathbb{I}_a \\ \mathbb{I}_b \end{pmatrix} \xrightarrow[\text{operation}]{\text{Special}} \begin{pmatrix} a + \infty \cdot \mathbb{I}_b + \infty \cdot \mathbb{I}_b + -\infty \cdot \mathbb{I}_b \\ b + \infty \cdot \mathbb{I}_a + \infty \cdot \mathbb{I}_a + -\infty \cdot \mathbb{I}_a \\ \mathbb{I}_a + 0 + 0 + -\mathbb{I}_a \\ \mathbb{I}_b + 0 + 0 + -\mathbb{I}_b \end{pmatrix}$$

Note that if  $\mathbb{I}_a = 1$ , then  $\mathbb{I}_b = 0$ , so the first coordinate isn't modified and the second is changed to 0 because of the finite representation properties. In the same way, if  $\mathbb{I}_a = 0$ , then  $\mathbb{I}_b = 1$ , so the second coordinate isn't modified and the first is changed to 0.

Let us show how to instantiate the **FF** layers. The first two have parameters:

$$W_1 = id \quad W_2 = \begin{pmatrix} 0 & 0 & 0 & \infty \\ 0 & 0 & \infty & 0 \\ 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 \end{pmatrix} \quad b_1 = b_2 = \begin{pmatrix} 0 \\ 0 \\ 0 \\ 0 \end{pmatrix}$$

Notice that the output of these **FF** layer when the input is  $(x, y, \mathbb{I}_a, \mathbb{I}_b)$  will be  $(\infty \cdot \mathbb{I}_b, \infty \cdot \mathbb{I}_a, 0, 0)$ .

The third **FF** layer has parameters:

$$W_1 = id \quad W_2 = \begin{pmatrix} 0 & 0 & 0 & -\infty \\ 0 & 0 & -\infty & 0 \\ 0 & 0 & -1 & 0 \\ 0 & 0 & 0 & -1 \end{pmatrix} \quad b_1 = b_2 = \begin{pmatrix} 0 \\ 0 \\ 0 \\ 0 \end{pmatrix}$$

which, when the input is  $(x, y, \mathbb{I}_a, \mathbb{I}_b)$ , outputs  $(-\infty \cdot \mathbb{I}_b, -\infty \cdot \mathbb{I}_a, -\mathbb{I}_a, -\mathbb{I}_b)$

**RR** intenté escribir cómo quedaban las **FF** instanciadas pero queda medio ilegible. No sé si vale la pena

### 3.1.7 Repeat a token once it's printed

We show how to make a transformer repeat a token until we decide it to make it stop. This construction is needed to argue that for every transformer **TF** there is a transformer **TF'** of the same  $\text{CoT}[T(n), d(n)]$  class that computes the same language and consumes all of its budget.

Assume that **TF** is normalized and let us define **TF'**. **TF'** is identical to **TF** except for the following modifications. Let  $\sigma$  be the token to repeat. We flag the vectors coming from  $\sigma$  with the token embedding. Since we are replicating a token that was printed, i.e., it is not part of the input, its corresponding vector is the last one. This is because we begin repeating it after the first time it is printed.

Since we flag the vectors coming from the token  $\sigma$  we can apply a constant function only to them. This function is applied in the last layer of the body of  $\text{TF}'$ . Since we assume  $\text{TF}$  was normalized, the **OUTPUT** matrix is just a projection. Then, the value of the constant function is what is outputted as the probability distribution (after softmax). If the constant function has value  $-\infty$  in every coordinate except in the one associated with the probability of  $\sigma$ , where it has  $\infty$ , the distribution outputted by the distribution generator will have all the probability concentrated in  $\sigma$ . Thus,  $\text{TF}'$  prints  $\sigma$ .

Observe that this only works if the token does not appear in the input word. This case can be handled by exploiting the non-uniformity of the model. If the input word is of size  $n$ , with the PE we can flag the first  $n$  vectors. Then, we can use that flag to set to 0 the coordinate used by the TE for flagging the token to be repeated, i.e., turning of the flag. We make them 0 using a constant function as described in the Section 3.1.3.

### 3.1.8 Force to halt after certain number of steps of CoT

Using a strategy similar to the one in Section 3.1.7, we use the PE to flag when a certain position is reached. Then, if the transformer is normalized, modifying the last vector we decide the outputted distribution. Particularly, we can force it to concentrate the probability in  $q_{\text{accept}}$  or  $q_{\text{reject}}$ . We change the last vector with a constant function. The flag inserted by the PE is to decide if the function has to be applied or not.

Notice that this primitive, together with the one described in Section 3.1.7, proves that for every transformer  $\text{TF}$  in  $\text{CoT}[T(n), d(n)]$  there is a transformer  $\text{TF}'$  in the same class that always consumes all of its budget.

We extend the alphabet with two dummy versions of the decision tokens and construct  $\text{TF}'$  over  $\text{TF}$ . We make  $\text{TF}'$  work as  $\text{TF}$ , but we make it print the dummy versions of the decision tokens when  $\text{TF}$  would have printed the real ones<sup>1</sup>. Then, we force  $\text{TF}'$  to repeat the dummy decision tokens until it reaches the last position of the tape (uses all the budget). At that point, we force  $\text{TF}'$  to print the real decision token corresponding to the one it was repeating.

### 3.1.9 Copy one vector into another with ATTN

For the disjunction of transformers we will need to move values from the  $i$ th vector to the  $j$ th vector. In this section we show how to do it. This operation can only be done if  $i < j$ , since the attention mechanism only allows for a vector to attend to vectors to its left, i.e.,  $h_k$  can attend to  $h_l$  if and only if  $k \leq l$ .

Let  $j$  be the number of the receiving vector and  $i$  the number of the copied, i.e., we copy from  $h_i$  to  $h_j$ . The operation is done in two steps.

First, we make 0 the coordinates of  $h_j$  that will receive the values. This is done with a constant function.

Second, we make  $h_j$  only attend to  $h_i$  in an attention layer. In said attention layer, we use  $W_V = id$  and  $W_O$  such that it outputs the desired coordinates of  $h_i$  in the target coordinates. Formally,  $W_O$  has only 0 in the row  $k$  if  $h_j$  is not receiving anything from  $h_i$  in the  $k$ th coordinate. If  $h_j$  is receiving the coordinate  $l$  of  $h_i$  in the position  $k$ , then the  $k$ th row of  $W_O$  has 0 in every position except the  $l$ th.

<sup>1</sup> This is achieved giving the probability of each decision token to its dummy version.

For example, if we want to copy the 2nd coordinate of  $h_i$  to the 3rd coordinate of  $h_j$ ,  $W_O$  would look like 3.2. Then, the vector coming out of the ATTN layer will have 0 in every coordinate, except in the 3rd, where it will be  $(h_i)_2$ .

$$W_O = \begin{pmatrix} 0 & 0 & 0 \\ 0 & 0 & 0 \\ 0 & 1 & 0 \end{pmatrix} \quad (3.2)$$

### 3.2 Normalization

For an easier handling of transformers, we define a notion of normalized transformer. This notion will specify properties of the OUTPUT matrix and of the vector consumed by this step.

**Definition 3.2.1** (Normalized deterministic transformer). We say a deterministic transformer is normalized if it satisfies two conditions.

- **N1:** The matrix of the OUTPUT has to be a projection, i.e., it is a diagonal matrix with ones on its diagonal.
- **N2:** The last vector after the last iteration has to only be composed of  $-\infty$  in every projected coordinate except one, which has to be  $\infty$ . The non projected coordinates are free to be any value.

The motivation for this definition is to simplify the OUTPUT layer (N1) and to have more control over the probability distribution generated (N2). The exact values of the last vectors are such that after projecting and applying softmax in the OUTPUT layer we get a deterministic distribution with 100% of the probability concentrated in only one token.

For every family of transformers  $\text{TF}_{nn \in \mathbb{N}} \in \text{CoT}[T(n), d(n)]$  computing the language  $\mathcal{L}$ , there is a family  $\text{TF}'_{nn \in \mathbb{N}}$  in the same class, computing the same language, where for all  $n$ ,  $\text{TF}'_n$  is normalized. In the next section we introduce an algorithm for normalizing transformers.

#### 3.2.1 Algorithm for normalizing

For normalizing transformers, first we make them meet N1 and then N2.

##### Algorithm for N1

Notice that every linear transformation  $M : \mathbb{R}^d \rightarrow \mathbb{R}^{|\Sigma|}$  with  $d \geq |\Sigma|$  can be decomposed into two steps.<sup>2</sup>

$$\begin{pmatrix} x_1 \\ \vdots \\ x_d \end{pmatrix} \xrightarrow{M'} \begin{pmatrix} (Mx)_1 \\ \vdots \\ (Mx)_{|\Sigma|} \\ v \end{pmatrix} \xrightarrow{\text{projection}} \begin{pmatrix} (Mx)_1 \\ \vdots \\ (Mx)_{|\Sigma|} \end{pmatrix}$$

<sup>2</sup> We can always add more coordinates to the embedding if  $d < |\Sigma|$ . Since the language is fixed, it will be a constant number of coordinates, thus not changing the class of the transformer.

where  $v \in \mathbb{R}^{d-|\Sigma|}$  can have any values.

Let **TF** be a transformer and  $M$  be the matrix of the **OUTPUT** layer. Using the decomposition just presented, and the strategy defined in Section 3.1.5 for applying linear transformations, we can change the matrix of the **OUTPUT** layer to be only a projection. The linear transformation  $M$  has to be applied to the last vector, since it is the one going from the body of **TF** to the **OUTPUT** layer.

#### Algorithm for N2

Let **TF** be a transformer satisfying **N1**, we give an algorithm for making it also meet **N2**.

We use the maximum operation described in Section 3.1.6 to make all the positions of the last vector 0 except the one that was the maximum before. Then, if the maximum is positive, multiplying by  $\infty$ , then subtracting 1 from every coordinate and multiplying by  $\infty$  again makes **TF** meet **N2**.

$$\begin{pmatrix} 0 \\ \vdots \\ 0 \\ x_{max} \\ 0 \\ \vdots \\ 0 \end{pmatrix} \xrightarrow[\infty]{\text{multiplying by}} \begin{pmatrix} 0 \\ \vdots \\ 0 \\ \infty \\ 0 \\ \vdots \\ 0 \end{pmatrix} \xrightarrow{\text{subtracting 1}} \begin{pmatrix} -1 \\ \vdots \\ -1 \\ \infty - 1 \\ -1 \\ \vdots \\ -1 \end{pmatrix} \xrightarrow[\infty]{\text{multiplying by}} \begin{pmatrix} -\infty \\ \vdots \\ -\infty \\ \infty \\ -\infty \\ \vdots \\ -\infty \end{pmatrix}$$

It could be the case that the maximum was negative, in that case the resulting vector is not what we presented. We show a mechanism for making the maximum strictly positive so we can insert it in the algorithm before subtracting 1.

At this point our vector looks like  $(0, \dots, 0, x_{max}, 0, \dots, 0)$ , we are going to transform it to  $(0, \dots, 0, |x_{max}|, 0, \dots, 0)$ . Recall that the coordinates we are interested in are only the first  $|\Sigma|$ . Let us call  $v$  to the block containing these coordinates. Since  $\text{relu}(x) + \text{relu}(-x) = |x|$  for all  $x$ , our algorithm does the following:

$$\begin{pmatrix} v \\ \vdots \end{pmatrix} \rightarrow \begin{pmatrix} v \\ -v \\ \vdots \end{pmatrix} \rightarrow \begin{pmatrix} \text{relu}(v) \\ \text{relu}(-v) \\ \vdots \end{pmatrix} \rightarrow \begin{pmatrix} \text{relu}(v) + \text{relu}(-v) \\ \vdots \end{pmatrix}$$

where  $\text{relu}$  is applied position wise.

The first step is done in the same way as the first half of the linear transformations. For the second step, we take max-coordinate (Section 3.1.6) in  $v$  and  $-v$  independently. This works since all the coordinates of  $v$  are 0 except one thus there are only two scenarios. The non-zero coordinate of  $v$  is positive, then it will be the maximum (so it won't be modified) and all the others will remain 0; or the non-zero coordinate of  $v$  is negative, then the maximum will be 0 and all the coordinates will become 0. Whichever happens in  $v$ , the other happens in  $-v$ . Then, for adding both vectors we can use one **FF** layer with the following parameters. Observe that this works since all the coordinates of  $\text{relu}(-v)$  are non-negative.

$$W_1 = id^{d \times d} \quad W_2 = \begin{pmatrix} 0^{|\Sigma| \times |\Sigma|} & id^{|\Sigma| \times |\Sigma|} & 0^{|\Sigma| \times (d-2|\Sigma|)} \\ 0^{d-|\Sigma| \times |\Sigma|} & 0^{d-|\Sigma| \times |\Sigma|} & 0^{d-|\Sigma| \times d-2|\Sigma|} \end{pmatrix} \quad b_1 = b_2 = (0^{d \times d})$$

SC hacer step by step esta FF

### 3.3 Boolean Operations

In this section we prove that  $\text{CoT}[T(n), d(n)]$  classes are closed under boolean operations. We show that the complement of a language in  $\text{CoT}[T(n), d(n)]$  is in  $\text{CoT}[T(n), d(n)]$  and that the disjunction of two language of the same class, stays in the same class.

#### 3.3.1 Complement of a language

Let  $\mathcal{L} \in \text{CoT}[T(n), d(n)]$ . Then there is a family of transformers  $\{\text{TF}_n\}_{n \in \mathbb{N}}$  of embedding size  $d(n)$  that after  $T(n)$  steps of CoT prints  $q_{\text{accept}}$  or  $q_{\text{reject}}$  depending on if the input word is in the language. For each transformer in the family, we construct a corresponding transformer that behaves identically (internally and externally) but prints the flipped decision token.

RR dibujito

Let  $\text{TF}$  be a transformer, we construct  $\text{TF}'$  that flips the decision token of  $\text{TF}$ , i.e., if  $\text{TF}^*(\alpha) = q_{\text{accept}}$ , then  $\text{TF}'^*(\alpha) = q_{\text{reject}}$  and if  $\text{TF}^*(\alpha) = q_{\text{reject}}$ , then  $\text{TF}'^*(\alpha) = q_{\text{accept}}$ . This is the only change made to  $\text{TF}$ , thus  $\text{TF}'^i(\sigma) = \text{TF}^i(\sigma)$  for every  $i$  when  $\sigma$  is not a decision symbol.

This is done by swapping the rows associated with  $q_{\text{accept}}$  and  $q_{\text{reject}}$  in the matrix of the OUTPUT layer of  $\text{TF}$ . Thus, when the largest probability is associated with  $q_{\text{accept}}$  in  $\text{TF}$ , that same number will be associated to  $q_{\text{reject}}$  in  $\text{TF}'$ .

Formally,  $\text{TF}'$  is defined exactly the same as  $\text{TF}$  except in the OUTPUT layer, where the matrix is different. Let  $M$  be the matrix of the OUTPUT layer of  $\text{TF}$ , let  $i$  be the position of the token  $q_{\text{accept}}$  and  $j$  the position of  $q_{\text{reject}}$ , then the matrix of the OUTPUT layer of  $\text{TF}'$  will be  $M' = PM$  where  $P$  is the matrix that flips the rows  $i$  and  $j$ .

$$M'_{kl} = \begin{cases} 1 & \text{if } k = l, k \neq i \text{ and } k \neq j \\ 1 & \text{if } k = j \text{ and } l = i \\ 1 & \text{if } k = i \text{ and } l = j \\ 0 & \text{otherwise} \end{cases}$$

Since the OUTPUT matrix is the only difference between  $\text{TF}$  and  $\text{TF}'$ , they have the same budget and embedding size as  $\text{TF}$ , therefore,  $\text{TF}'$  belongs to the same class as  $\text{TF}$ .

We can define the family  $\{\text{TF}'_n\}_{n \in \mathbb{N}}$  where  $\text{TF}'_i$  is the negation of  $\text{TF}_i$  constructed as explained in this section. Then  $\{\text{TF}'_n\}_{n \in \mathbb{N}}$  computes the complement of  $\{\text{TF}_n\}_{n \in \mathbb{N}}$ . Since for every  $i \in \mathbb{N}$ ,  $\text{TF}'_i$  has the same budget and embedding size as  $\text{TF}_i$ ,  $\bar{\mathcal{L}}$  is in  $\text{CoT}[T(n), d(n)]$ .



### 3.3.2 Disjunction of two languages

Let  $\mathcal{L}_1$  and  $\mathcal{L}_2$  be languages in  $\text{CoT}[T(n), d(n)]$ , then there is a family of transformers  $\{(\text{TF}_1)_n\}_{n \in \mathbb{N}}$  (resp.  $\{(\text{TF}_2)_n\}_{n \in \mathbb{N}}$ ) of embedding size  $d_1(n)$  (resp.  $d_2(n) \in O(d(n))$ ) that after  $T_1(n)$  (resp.  $T_2(n) \in T(n)$ ) steps of CoT computes  $\mathcal{L}_1$  (resp.  $\mathcal{L}_2$ ).

We construct a family  $\{\text{TF}_n\}_{n \in \mathbb{N}}$  that computes  $\mathcal{L}_1 \cup \mathcal{L}_2$ . It does it in  $T_1(n) + T_2(n) + 1 \in O(T(n))$  steps of CoT and it has embedding size  $O(d(n))$ .

The idea behind the construction is to run first one of the transformers, then run the other and finally compute the disjunction of both results.

For every input  $\alpha$ , the tape of the transformer, at the end of the computation, ends up looking as  $\alpha\beta q$ , where  $q$  is the decision token printed by the transformer and  $\beta \in \Sigma^*$  is the sequence of intermediate tokens printed during the chain of thought. We denote by  $\alpha\beta q_1$  the content of the tape of  $\text{TF}_1$  at the end of the computation when the input word is  $\alpha$ . Similarly, we denote by  $\alpha\gamma q_2$  the content of the tape of  $\text{TF}_2$  at the end of the computation when the input word is  $\alpha$ . The tape of  $\text{TF}$ , computing the disjunction of  $\text{TF}_1$  and  $\text{TF}_2$ , at the end of the computation has content  $\alpha\beta q_1 \gamma q_2 q$ , where  $q$  is the decision token representing the disjunction of  $q_1$  and  $q_2$ .

Let us construct  $\text{TF}$  from  $\text{TF}_1$  and  $\text{TF}_2$ . Since we cannot change the body of the transformer during the iterations, we are required to use some tricks.

First, we grow the dimension of the embedding for running the body of both transformers at the same time over the same vectors but independently. In the final layers we choose which probability distribution will go to the output layer, this makes  $\text{TF}$  print the token generated by  $\text{TF}_1$  or  $\text{TF}_2$ . As always, we assume that the transformers are normalized.

We use the first half of the embedding for the transformer  $\text{TF}_1$  and the second for the transformer  $\text{TF}_2$ . We have some extra coordinates for flags at the end, but we will ignore them for now.

$$\text{TE}(\sigma)_i = \begin{cases} \text{TE}_1(\sigma)_i & \text{if } i \leq d_1(|\alpha|) \\ \text{TE}_2(\sigma)_{i-d_1(|\alpha|)} & \text{if } d_1(|\alpha|) < i \leq d_2(|\alpha|) \\ 0 & \text{if } d_2(|\alpha|) < i \end{cases}$$

$$\text{PE}(n)_i = \begin{cases} \text{PE}_1(n)_i & \text{if } i \leq d_1(|\alpha|) \\ \text{PE}_2(n)_{i-d_1(|\alpha|)} & \text{if } d_1(|\alpha|) < i \leq d_2(|\alpha|) \\ \text{flags} & \text{if } d_2(|\alpha|) < i \end{cases}$$

**RR** Me gustaría en vez de poner  $\text{PE}_2$  poner algo que se entienda que es más o menos  $\text{PE}_2$ . Formalmente  $\text{PE}_2$  está mal porque no está definida en el rango que se la usa. Debería ser  $\text{shiftedPE}_2$  que lo defino más adelante. Hacer esa explicación ahora me parece que mete mucho ruido.

The transformer  $\text{TF}$  has one layer for each layer of  $\text{TF}_1$  and each of  $\text{TF}_2$ . The first layers correspond to the body of  $\text{TF}_1$ . They are constructed ignoring the half of the embedding storing the coordinates used by  $\text{TF}_2$ . The second half of layers are the other way around, they compute the layers of  $\text{TF}_2$  and ignore the  $\text{TF}_1$  coordinates. Then, all the matrices in the layers of the first part of the body of  $\text{TF}$  look like  $V_1$  and the ones in the second layers will look like  $V_2$  where  $W_1$  and  $W_2$  are the matrices of the corresponding layers in  $\text{TF}_1$  and  $\text{TF}_2$ .

$$V_1 = \begin{pmatrix} W_1 & 0 \\ 0 & 0 \end{pmatrix} \quad V_2 = \begin{pmatrix} 0 & 0 \\ 0 & W_2 \end{pmatrix}$$

It is important to notice that the values of the coordinates corresponding to one transformer do not affect the values of the other. Thus, if the tape of the transformer has content  $\gamma$ , then at the end of the body of  $\text{TF}$ , the last vector would look like  $(x, y, \text{flags})$  where  $x$  is the last vector at the end of the body of  $\text{TF}_1$  when its tape content is  $\gamma$ , and  $y$  the one at the end of  $\text{TF}_2$  when its tape content is  $\gamma$ . Therefore, if we always choose to transform said vector into  $(x, 0, \dots, 0)$ , then  $\text{TF}$  will produce the same tape as  $\text{TF}_1$ . This works for generating the first  $T_1(|\alpha|)$  tokens which are  $\beta q_1$ . Thus, we would like to apply this transformation only to the vectors generating them. Therefore, we need to flag them with the position embedding. Since  $h_{|\alpha|+i}$  carries the distribution to output in the  $(i+1)$ -th step, i.e., the  $(|\alpha|+i+1)$ -th token of the tape, the vectors that we need to flag will be  $h_{|\alpha|}, h_{|\alpha|+1}, \dots, h_{|\alpha|+T_1(|\alpha|)-1}$ .

After running  $T_1(|\alpha|)$  steps of this construction, the tape will be  $\alpha\beta q_1$ . At this point, we want to start generating  $\gamma$ , i.e., the tape of  $\text{TF}_2$ . Just changing what we choose before the **OUTPUT** layer does not work, since what the construction will produce would be  $\text{TF}_2(\alpha\beta q_1)$ , and we need  $\text{TF}_2(\alpha)$ .

Now we show how to run  $\text{TF}_2$  in this construction ignoring the tokens produced by  $\text{TF}_1$ . We make  $\text{TF}_2$  not attend to the tokens in positions  $|\alpha|+1$  to  $T_1(|\alpha|)$  (this can be done using the primitive presented in the Section 3.1.4). We also need for them to be invisible to the portion of the position encoding corresponding to  $\text{TF}_2$ , thus we "shift" it in the following way.

$$\text{PE}(n)_i = \begin{cases} \text{PE}_1(n)_i & \text{if } i \leq d_1(|\alpha|) \\ \text{shiftedPE}_2(n)_{i-d_1(|\alpha|)} & \text{if } d_1(|\alpha|) < i \leq d_2(|\alpha|) \\ \text{flags} & \text{if } d_2(|\alpha|) < i \end{cases}$$

$$\text{shiftedPE}_2(n) = \begin{cases} \text{PE}_2(n) & \text{if } n \leq |\alpha| \\ 0 & \text{if } |\alpha| < n < T_1(|\alpha|) \\ \text{PE}_2(n - T_1(|\alpha|)) & \text{if } T_1(|\alpha|) < n \end{cases}$$

The first case is for the positions occupied by the input, we do not want to change their encoding. The second case is for the positions occupied by the tokens generated by  $\text{TF}_1$ , we do not care about them. The third case is for the tokens that  $\text{TF}_2$  will generate. It should believe that they are in their natural positions, i.e., shifted  $T_1(|\alpha|)$  positions left (the positions occupied by the tokens generated by  $\text{TF}_1$ ). We also need to change the flag used for choosing the distribution in the vectors at positions after  $|\alpha| + T_1(|\alpha|)$  so  $\text{TF}$  chooses the distribution corresponding to  $\text{TF}_2$  instead of  $\text{TF}_1$ .

At this point, the construction almost produces the tape of  $\text{TF}_1$  followed by the tape of  $\text{TF}_2$  except for one token. We need to do something special for the first token generated by  $\text{TF}_2$ . The distribution corresponding to it is in the vector  $h_{|\alpha|}$  but the last vector of the first run computing  $\text{TF}_2$  tokens is  $h_{|\alpha|+T_1(|\alpha|)}$ , the one coming from the last token of  $\text{TF}_1$ . Therefore, before choosing the distribution, we need to copy the vector  $h_{|\alpha|}$  to  $h_{|\alpha|+T_1(|\alpha|)}$ . Notice that nothing more is needed, since the attention mechanism only attends to vectors

to the right, so  $h_{|\alpha|}$  is not affected by the vectors coming from the tokens produced by  $\text{TF}_1$ . This is done with the operation presented in Section 3.1.9.

With this last fix,  $\text{TF}$  generates the tape of  $\text{TF}_1$  followed by the tape of  $\text{TF}_2$ . We now show how to compute  $q$ , i.e., the disjunction of  $q_1$  and  $q_2$ . This is easily done making  $h_{|\alpha|+T_1(|\alpha|)+T_2(\alpha)}$  only attend to itself and to  $h_{|\alpha|+T_1(|\alpha|)}$ , i.e., to the vectors coming from  $q_1$  and  $q_2$ .

We add a flag in the token embedding of  $\text{TF}$  to recognize the decision tokens of  $\text{TF}_1$  and  $\text{TF}_2$ . The flag is in the coordinate  $f_c$ .

$$\text{TE}(\sigma)_{f_c} = \begin{cases} 1 & \text{if } \sigma = q_{\text{accept}} \\ 0 & \text{otherwise} \end{cases}$$

Then, adding  $h_{|\alpha|+T_1(|\alpha|)}$  to  $h_{|\alpha|+T_1(|\alpha|)+T_2(\alpha)}$ , we have 1 or 2 in the coordinate  $f_c$  iff  $\text{TF}_1$  or  $\text{TF}_2$  accepted  $\alpha$ . Now we make all of  $h_{|\alpha|+T_1(|\alpha|)+T_2(\alpha)}$  0, except the  $f_c$  coordinate (using Section 3.1.3).

Before going to the **OUTPUT** layer we move the value of  $f_c$  to the coordinate that will project to the probability of  $q_{\text{accept}}$  and add some number larger than 0 but smaller than 1 to the coordinate that will define the probability of  $q_{\text{reject}}$ . Therefore, the token with the maximum probability will be  $q_{\text{accept}}$  iff  $\text{TF}_1$  or  $\text{TF}_2$  accepted  $\alpha$ , and it will be  $q_{\text{reject}}$  in the other case.

Then,  $\text{TF}$  computes the union of the languages of  $\text{TF}_1$  and  $\text{TF}_2$ . The number of steps of **CoT** that it needs is exactly the ones made by  $\text{TF}_1$  plus the ones made by  $\text{TF}_2$  plus 1 extra for taking the disjunction, thus in  $O(\cdot)$  notation it takes the same as the maximum of  $\text{TF}_1$  and  $\text{TF}_2$ . The embedding dimension of  $\text{TF}$  is the same as the dimension of  $\text{TF}_1$  plus the dimension of  $\text{TF}_2$  plus some constant number of flags, thus in  $O(\cdot)$  notation it is the same as the maximum of  $\text{TF}_1$  and  $\text{TF}_2$ .

Therefore, the language computed by the family of transformers  $\{\text{TF}_n\}_{n \in \mathbb{N}}$  is in  $\text{CoT}[T(n), d(n)]$ . Thus,  $\mathcal{L}_1 \cup \mathcal{L}_2 \in \text{CoT}[T(n), d(n)]$ .

SC agregar un diagrama

## 4. MAIN RESULTS

In this chapter, we present our results on probabilistic transformers. We show that probabilistic transformers are more powerful than deterministic ones, as they are capable of simulating them. We prove that the probabilistic classes support error reduction by means of a Monte Carlo simulation. Furthermore, we also demonstrate that RCoT classes are closed under boolean operations.

Finally, we prove upper and lower bounds on the expressive power of probabilistic transformers. These theorems are extensions of [?]’s results for deterministic transformers.

### 4.1 Probabilistic can simulate det

In this section we show how to simulate a deterministic transformer with a probabilistic one. Thus, we obtain the following theorem.

**Theorem 4.1.1.**  $\text{CoT}[T(n), d(n)] \subseteq \text{RCoT}[T(n), d(n)]$  for all  $T(\cdot)$  and  $d(\cdot)$

Let TF be a deterministic transformer. We construct a probabilistic transformer  $\text{TF}'$  that accepts the same words in the same number of steps of CoT and with the same embedding size.

As per Section 3.2, we assume TF to be normalized. We define  $\text{TF}'$  exactly as TF, except for the token selector which by definition is the probabilistic. Now, TF and  $\text{TF}'$  behave exactly the same. Since TF is normalized, the distribution  $p$  generated in every step satisfies the property that there exists  $\sigma \in \Sigma$  such that:

$$p(\tau) = \begin{cases} 1 & \text{if } \tau = \sigma \\ 0 & \text{otherwise} \end{cases}$$

Then, sampling this distribution  $p$  always returns the token  $\sigma$ , the same token that argmax would return. Next, the deterministic token selector and the probabilistic token selector behave the same. Thus, for all  $i$  and  $\alpha$ ,  $\text{TF}^i(\alpha) = \text{TF}'^i(\alpha)$ . In particular, this means that  $\text{TF}^*(\alpha) = \text{TF}'^*(\alpha)$ . Therefore, TF and  $\text{TF}'$  accept the same words (and in the same number of steps). By construction, the embedding size of TF and  $\text{TF}'$  is the same.

Finally, let  $\mathcal{L} \in \text{CoT}[T(n), d(n)]$ . Then, there is a family of deterministic transformers  $\{\text{TF}_n\}_{n \in \mathbb{N}}$  of embedding size  $O(d(n))$  that after  $O(T(n))$  steps of CoT computes  $\mathcal{L}$ .

We define the family  $\{\text{TF}'_n\}_{n \in \mathbb{N}}$  of probabilistic transformers by applying the construction just presented to each  $\text{TF}_n$ . Then, the transformers in  $\{\text{TF}'_n\}_{n \in \mathbb{N}}$  have embedding size  $O(d(n))$  and compute  $\mathcal{L}$  using  $O(T(n))$  steps of CoT.

Therefore,  $\mathcal{L} \in \text{RCoT}[T(n), d(n)]$ , and consequently  $\text{CoT}[T(n), d(n)] \subseteq \text{RCoT}[T(n), d(n)]$ .

### 4.2 Normalization

Recall the definition of normalization for deterministic transformers. We define a deterministic transformer as normalized when it satisfies the conditions N1 and N2.

**RR:** no estoy seguro cómo llamar a esta sección

If we make a probabilistic transformer satisfy N2, then, as seen in Section 4.1, it becomes deterministic. This changes the behavior of the transformer, so we cannot claim that for every probabilistic transformer there exists another with the same characteristics that computes the same language and satisfies N2. Thus, for probabilistic transformers, we define normalization solely as the condition N1.

**Definition 4.2.1** (Normalized probabilistic transformer). We say a probabilistic transformer is normalized if it satisfies the following condition.

- **N1:** The matrix of the OUTPUT has to be a projection, i.e., it is a diagonal matrix with ones on its diagonal.

Observe that the algorithm presented in Section 3.2.1 for constructing equivalent transformers satisfying N1 in the deterministic setting also applies to probabilistic transformers.

### 4.3 Robustness of our Definition

Recall the definition of language recognition by probabilistic transformers. We say that a family of probabilistic transformers  $\{\text{TF}_n\}_{n \in \mathbb{N}}$  recognizes a language  $\mathcal{L}$  when

$$\alpha \in \mathcal{L} \iff \Pr[\text{TF}_{|\alpha|}^*(\alpha) = q_{\text{accept}}] > 2/3$$

The choice of the constant  $2/3$  may seem somewhat arbitrary. We now show that we can replace  $2/3$  with any constant in the interval  $[1/2, 1)$ .

Let  $\{\text{TF}_n\}_{n \in \mathbb{N}}$  be a family of probabilistic transformers such that  $\alpha \in \mathcal{L} \iff \Pr[\text{TF}_{|\alpha|}^*(\alpha) = q_{\text{accept}}] > c$  for some fixed  $c$  in the interval  $[1/2, 1)$  and show that is equivalent to a family  $\{\text{TF}'_n\}_{n \in \mathbb{N}}$  that answers correctly with probability  $d < 1$ .

For each  $\text{TF}_n$  we construct  $\text{TF}'_n$  such that for all words of length  $n$ ,  $\alpha \in \mathcal{L} \iff \Pr[\text{TF}_n'^*(\alpha) = q_{\text{accept}}] > d$ . First, if  $d \leq c$ , then  $\text{TF}_n = \text{TF}'_n$ . Let us explore the case where  $c < d$ .

We construct  $\text{TF}'_n$  based on  $\text{TF}_n$  using a Monte Carlo simulation. The idea is similar to the disjunction of languages.  $\text{TF}'_n$  runs the transformer  $\text{TF}_n$  sequentially  $r$  times and then analyzes whether the majority of the results accept.

The construction presented in Section 3.3.2 for the disjunction of two deterministic transformers can be easily extended to run a larger number of transformers sequentially. This construction also applies for probabilistic transformers.

To compute the disjunction, we analyzed the coordinate  $f_c$  at the last step of the computation. It worked as a counter of how many transformers accepted the word. We replicate this strategy to check whether the majority of the  $r$  runs of  $\text{TF}_n$  accepted the word, i.e.,  $f_c > r/2$ . If that is the case, we force  $\text{TF}'_n$  to accept; otherwise, we force it to reject.

The question is how many runs of  $\text{TF}_n$  are needed, given that it answers correctly with probability  $c > 1/2$ , in order to obtain the correct answer with probability at least  $d$ . We estimate the necessary  $r$  using Chernoff's bound.

**Theorem 4.3.1** (Chernoff's bound). *Let  $y_1, \dots, y_r$  be independent Boolean random variables with  $\Pr[y_i = 1] = c$  for every  $1 \leq i \leq r$ . Let  $Y = \sum_{i=1}^r y_i$  and let  $\delta \in (0, 1)$ . Then  $\Pr[Y \leq (1 - \delta)rc] \leq e^{-\frac{\delta^2}{4}rc}$  and  $\Pr[Y \geq (1 + \delta)rc] \leq e^{-\frac{\delta^2}{4}rc}$ .*

We need to find  $\delta$  and large enough  $r$  such that  $\Pr[Y > r/2] > d$ . Let us analyze how the bound applies to our case. Each  $y_i$  represents a run of  $\text{TF}_n$ .  $y_i = 1$  indicates that the  $i$ th run of  $\text{TF}_n$  answered correctly, and  $y_i = 0$  that it answered incorrectly. Since  $\text{TF}_n$  answers correctly with probability higher than  $c$ , we have  $\Pr[y_i = 1] > c$ .

Applying Chernoff's bound, we require  $\delta$  and  $r$  such that  $(1+\delta)rc > r/2$  and  $e^{-\delta^2 rc/4} \geq d$ .

$$\begin{aligned} (1+\delta)rc &> \frac{r}{2} & e^{-\frac{\delta^2}{4}rc} &\geq d \\ (1+\delta)c &> \frac{1}{2} & -\frac{\delta^2}{4}rc &\geq \log_e(d) \\ c + \delta c &> \frac{1}{2} & r &\geq -4 \frac{\log_e(d)}{c\delta^2} \\ \delta &> \frac{1/2 - c}{c} \end{aligned}$$

Since  $c \geq 1/2$ , we can bound  $(1/2 - c)/c$  by 0; thus, choosing any  $\delta \in (0, 1)$  satisfies the first equation. Then, with a fixed  $\delta$ , choosing  $r = -4 \log_e(d)/(c\delta^2)$  makes  $\Pr[Y > r/2] > d$ . Therefore,  $\text{TF}'_n$  runs  $\text{TF}_n$   $\lceil -4 \log_e(d)/(c\delta^2) \rceil$  times and accepts the word if more than half of the runs accepted; otherwise, it rejects. Finally,  $\text{TF}'_n$  answers correctly with probability at least  $d$ .

For  $|\alpha| = n$ , let  $T(n)$  be the number of steps the family  $\{\text{TF}_n\}_{n \in \mathbb{N}}$  needs to compute  $\mathcal{L}(\alpha)$  and  $d(n)$  its embedding dimension. Then, the family  $\{\text{TF}'_n\}_{n \in \mathbb{N}}$  takes  $O(r \cdot T(n))$  steps to compute  $\mathcal{L}(\alpha)$  and uses embedding size  $O(r \cdot d(n))$ . Thus, the  $\text{RCoT}[T(n), d(n)]$  classes are invariant to the choice of the constant in the definition of language recognition by probabilistic transformers.

#### 4.4 Closed by Boolean Operations

We showed in Section 3.3 that deterministic transformers can perform boolean operations with only a constant overhead; thus,  $\text{CoT}[T(n), d(n)]$  classes are closed under boolean operations. Similar constructions apply to probabilistic transformers. Therefore, we show that  $\text{RCoT}[T(n), d(n)]$  classes are closed under complement and disjunction, which together imply closure under boolean operations.

**Theorem 4.4.1** ( $\text{RCoT}[T(n), d(n)]$  is closed under complement).  $\mathcal{L} \in \text{RCoT}[T(n), d(n)] \implies \mathcal{L}^c \in \text{RCoT}[T(n), d(n)]$

Let  $\mathcal{L} \in \text{RCoT}[T(n), d(n)]$  and let  $\{\text{TF}_n\}_{n \in \mathbb{N}}$  be the family of probabilistic transformers that computes  $\mathcal{L}$ . Let us assume all the transformers in the family are normalized. Then, similarly to the deterministic case, we define the family  $\{\text{TF}'_n\}_{n \in \mathbb{N}}$  as  $\{\text{TF}_n\}_{n \in \mathbb{N}}$ , except the OUTPUT matrices have the rows corresponding to  $q_{\text{accept}}$  and  $q_{\text{reject}}$  swapped.

Next, for all  $\alpha \in \Sigma^*$ , if  $\Pr[\text{TF}_{|\alpha|}^*(\alpha) = q_{\text{accept}}] > 2/3$ , then  $\Pr[\text{TF}'_{|\alpha|}(\alpha) = q_{\text{reject}}] > 2/3$ , since every time  $\text{TF}_{|\alpha|}$  prints  $q_{\text{accept}}$ ,  $\text{TF}'_{|\alpha|}$  prints  $q_{\text{reject}}$ . Therefore,  $\{\text{TF}'_n\}_{n \in \mathbb{N}}$  computes  $\mathcal{L}^c$  using the same number of CoT steps and the same embedding size  $O(d(n))$  as  $\{\text{TF}_n\}_{n \in \mathbb{N}}$  to compute  $\mathcal{L}$ . Thus,  $\mathcal{L} \in \text{RCoT}[T(n), d(n)]$  implies  $\mathcal{L}^c \in \text{RCoT}[T(n), d(n)]$ .

**Theorem 4.4.2** ( $\text{RCoT}[T(n), d(n)]$  is closed under disjunction).  $\mathcal{L}_1, \mathcal{L}_2 \in \text{RCoT}[T(n), d(n)] \implies \mathcal{L}_1 \cup \mathcal{L}_2 \in \text{RCoT}[T(n), d(n)]$

Let  $\{\text{TF}_n^1\}_{n \in \mathbb{N}}$  be the family of probabilistic transformers recognizing  $\mathcal{L}_1$  and  $\{\text{TF}_n^2\}_{n \in \mathbb{N}}$  the one recognizing  $\mathcal{L}_2$ . We construct the family  $\{\text{TF}_n\}_{n \in \mathbb{N}}$  following the same construction as for the disjunction of deterministic transformers (Section 3.3.2). By the result of Section 4.3, we can assume that the probability of error of each family is arbitrary small. Let  $p_1$  be the probability of  $\{\text{TF}_n^1\}_{n \in \mathbb{N}}$  answering correctly and  $p_2$  the probability of  $\{\text{TF}_n^2\}_{n \in \mathbb{N}}$  answering correctly. Then, for each word  $\alpha$ , we have four possible cases. In Figure 4.1, we show the probability of  $\text{TF}^1$  and  $\text{TF}^2$  accepting or rejecting  $\alpha$  in each of the cases.

|                       | $\text{TF}^1$ accepts | $\text{TF}^1$ rejects |                       | $\text{TF}^1$ accepts | $\text{TF}^1$ rejects |
|-----------------------|-----------------------|-----------------------|-----------------------|-----------------------|-----------------------|
| $\text{TF}^2$ accepts | $p_1 p_2$             | $(1 - p_1) p_2$       | $\text{TF}^2$ accepts | $p_1 (1 - p_2)$       | $(1 - p_1) (1 - p_2)$ |
| $\text{TF}^2$ rejects | $p_1 (1 - p_2)$       | $(1 - p_1) (1 - p_2)$ | $\text{TF}^2$ rejects | $p_1 p_2$             | $(1 - p_1) p_2$       |

Case 1:  $\alpha \in \mathcal{L}_1 \wedge \alpha \in \mathcal{L}_2$

Case 2:  $\alpha \in \mathcal{L}_1 \wedge \alpha \notin \mathcal{L}_2$

|                       | $\text{TF}^1$ accepts | $\text{TF}^1$ rejects |                       | $\text{TF}^1$ accepts | $\text{TF}^1$ rejects |
|-----------------------|-----------------------|-----------------------|-----------------------|-----------------------|-----------------------|
| $\text{TF}^2$ accepts | $(1 - p_1) p_2$       | $p_1 p_2$             | $\text{TF}^2$ accepts | $(1 - p_1) (1 - p_2)$ | $p_1 (1 - p_2)$       |
| $\text{TF}^2$ rejects | $(1 - p_1) (1 - p_2)$ | $p_1 (1 - p_2)$       | $\text{TF}^2$ rejects | $(1 - p_1) p_2$       | $p_1 p_2$             |

Case 3:  $\alpha \notin \mathcal{L}_1 \wedge \alpha \in \mathcal{L}_2$

Case 4:  $\alpha \notin \mathcal{L}_1 \wedge \alpha \notin \mathcal{L}_2$

Fig. 4.1: Probabilities of all possible behaviors of  $\text{TF}^1$  and  $\text{TF}^2$

TF accepts whenever  $\text{TF}^1$  or  $\text{TF}^2$  accepts. Let us analyze case by case the probability of TF answering correctly.

- Case 1:  $p_1 + p_2 - p_1 p_2 = p_1 p_2 + (1 - p_1) p_2 + p_1 (1 - p_2)$
- Case 2:  $1 - p_2 + p_1 p_2 = p_1 (1 - p_2) + (1 - p_1) (1 - p_2) + p_1 p_2$
- Case 3:  $1 - p_1 + p_1 p_2 = (1 - p_1) p_2 + p_1 p_2 + (1 - p_1) (1 - p_2)$
- Case 4:  $p_1 p_2$

Then, if  $p_1 = p_2 = 4/5$ , the probability of TF answering correctly in each case is larger than  $1/2$ .

- Case 1:  $24/25 > 1/2$
- Case 2:  $21/25 > 1/2$
- Case 3:  $21/25 > 1/2$
- Case 4:  $16/25 > 1/2$

Thus, the probability of  $\{\text{TF}_n\}_{n \in \mathbb{N}}$  answering  $\mathcal{L}_1 \cup \mathcal{L}_2$  correctly is at least  $1/2$ . Therefore, by the construction in Section 4.3, there is a family of transformers in  $\text{RCoT}[T(n), d(n)]$  that recognizes  $\mathcal{L}_1 \cup \mathcal{L}_2$  in  $\text{RCoT}[T(n), d(n)]$ .

## 4.5 Upper Bound

## 4.6 Lower Bound

## 5. CONCLUSIONS